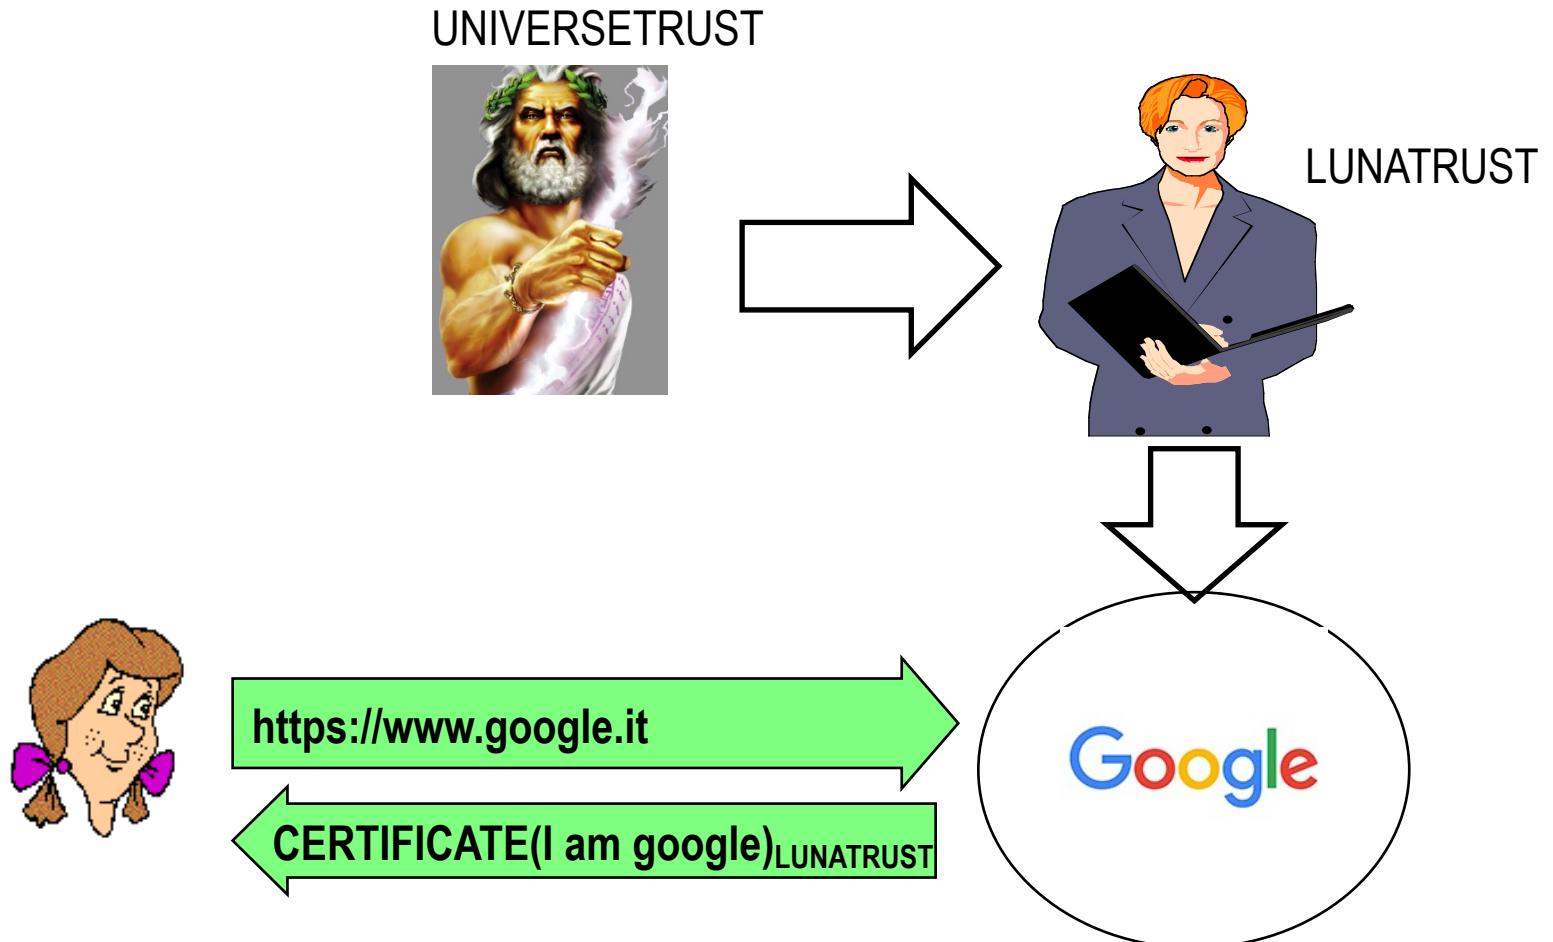


The failure of PKI → Certificate Transparency

(with insights on hash based data structures, specifically Merkle trees)

a **VERY SERIOUS** real world
problem: fake certificates

Web security pillar: Certificate Authorities ARE trusted!



Fact: trusted CA assumption at stake

→ <Google, PK>
↳ non di Google

Fake SSL certificates pose threats, say



Home News Anti-Phishing

Fake SSL certificates pose threats, say

by George Leopold

Published: 09 Jan 2015

Ex
re
SS
at

Netcraft has found dozens of fai
Some of these certificates may
customers. Successful attacks
and forwarding it to the bank.
authentication credentials, or r

The fake certificates bear com
As the certificates are not sign

Google Online Security Blog

The latest news and insights from Google on security and safety on the Internet

Enhancing digital certificate security

Posted: Thursday, January 3, 2013



183



300



{

→ Google's VALID fake Certificates mistakenly (?) issued

→ by TurkTrust (2012), ANSSI France (2013), etc

→ Smaller CAs: compromised

- ⇒ Holland: Dgnotar
- ⇒ Malaysia: DigiCert sdn. Bhd.
- ⇒ etc

Engineer

ected and blocked an unauthorized digital certificate for the "*google.com" domain and found the certificate was issued by an intermediate certificate authority, a Turkish certificate authority. Intermediate CA certificates carry the full name of the organization that has one can use it to create a certificate for any website they wish to

certificate revocation metadata on December 25 to block that intermediate CA, other browser vendors. TURKTRUST told us that based on our information,

TLS Proxies: Friend or Foe?

Mark O'Neill, Scott Ruoti, Kent Seamons, Daniel Zappala
Brigham Young University
Department of Computer Science
Provo, UT 84602

Online banking apps for mobile devices is far from trivial, and mobile devices are not the only threat. iOS-based banking apps test the authenticity of SSL certificates through manual tests by Leibniz University. These may also be vulnerable if a user

Our actions add Chrome again in though connecti

Since our priority further discussion and careful consideration.

===== Giuseppe Bianchi :)

A serious problem

→Fake VALID certificates

- ⇒ Various small/moderate CA «releasing» valid certificates for major sites
 - Asked by governments? Mistakenly... lost?
 - Many cases:
 - google fake certificates issued by
 - » Turkey's Turktrust, December 2012
 - » ANSSI in France on December 2013
 - » etc

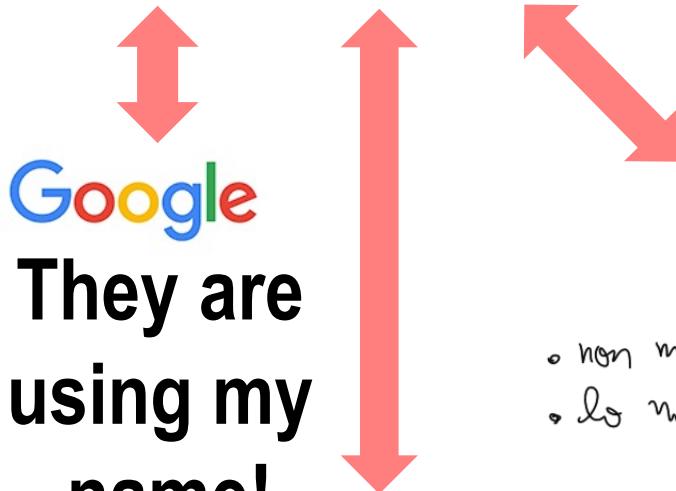
→Compromised CAs

- ⇒ Holland: DGnotar
- ⇒ Malaysia: DigiCert sdn. Bhd.
- ⇒ etc

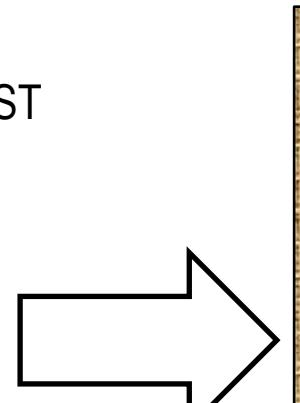
With powerful threats (governments), and many players which can «make mistakes»
the security of the PKI model is getting weaker and weaker

How to cope with malicious CAs?

Idea: **gigantic worldwide DB which anyone can check!**



UNIVERSETRUST



LUNATRUST



- non metto certificato in DB? non mi fido
- lo metto? Tutti lo vedono, e Google lo toglierà

- 1) validation certificate
- 2) proof SK
- 3) check certification in DB



<https://www.google.it>

Fake!!

CERTIFICATE(I am google)

LUNATRUST

OK, c'e' universetrust

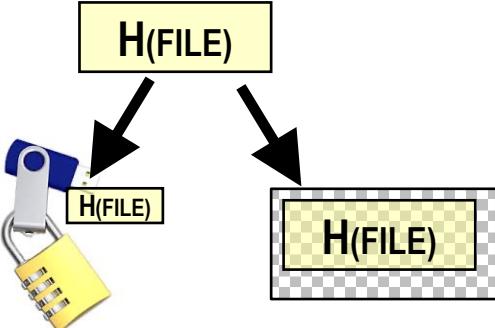


How to implement such gigantic Database of certificates?

A parenthesis: Merkle Trees

How to secure a file

(when you don't have a large secure storage available)



(HASH, SHA256)

→ Step 1: make “fingerprint” of the file

- ⇒ Via a cryptographic hash function
- ⇒ large file → small fingerprint!

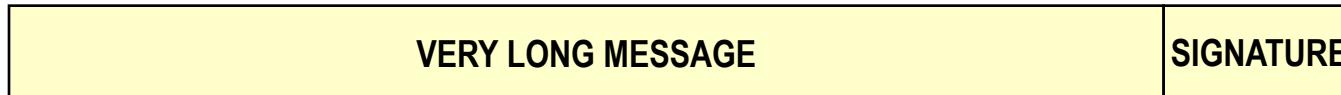
→ Step 2: “secure” the fingerprint!

(NON SERVE
FARLO SU TUTTO
IL FILE)

- ⇒ By copying it in a secure storage
- ⇒ ... or...
- ⇒ By digitally signing the hash fingerprint

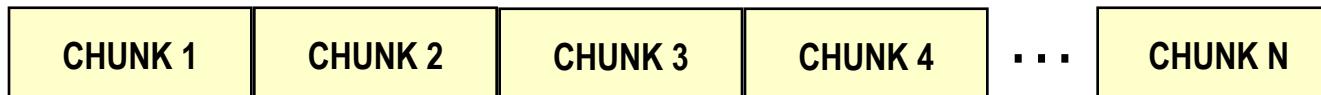
(divide in più pezzi e dato a persone diverse)

What about «chunked» messages?



→ Peer-to-peer delivery

- ⇒ Message divided into independent chunks
- ⇒ Frequently received out of order, and from distinct peers

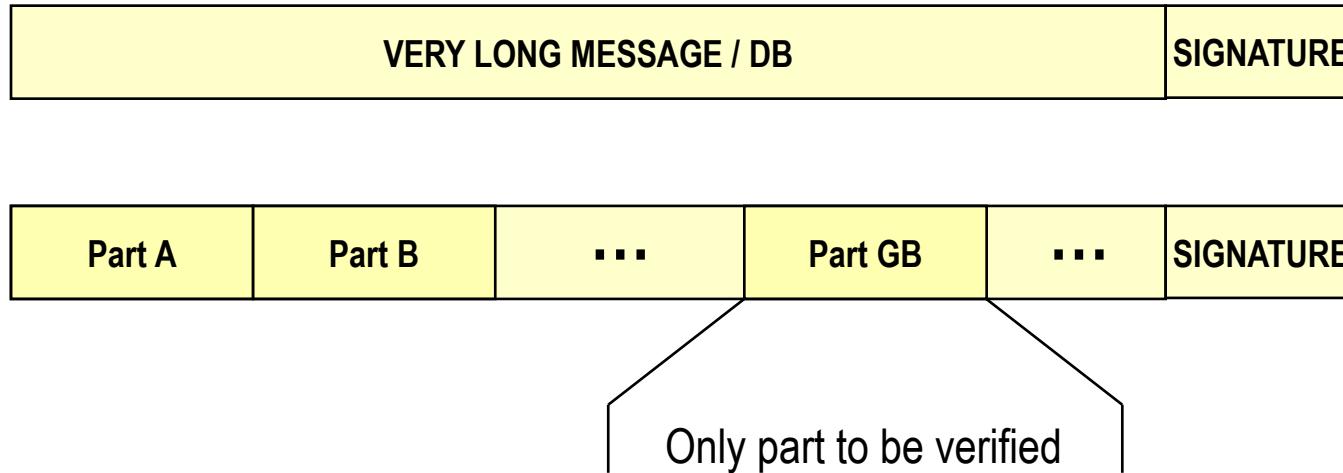


→ Signature verification: needs to wait until COMPLETE message reconstruction

- ⇒ But what about fake injected chunks?

Copisco solo alla fine

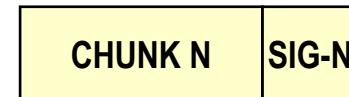
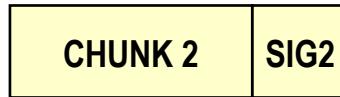
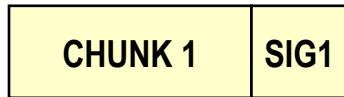
What about «partial» verification?



**Signature verification:
need to retrieve the whole database??**

Per-chunk signatures?

A chunk



→ OK but...

→ Computational overhead

⇒ Signature = asymmetric crypto = $\Theta(1000+)$ x hash computation cost

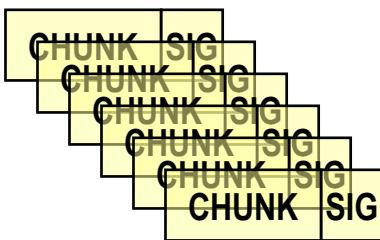
→ Storage overhead

⇒ RSA2048 = 256 bytes per signature

⇒ Bitcoin's ECDSA256 = 2×256 bit = 64 bytes

→ still quite a lot if you need MANY of them (small chunks)

→ No more integrity check for whole file (strip-type attacks)



Per-chunk signatures?

CHUNK 1	SIG1
---------	------

CHUNK 2	SIG2
---------	------

CHUNK N	SIG-N
---------	-------

→ OK but...

→ Computational overhead

⇒ Signature = asymmetric crypto = $\Theta(1000+)$ x hash computation cost

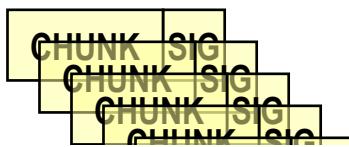
→ Storage overhead

⇒ RSA2048 = 256 bytes

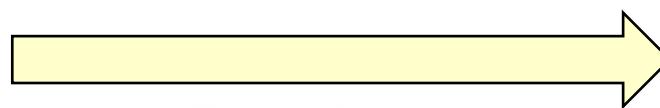
⇒ Bitcoin's ECDSA256 = 2×256 bit = 64 bytes

→ still quite a lot if you need MANY of them (small chunks)

→ No more integrity check for whole file (strip-type attacks)



Can we find
a better solution?



me riceve
di meno!

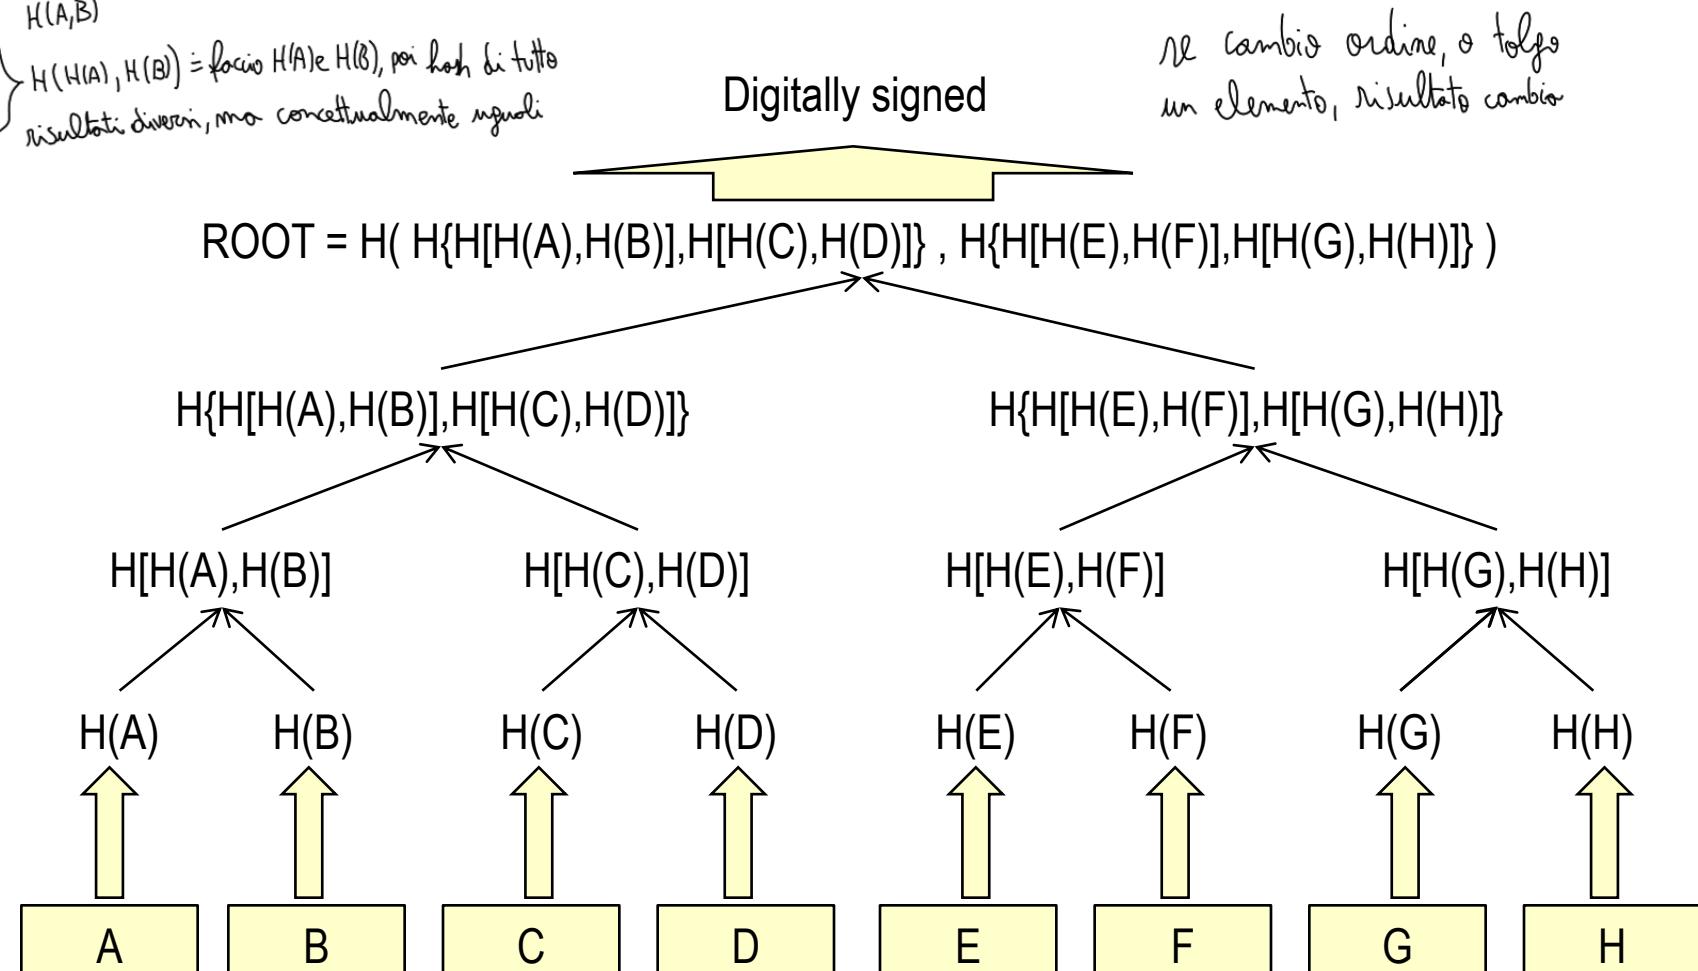
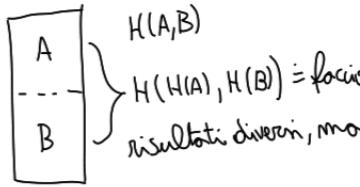


for riconoscere
alcuni
CHUNKS SIG

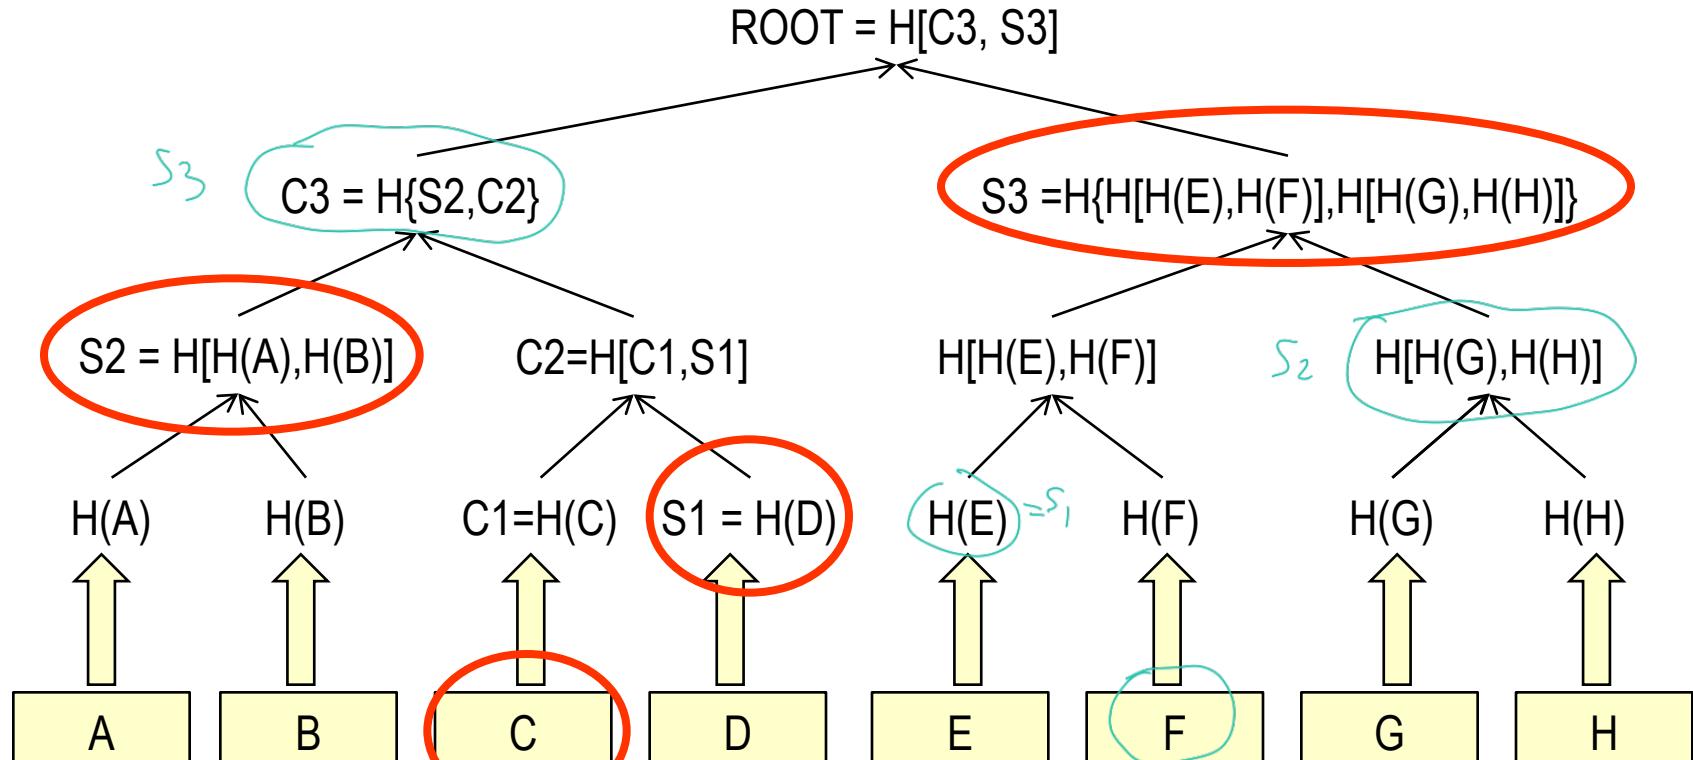


Ops... did I
receive all?

Merkle's idea: Hash tree (1979)



Single chunk verification: use “siblings”



dato C:

- 1) $H(C)$
 - 2) $H(H(C), S_1) = C_2$
 - 3) $H(S_2, C_2) = C_3$
 - 4) $H(C_3, S_3) = \text{ROOT}$
- $\xrightarrow{\text{si, ok}}$
- $\xrightarrow{\text{NO, errore nel DB}}$

C è modificato? dato C , e S_j cerchiati, uno ricomporre!!

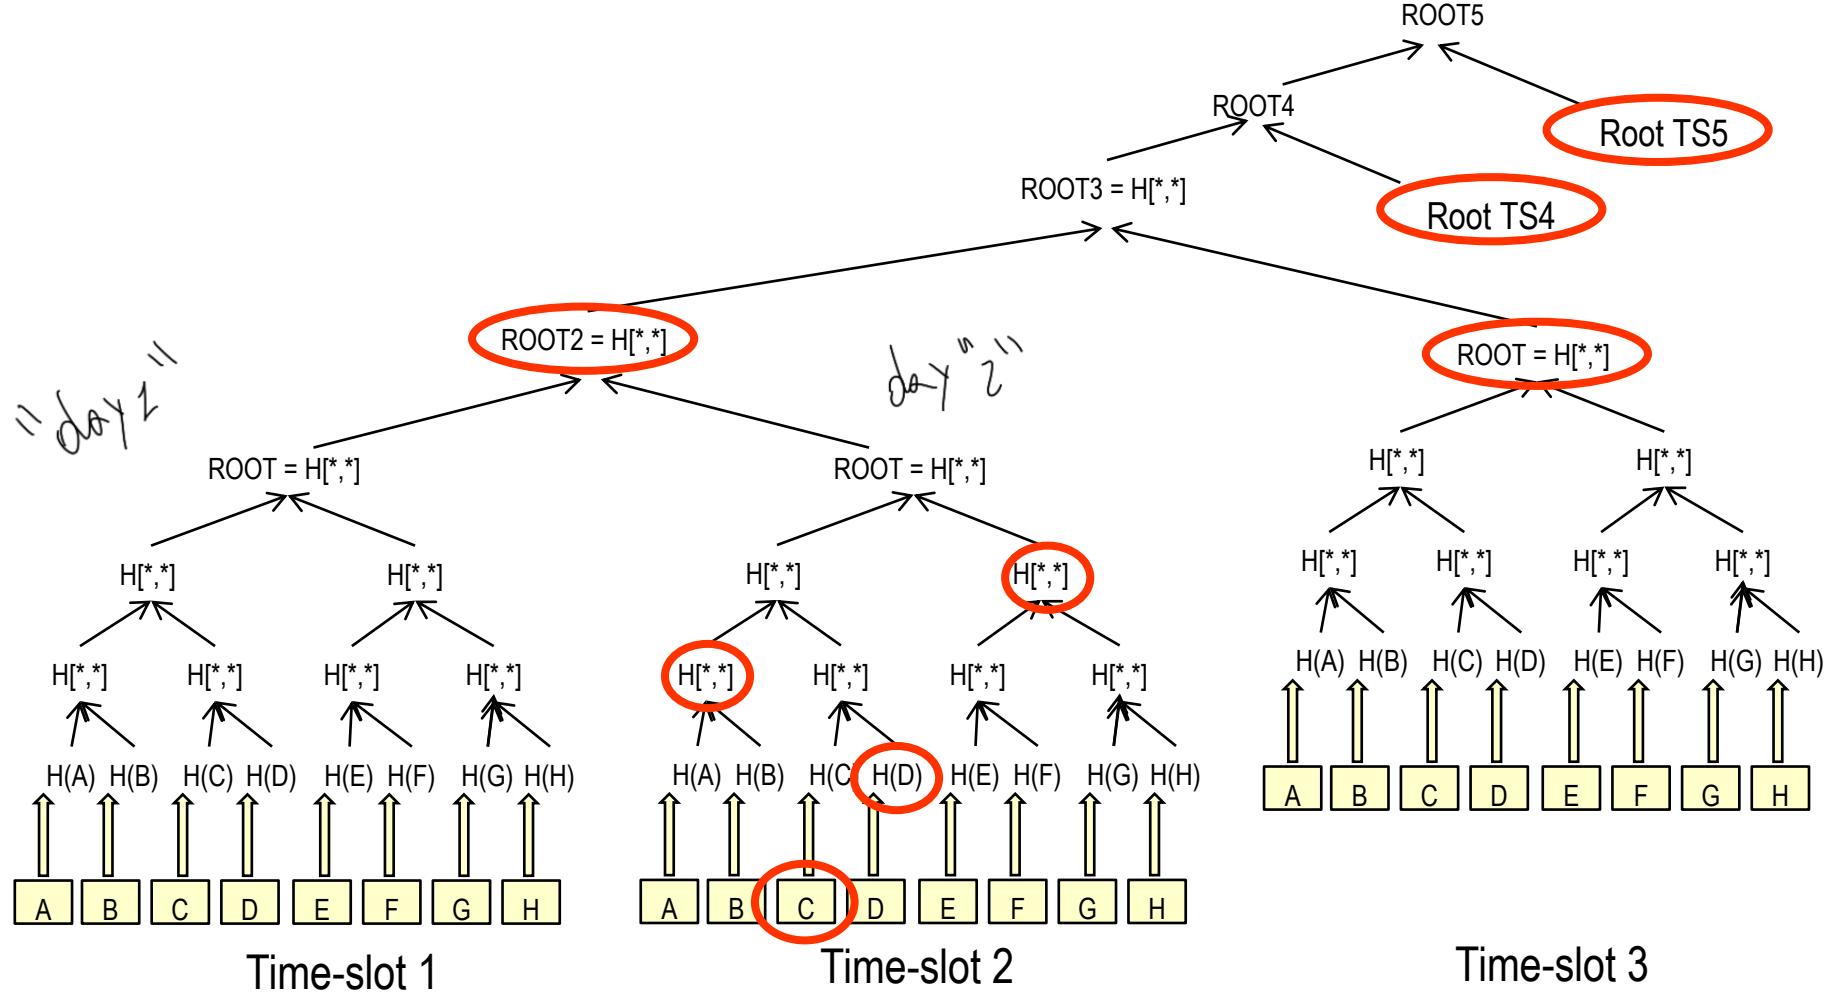
Log(N) siblings needed

per ottenere C con C' , dovei generare $S'_1, S'_2 \dots$ che va contro le proprietà della crypto hash

P

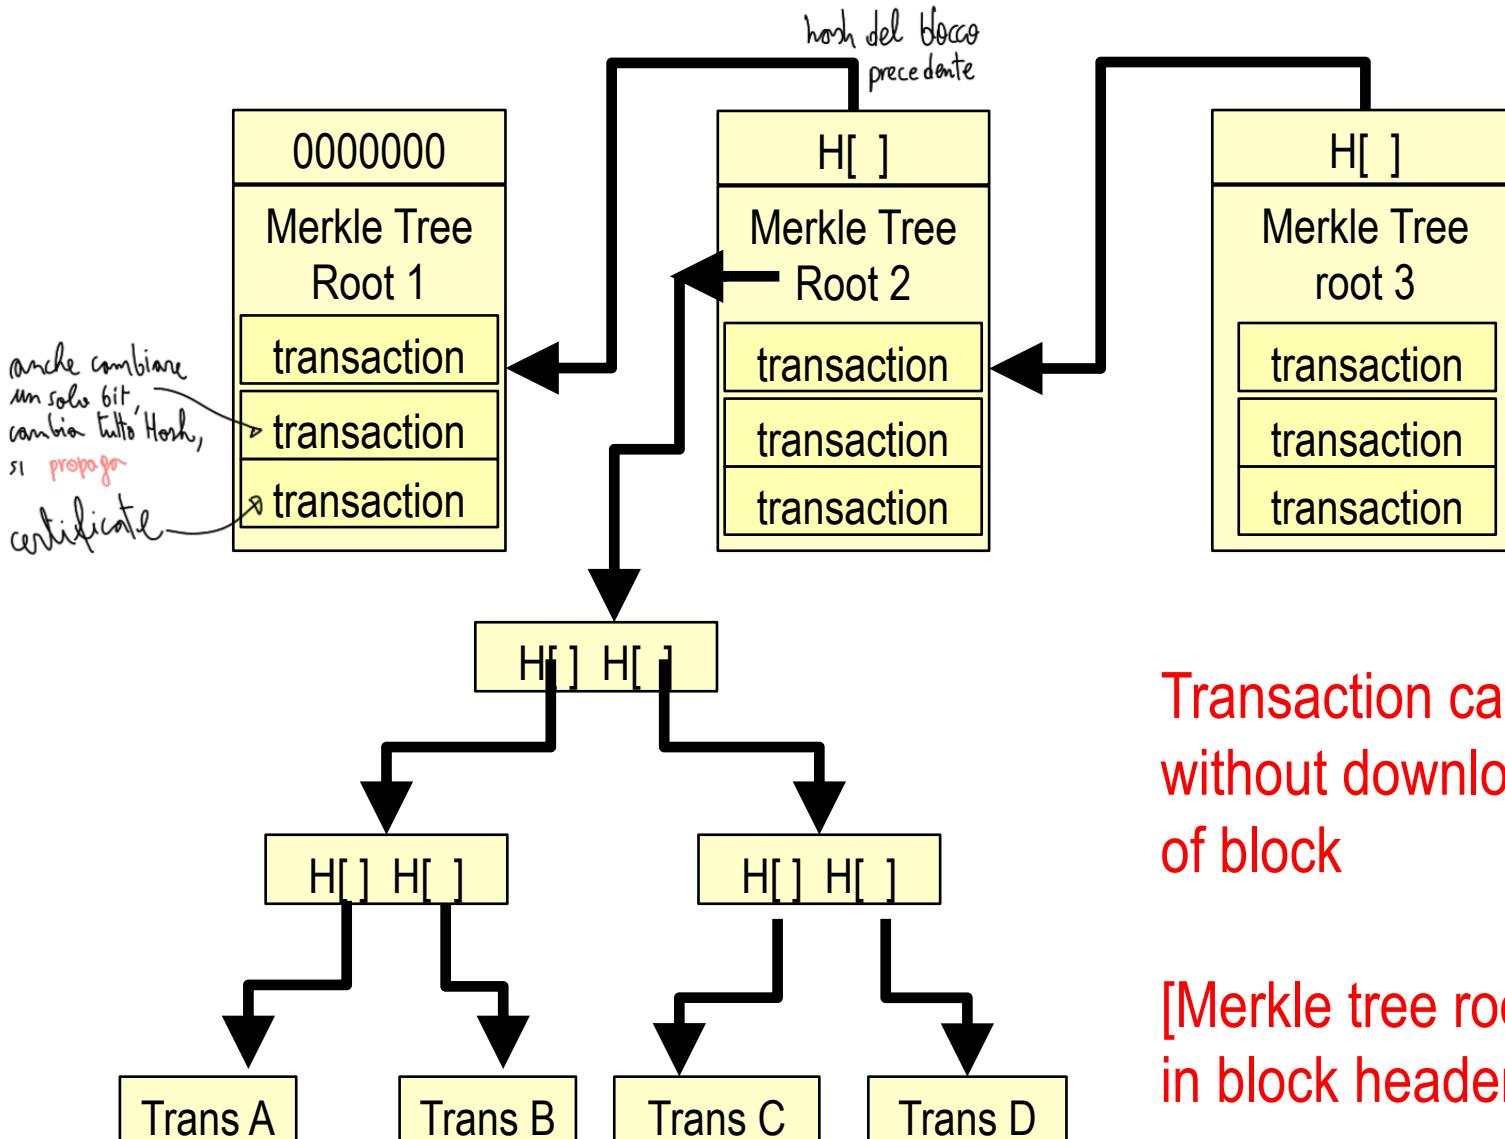
* esempio con F

Extending merkle's trees w. time



VALIDATION: Siblings for «your» tree + root of previous timeslot + roots of next timeslots

Merkle trees and blockchains



Transaction can be verified without downloading content of block

[Merkle tree root stored in block header]

Merkle Trees: Applications

→ Efficient storage in blockchains

- ⇒ Bitcoin, Ethereum, etc
- ⇒ Used in many parts of many blockchains!

→ Google's Certificate transparency

→ File Chunk validation

→ And many more

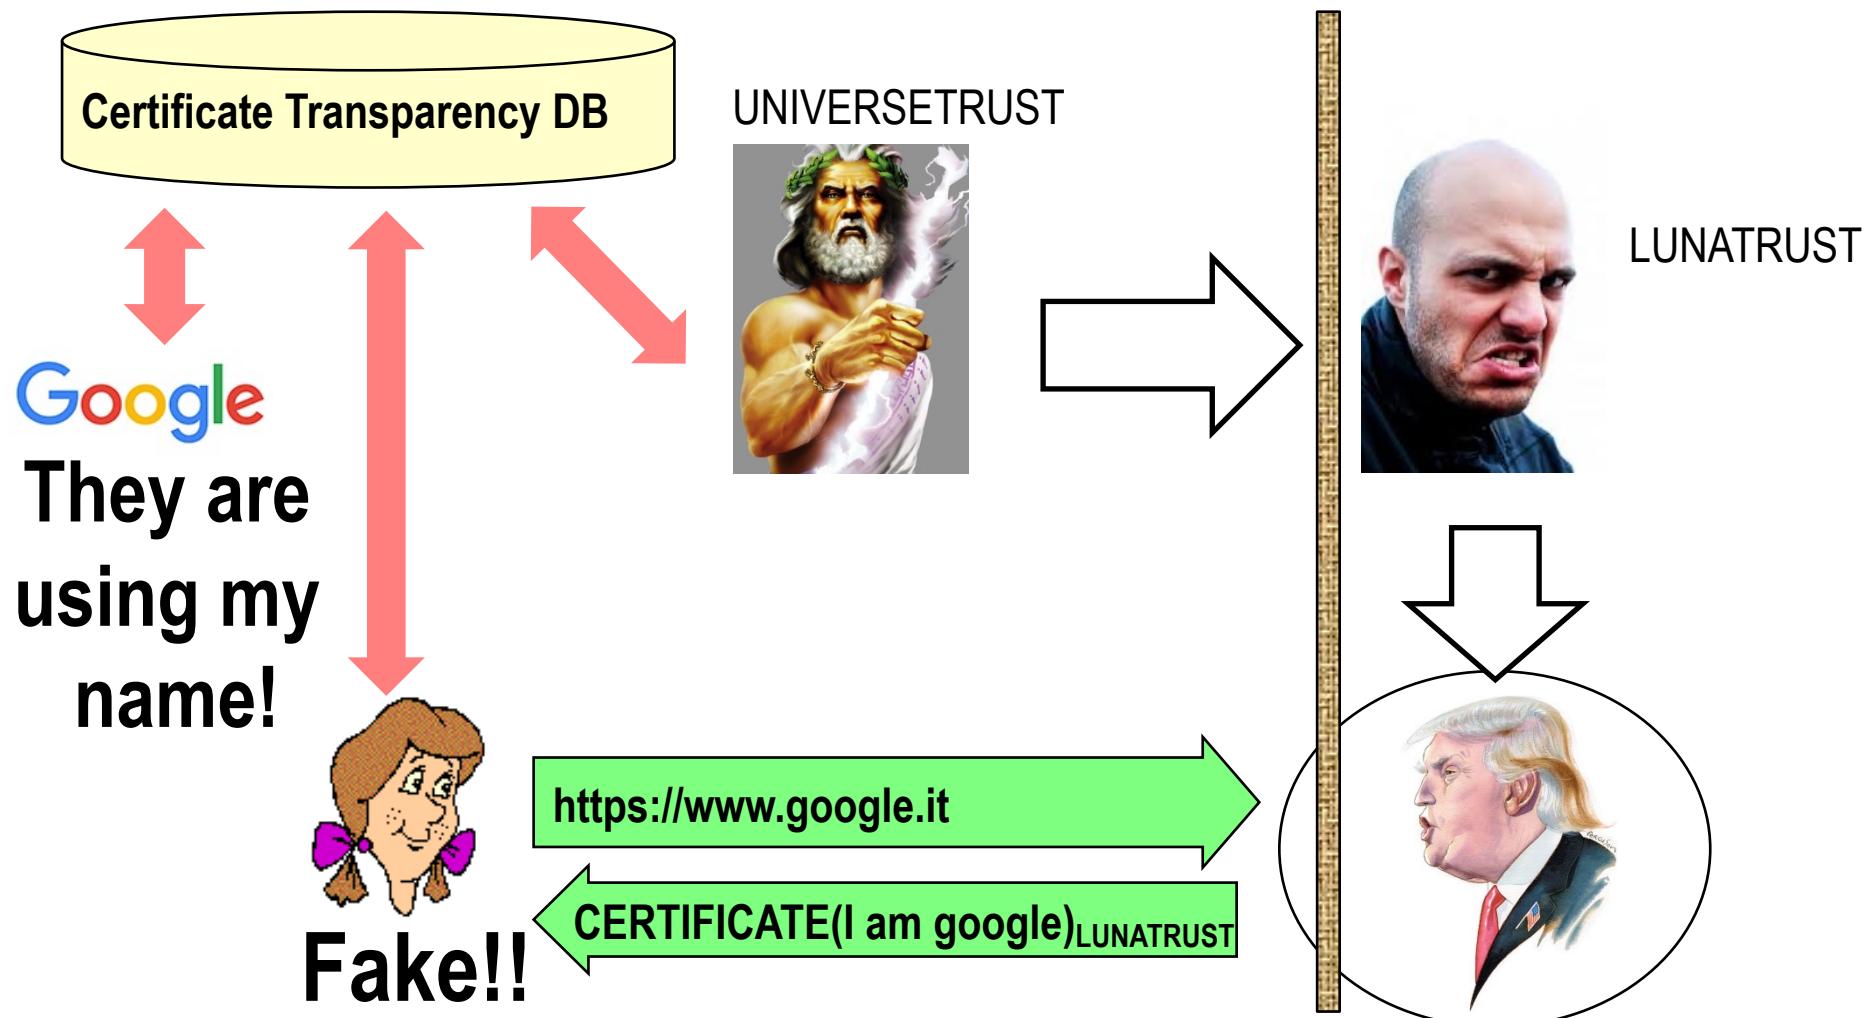
- ⇒ One-time signatures
- ⇒ Node authentication
- ⇒ Memory & storage authentication
- ⇒ Signature spreading
- ⇒ Etc etc ...

Merkle Trees' application:

Certificate Transparency

How to cope with malicious CAs?

Idea: gigantic worldwide DB which anyone can check!



Why certificate transparency?

- Make it almost impossible for a CA to issue a SSL certificate for a domain without the certificate being visible to the owner of that domain.
- Idea: put all certificates in a gigantic public list, that everyone can access and consult
- Further idea: make sure that we all see the “same” list!!
 - ⇒ New threat: “split world”: the public list you see is different (and crafted!) from what the rest of the world sees

Certificate Transparency in a nutshell

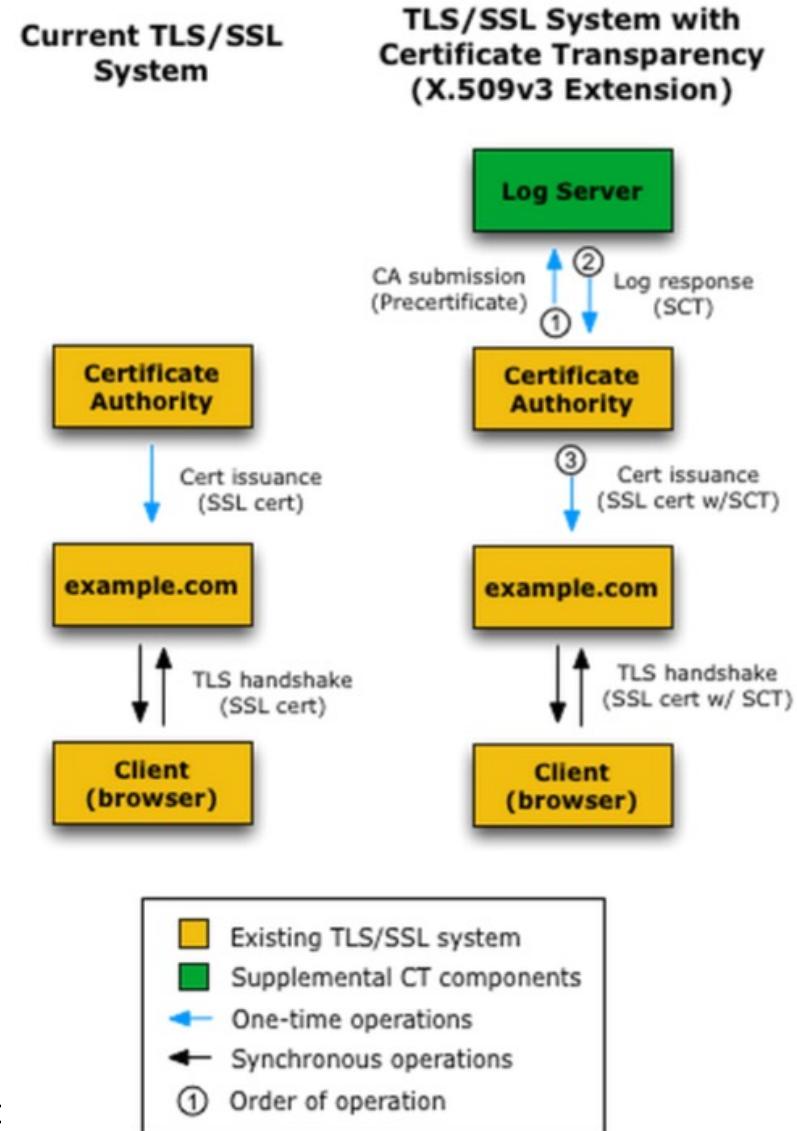
→ When issuing certificate: add it to the log server

⇒ Returns a Signed Certificate Timestamp (SCT)

→ During TLS handshake:

⇒ Check certificate validity (as usual)

⇒ Follow up with an additional Extended Validation (EV) by verifying that the certificate is in the public list



Certificate Transparency

→ Google's initiative

⇒ started 2013

→ First stated in experimental RFC 6962

⇒ July 2013

→ Endorsed and supported by IETF WG «trans»

⇒ Goal: standardize CT

⇒ RFC 6962bis released in December 2014

→ Currently integrated in major browsers

⇒ Chrome 1+, IE7+, Firefox 3+, Safari 3.2+, Opera 9.5+

⇒ Although not yet consistently used by all web sites

⇒ Plan to improve support for Extended Validation by 2015

→ Details @

<http://www.certificate-transparency.org/>

CT: how it works?

→Uses Merkle trees!

⇒ Very scalable, $\log(n)$ verification complexity

→A special type (sometimes called chron tree)

⇒ «append only» Merkle tree

⇒ Timestamped entries

→ Cannot retroactively insert a certificate

Following figures: source <http://www.certificate-transparency.org/log-proofs-work>

CT: how it works?

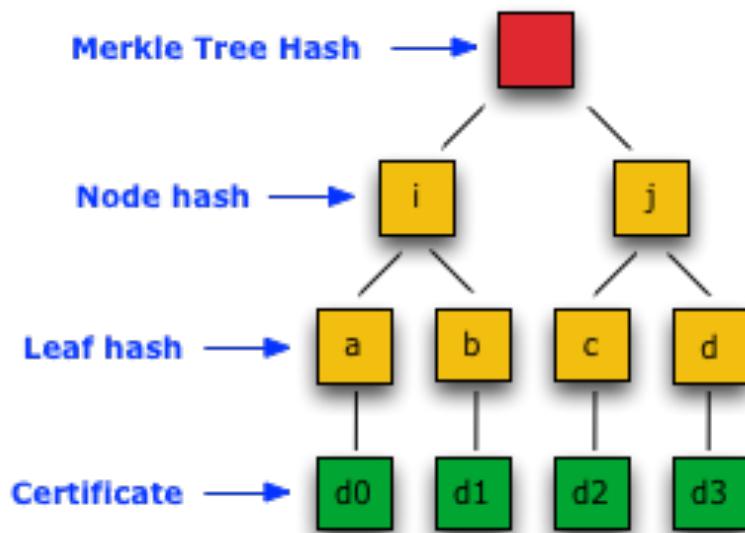


Figure 1

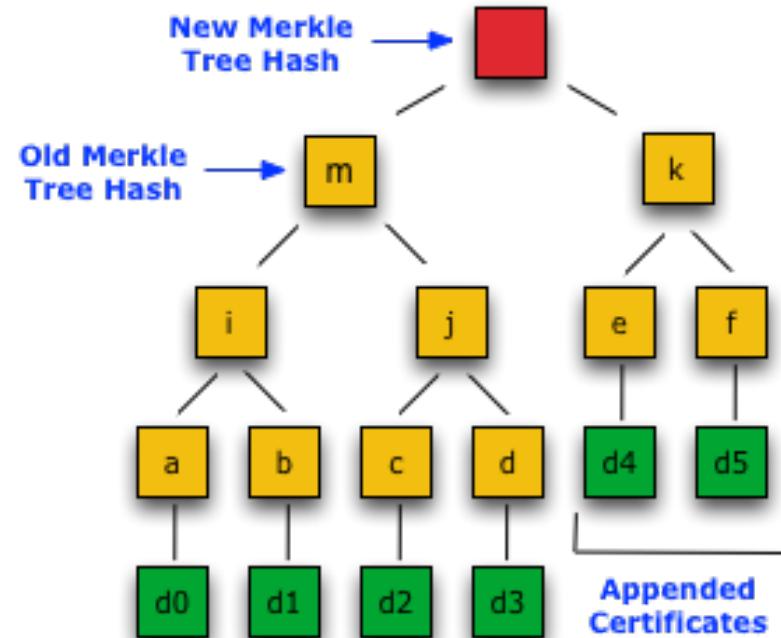


Figure 2

Every X hours, tree updated (currently X=24h)

Merkle Consistency proof

Verify that two versions of a log are consistent (i.e., the later includes everything of the earlier).
Consistency proof → no certificates have been modified or back-dated and inserted into the log.
Proof = need only root of the subtrees – includes past “root”

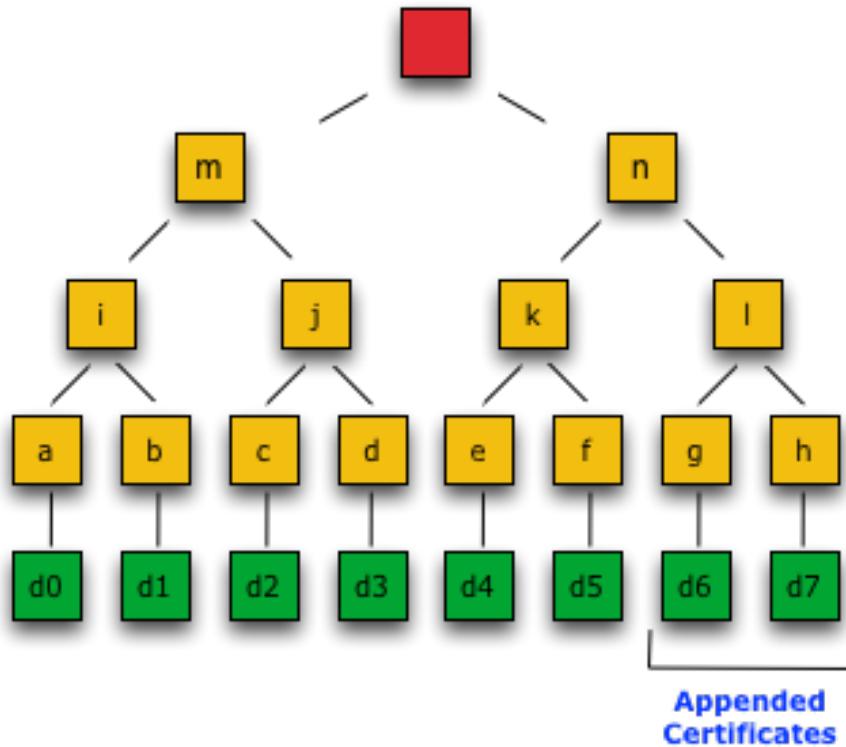


Figure 3

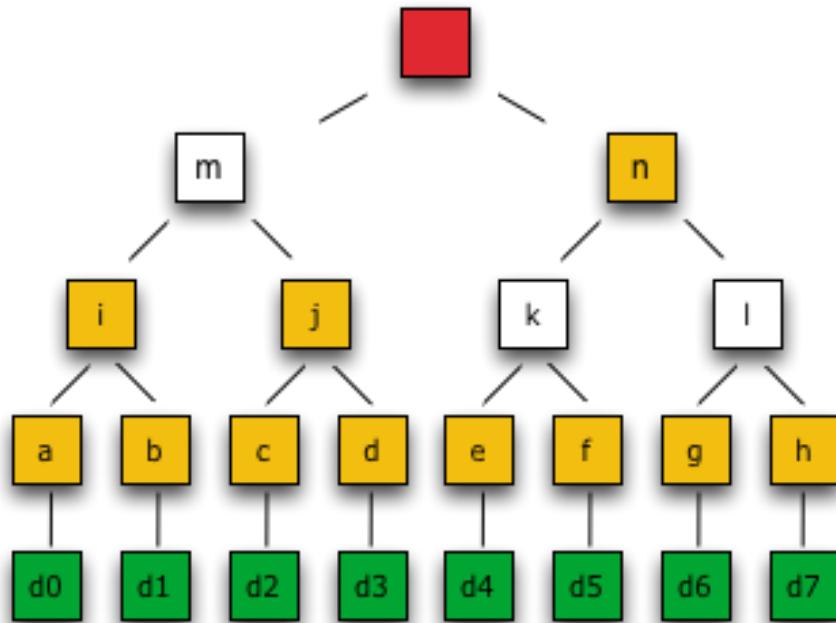


Figure 4

Merkle Audit Proof

→ Usual: via certificate + siblings

→ Siblings obtained from log

- ⇒ Issue: online operation, visible to all...
- ⇒ Remember who are the threats today...

→ Alternative: use «gossiping» protocols

- ⇒ Work in progress @ 2015!
- ⇒ No time for (very nice) details...

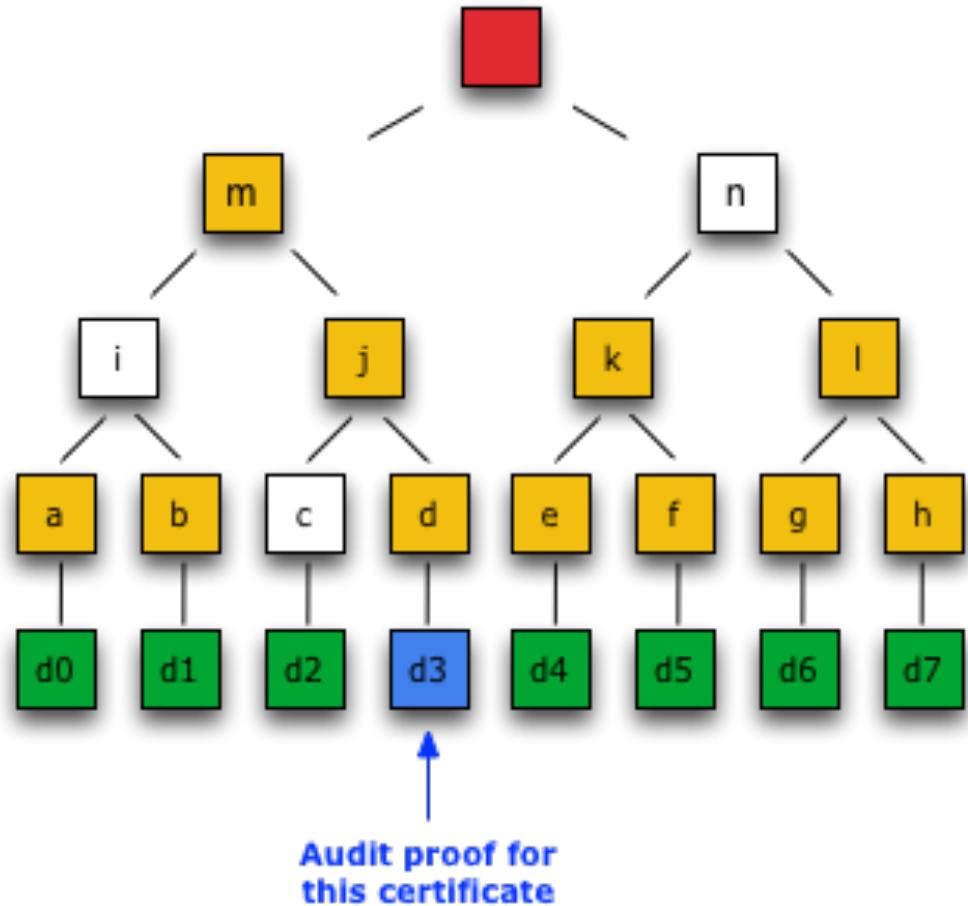


Figure 5

EV in the real world

PayPal, Inc. [US] <https://www.paypal.com/it/webapps/mpp/home>

PayPal, Inc.
La tua connessione a questo sito è privata.

Autorizzazioni **Connessione**

L'identità di PayPal, Inc. a San Jose, California US è stata verificata da Symantec Class 3 EV SSL CA - G2 e può essere controllata pubblicamente.

[Informazioni certificato](#) [Informazioni sulla trasparenza](#)

La connessione a www.paypal.com è crittografata tramite crittografia obsoleta.

La connessione utilizza TLS 1.2.

La connessione è stata crittografata utilizzando AES_256_CBC, con SHA1 per l'autenticazione dei messaggi e RSA come meccanismo principale di scambio delle chiavi.

i Informazioni sito
Non hai mai visitato questo sito prima di oggi.

[Che cosa significano?](#)

il nostro uso dei **cookie**, per una tua migliore esperienza di navigazione.

società indipendenti e non faranno più parte dello stesso gruppo. Per questo motivo, le rela

Visualizzatore timestamp certificato firmato

1: Integrato, Verificato

Stato di convalida	Verificato
Origine	Integrato
Versione	V1
Nome log	Google 'Pilot' log
ID log	A4 B9 09 90 B4 18 58 14 87 BB 13 A2 CC 67 70 0A 3C 35 98 04 F9 1B DF B8 E3 77 CD 0E C8 0D DC 10
Emesso alle ore	giovedì 23 aprile 2015 00:28:58
Algoritmo di hash	SHA-256
Algoritmo di firma	ECDSA
Dati della firma	30 45 02 20 5B E3 52 37 1B 84 B4 48 D1 CD 8F 53 34 D5 31 22 2D 46 C1 91 B2 86 6A 77 3E 37 DF FD CC 9D 78 A5 02 21 00 DE 96 80 20 B6 82 09 1E 8E 4A 9F C7 EE 3B 35 47 82 31 B1 D0 B1 63 F7 7D 52 1C C5 C5 41 5E BF 1A

Chiudi