

*Machine Learning*

# ML with Python: Setup & Introduction

Gabriele Russo Russo    Francesco Lo Presti

Laurea Magistrale in Ingegneria Informatica – A.Y. 2023/24



**TOR VERGATA**  
UNIVERSITÀ DEGLI STUDI DI ROMA

Dipartimento di Ingegneria Civile e  
Ingegneria Informatica

# Outline

- ▶ Environment setup
- ▶ NumPy & Pandas: Basics
- ▶ Scikit-learn and ML workflow example

# Environment Setup

# Requirements

We will use **Python** as the reference programming language for the course.

Software requirements:

- ▶ Python 3.x
- ▶ Some libraries, e.g.:
  - ▶ `scikit-learn`
  - ▶ `tensorflow`
- ▶ Jupyter (Lab)

# Jupyter Notebooks

- ▶ **Jupyter Notebook**: web application for creating computational documents (i.e., notebooks), which may contain Python code, text, figures
- ▶ **Jupyter Lab**: evolution of the original web app, providing a development environment to work with notebooks
- ▶ Pick the one you prefer

# Setup Methods

- ▶ Method 0: naive approach (not recommended)
- ▶ Method 1: venv
- ▶ Method 2: conda
- ▶ Method 3: Google Colab

## Method 0 (Not recommended)

- ▶ Install Python and Jupyter following the instructions for your OS
- ▶ Install the libraries you need using `pip`
  - ▶ e.g., `pip install scikit-learn`
- ▶ Difficult to install specific library versions, if required
- ▶ System upgrades can break your project

# Isolated Environments

- ▶ Isolated Python environments avoid the issue of conflicting Python/library versions across different projects
- ▶ Libraries installed within an environment not visible outside
- ▶ Option 1: Lightweight environments
  - ▶ A directory with all the install libraries
  - ▶ Default Python interpreter of the system is used
- ▶ Option 2: conda environments
  - ▶ Libraries + Python interpreter



# Method 1: venv

- ▶ Install Python following the instructions for your OS
- ▶ Create a virtual environment using `venv`
- ▶ Install the libraries you need using `pip`
- ▶ Jupyter can be either installed at system-level or using `pip` in the env.
- ▶ No third-party tools to install
- ▶ Same Python version across environments
- ▶ Python upgrades can break your environments

# Method 1: Example

```
# Create environment (in the current directory)
python -m venv my_env
# Activate it
source ./my_env/bin/activate
# Install libraries...
pip install scikit-learn
# ...
deactivate
```

# Requirements file

- ▶ It is a good practice to list the required libraries (and their version) in a `requirements.txt` file

## Example requirements.txt

```
matplotlib==3.6.1  
numpy==1.23.4  
packaging==21.3
```

- ▶ You can easily reinstall them with a single command

```
pip install -r requirements.txt
```

## Method 2: conda

- ▶ Install conda, picking one of the available distributions, e.g.:
  - ▶ miniforge
  - ▶ miniconda
  - ▶ ...
- ▶ Create an environment using conda
- ▶ Install libraries and Jupyter using `conda install`
- ▶ Each environment can use its own Python version and libraries

## Method 2: Example

```
# Create environment
conda create -n ENVNAME python=3.10
# Activate it
conda activate ENVNAME
# Install stuff
conda install jupyterlab
conda install scikit-learn

jupyter lab

conda deactivate
```

## Note: base environment

- ▶ When installing a conda distribution, a **base** environment is created
- ▶ By default, the base environment is activated automatically when no other env is in use
- ▶ If you want to disable this behavior (recommended):

```
conda config --set auto_activate_base false
```

# Importing/exporting environments

```
# Export to file
conda env export > ENV.yml

# Cross-platform export
conda env export --from-history > ENV.yml

# Import
conda env create -n ENVNAME --file ENV.yml
```

We will share a `env.yml` file that you can import

## Method 3: Google Colab

- ▶ **Google Colaboratory** (or, **Colab**) is a cloud-based platform to run Jupyter notebooks in your browser
- ▶ Usage (with some limitations) is free and only requires a Google account
- ▶ Colab can connect to Google Drive and GitHub to load notebooks
- ▶ No need to install anything on your PC, except for a browser
- ▶ Effortless access to GPUs
- ▶ `https://colab.research.google.com`



# NumPy & Pandas: Basics

# NumPy

NumPy (Numerical Python) provides data structures and functions to efficiently work on large N-dimensional numerical arrays. Key features:

- ▶ **ndarray** (N-dimensional array)
- ▶ functions to operate on ndarrays
- ▶ I/O functions

NumPy is conventionally imported as follows:

```
import numpy as np
```

# ndarray

- ▶ N-dimensional array of numbers (either integers or floats)
- ▶ Support for efficient and easy manipulation (you can work on vectors and matrices without loops)

```
import numpy as np  
A = np.array([1,2,5])  
print(A)  
print(A*2)
```

# Creating a ndarray

## Converting from a different collection (e.g., list)

```
A = np.array([1,2,3,4])  
M = np.array([[1,2,3], [7,8,9]])
```

## Specifying arbitrary dimensions

```
A = np.zeros((5,3)) # not zeros(5,3)!  
B = np.ones(15)
```

## From file

```
A = np.loadtxt("filename.txt")
```

# Basic operations

```
A = np.ones((2,2))  
B = np.array([[1,2],[0,1]])  
print(A+B)  
print(A+5*B)  
print(np.matmul(A,B.transpose()))
```

# shape

- ▶ The `shape` field is a tuple indicating the dimensions

```
A = np.array([[1,2], [4,5]])  
print(A.shape) # (2,2)
```

- ▶ Sometimes, it is useful to “reshape” a ndarray. For instance, if you have a 1-D array of length N, you may want to reshape it to a 1xN matrix.

```
A = np.array([1,2,3])  
B = np.array([[1,2,3]])  
print(np.matmul(A.T,B)) # error  
A = A.reshape(1,3)  
print(np.matmul(A.T,B)) # 3x3 matrix
```

# Indexing and slicing

You can easily access single elements or slices:

```
print(A[4])  
print(A[4:6])  
A[4:6] = -1.0  
A[:] = 5 # assign 5 to all elements  
# For a matrix:  
print(M[1,2]) # or M[1][2]  
print(M[:1, :])
```

# References

- ▶ Official docs:  
`https://numpy.org/doc/stable/user/index.html`



# Pandas

Pandas is a Python library providing data structures and functions to efficiently manipulate data. Key data structures:

- ▶ Series
- ▶ DataFrame

Differently from NumPy, Pandas can handle tabular data with heterogeneous data types (i.e., not limited to numerical values).

```
import pandas as pd
```

# Series

- ▶ **Series** is a one-dimensional labeled array capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.).
- ▶ The axis labels are collectively referred to as the **index**
  - ▶ e.g., in a time series, values might be labeled with a date
  - ▶ If not specified, the default index consists of integers from 0 to  $(N - 1)$  (where  $N$  is the number of values)
- ▶ Similar to NumPy's ndarray, but with labels
  - ▶ you can pass a Series to many functions acting on ndarrays

```
A = pd.Series([4, 7, 5], index=["a", "b", "c"])  
print(A)
```

# DataFrame

- ▶ **DataFrame** is a 2-dimensional labeled data structure with columns of potentially different types.
- ▶ You can think of a DataFrame as a dictionary of Series objects (each representing a column)
- ▶ Both rows and columns have a corresponding array of labels: **index** for rows, **columns** for columns

```
df = {"one": [1.0, 2.0, 3.0, 4.0],  
      "two": [4.0, 3.0, 2.0, 1.0]}  
pd.DataFrame(df)  
#    one  two  
#0  1.0  4.0  
#1  2.0  3.0  
# ...
```

# Value Selection

- ▶ Two fundamental functions to select values based on their integer location (`.iloc()`) or their label (`.loc()`)

```
df.columnName # -> Series
df["columnName"] # -> Series
df.loc[:, "columnName"] # -> Series
df.loc[:, ["colA", "colB"]] # -> DataFrame

df.iloc[3, :] # -> Series
df.iloc[3, 2:5] # -> DataFrame

df.loc[bool_vector] # -> DataFrame
df.loc[df.columnName > 3]
```

# Reading a CSV File

```
df = pd.read_csv("nomefile.csv")
```

- ▶ By default, first line is used to extract column names, unless `header=None` is specified
- ▶ `skiprows=N` allows to skip first  $N$  lines in the file
- ▶ `names=["a", ...]` allows manual column naming
- ▶ `index_col="a"` specifies the column to use as index

# scikit-learn and ML workflow

# Scikit-learn

- ▶ Scikit-learn is an open source ML library that supports supervised and unsupervised learning
- ▶ It also provides various tools for model fitting, data preprocessing, model selection, model evaluation, and many other utilities
- ▶ <https://scikit-learn.org/>
- ▶ We will see scikit-learn in action by presenting an end-to-end example ML project

# Scikit-learn: Key Concepts

- ▶ Scikit-learn provides dozens of built-in algorithms and models, called **estimators**, which implement a common interface
- ▶ Each estimator can be trained against some data using its **fit()** method
- ▶ Once the estimator is fitted, it can be used for predicting target values of new data, using the **predict()** method
- ▶ Data are provided to scikit-learn as array-like objects, usually consisting of Numpy's ndarrays or Pandas Series/DataFrame



# End-to-End ML Workflow Example

- ▶ We consider an example project, described in greater detail in Chapter 2 of “Hands-on ML with Scikit-learn, ...”
- ▶ Our task is **predicting median house prices** in California districts, based on the well-known *California Housing Prices* dataset
  - ▶ Data from 1990 California census
  - ▶ Actually, we will use a slightly simplified version

# Understanding the Task

- ▶ As a ML engineer, your first step is understanding the business needs and goals behind a project
- ▶ For instance, we are requested to predict house prices. How will be the prediction used?
  - ▶ e.g., to help potential buyers estimate prices
  - ▶ e.g., to help investors deciding whether to buy or not
- ▶ This allows you to understand how critical the task is and choose a suitable performance metric
  - ▶ e.g., in a classification task, are false positives less harmful than false negatives?


# Understanding the Task

- ▶ As a ML engineer, your first step is understanding the business needs and goals behind a project
- ▶ For instance, we are requested to predict house prices. How will be the prediction used?
  - ▶ e.g., to help potential buyers estimate prices
  - ▶ e.g., to help investors deciding whether to buy or not
- ▶ This allows you to understand how critical the task is and choose a suitable performance metric
  - ▶ e.g., in a classification task, are false positives less harmful than false negatives?
- ▶ We use a typical performance measure for regression

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (t^{(i)} - y^{(i)})^2}$$

# Hands-on Example

The example is presented in a Jupyter notebook.

 end2end.ipynb