

Machine Learning

Neural Networks and Deep Learning: Fundamentals

Gabriele Russo Russo Francesco Lo Presti

Laurea Magistrale in Ingegneria Informatica – A.Y. 2023/24



TOR VERGATA
UNIVERSITÀ DEGLI STUDI DI ROMA


Dipartimento di Ingegneria Civile e
Ingegneria Informatica

From Human Brain to ANNs

- ▶ Science and technology often inspired by nature
 - ▶ e.g., birds inspired us to build aircrafts
- ▶ From human brain to intelligent machines?
- ▶ Challenging task: computational power of brain much higher than modern computers
- ▶ Artificial neural network (ANN): popular ML technique that simulates the mechanism of learning in biological organisms


Biological Neurons

- ▶ Cell body + branching extensions (**dendrites**) + very long extension (**axon**)
- ▶ Axon splits off into branches (**telodendria**) at its extremity
- ▶ **Synaptic terminals** (or, **synapses**) at the tip of these branches, connected to dendrites or cell bodies of other neurons



Biological Neurons (2)


- ▶ Neurons produce short electrical impulses, which make synapses release chemical signals called **neurotransmitters**
- ▶ When a neuron receives a sufficient amount of specific neurotransmitters within a few milliseconds, it fires its own electrical impulses
- ▶ Single neurons are pretty simple, but they are organized in a network of billions
 - ▶ Each connected to 1,000+ neurons (often organized in layers)
 - ▶ Highly complex computations performed by such networks



History: McCulloch and Pitts (1943)


A Logical Calculus of Ideas Immanent in Nervous Activity

- ▶ Simplified computational model of how biological neurons might work (first ANN architecture)
- ▶ Neuron: 1+ binary inputs, 1 binary output
- ▶ Output activated when at least X inputs are active
- ▶ **Example:** activation with at least 2 active inputs




Perceptron

- ▶ Different model, proposed by F. Rosenblatt in 1957
- ▶ Inputs and outputs are real numbers
- ▶ Each input has an associated **weight**
- ▶ Output computed through an **activation function ϕ** (e.g., the sign function)




$$y = \phi\left(\sum_{j=1}^d w_j x_j\right) = \phi(\mathbf{w}^T \mathbf{x}) = \text{sgn}(\mathbf{w}^T \mathbf{x}) \quad (1)$$



Perceptron with Bias

- ▶ There is often an invariant part of the prediction, called **bias**
 - ▶ e.g., you need to predict a positive value when $x_j = 0 \forall j$
- ▶ We consider an additional input node that always transmits a constant value 1 with connection weight b (bias variable)







$$y = \phi\left(\sum_{j=1}^d w_j x_j + b\right) = \phi(\mathbf{w}^T \mathbf{x} + b) \quad (2)$$

Note: Hereafter, to simplify notation, we will often omit b in the equations

Perceptron for Classification

- ▶ Perceptron can be seen as a binary classifier
- ▶ It suffices to map the output values $\{-1, 1\}$ to the target classes



Perceptron: Training

Rosenblatt proposed a heuristic algorithm for training

- ▶ Weights initialized arbitrarily (e.g., $\mathbf{w} = 0$)
- ▶ Perceptron is given one training instance $\mathbf{x}^{(i)} \in \mathcal{D}$ at a time to make a prediction y
- ▶ In case of error ($y^{(i)} \neq t^{(i)}$), “reinforce” connections that would contribute to a correct prediction

$$w_j \leftarrow w_j + (t^{(i)} - y^{(i)})x_j^{(i)} \quad (3)$$

- ▶ Iterate for a fixed number of epochs (or, other stopping criteria)

Example

Classifying Iris Setosa flowers using a Perceptron.



perceptron.ipynb

Perceptron and SGD

- ▶ The original paper by Rosenblatt did not consider any explicit loss function to optimize and only proposed a training heuristic

$$\mathbf{w} \leftarrow \mathbf{w} + (t^{(i)} - y^{(i)})\mathbf{x}^{(i)} \quad (4)$$

- ▶ Later, the algorithm was reverse-engineered and interpreted as an application of SGD
- ▶ Keep in mind that the original algorithm is not SGD though!
 - ▶ e.g., Perceptron can be fed training data without any randomization, it is not **stochastic**!


Perceptron and SGD (2)

- ▶ Consider a 0-1 loss function for the instance $(\mathbf{x}^{(i)}, t^{(i)})$

$$\mathcal{L}_{0/1}^{(i)} = \frac{1}{2}(t^{(i)} - \text{sgn}\{\mathbf{w}^T \mathbf{x}^{(i)}\})^2 = (1 - t^{(i)} \text{sgn}\{\mathbf{w}^T \mathbf{x}^{(i)}\}) \quad (5)$$

- ▶ The sign function is a problem for differentiability and, thus, SGD application
- ▶ We consider a **smoothed surrogate loss function**:

$$\mathcal{L}^{(i)} = \max\{0, -t^{(i)}(\mathbf{w}^T \mathbf{x}^{(i)})\} \quad (6)$$

 perceptron-sgd.ipynb

Perceptron and SGD (3)

$$\mathcal{L} = \sum_{i=1}^N \max \{0, -t^{(i)}(\mathbf{w}^T \mathbf{x}^{(i)})\} \quad (7)$$

- ▶ We can compute the gradient and the SGD update of the smoothed surrogate loss function:


$$\nabla_{\mathbf{w}} \mathcal{L}^{(i)} = (y^{(i)} - t^{(i)}) \mathbf{x}^{(i)} \quad (8)$$

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} \mathcal{L}^{(i)} = \mathbf{w} + \eta (t^{(i)} - y^{(i)}) \mathbf{x}^{(i)} \quad (9)$$


With $\eta = 1$, we get back the Perceptron update.

Beyond the Perceptron

- ▶ The Perceptron is a neural network with a single computational layer
- ▶ It can only learn **linear** decision boundaries


 perceptron-xor.ipynb

- ▶ Much more powerful models can be obtained **composing** neurons into an **artificial neural network**
 - ▶ **Feedforward** (no cycles) vs **Recurrent** neural networks



Feedforward Neural Network


- ▶ Neurons (or, “units”) organized into layers:
 - ▶ A passthrough **input layer**
 - ▶ 1+ **hidden layers**
 - ▶ 1 **output layer**
- ▶ **Feedforward**: information flows in one direction, from the input layer to the output layer; no **feedback** connections
- ▶ We assume **fully connected** layers (we will generalize later...)



Note: MLP


- ▶ Feedforward NNs have been originally referred to as **multilayer perceptrons (MLP)**
 - ▶ e.g., scikit-learn provides the `MLPClassifier` class
- ▶ While still used, MLP is considered to be a *misnomer*
- ▶ Besides the use of artificial neurons, modern feedforward NNs profoundly differ from the Perceptron model w.r.t. activation functions, loss function, training algorithm, regularization strategies, ...

Notation and Terminology




- ▶ \mathbf{x} is the input vector
- ▶ $\mathbf{h}^{(i)}$ is the output of the i -th hidden layer
- ▶ \mathbf{y} is the output vector

Notation and Terminology (2)



- ▶ the number of layers (excluding the input layer) L is the **depth** of the network
 - ▶ **deep** learning involves NNs with many layers
- ▶ the number of units in a layer is known as the **width**


Activation Functions



- ▶ As for the Perceptron, units in the input layer simply output their own input unmodified
- ▶ Units in the hidden and output layers instead apply an **activation function** to their input
- ▶ We assume activation functions to be identical for all the units in the same layer, but they can differ across layers

Hidden Units

Let's consider the case of a single hidden layer:



$$h_1 = \phi\left(\sum_{j=1}^d w_{1,j} * x_j + b_i\right)$$

For the whole layer:

$$\mathbf{h} = \phi(\mathbf{W}\mathbf{x} + \mathbf{b})$$

where $\mathbf{W} \in \mathbb{R}^{m \times d}$ is a matrix of **weights**, and $\mathbf{b} \in \mathbb{R}^m$ is the **bias** vector.

Hidden Units (2)

In general, we have multiple hidden layers:

- ▶ $h_i^{(\ell)}$ denotes the i -th unit of the ℓ -th layer
- ▶ $w_{i,j}^{(\ell)}$ denotes the weight of the connection from the j -th unit of the $(\ell - 1)$ -th layer to the i -th unit of the ℓ -th layer

So, for $\ell > 1$:


$$\mathbf{h}^{(1)} = \phi^{(1)}(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^1)$$

$$\mathbf{h}^{(\ell)} = \phi^{(\ell)}(\mathbf{W}^{(\ell)}\mathbf{h}^{(\ell-1)} + \mathbf{b}^\ell)$$

Activation Function: Examples

ReLU (Rectified Linear Unit)


$$\phi(z) = \max \{0, z\}$$



Parametric ReLU (PReLU)

$$\phi(z) = \begin{cases} z & z > 0 \\ pz & z \leq 0 \end{cases}$$


e.g., $p = 0.01$



Activation Function: Examples (2)


Logistic sigmoid

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Hyperbolic Tangent

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$



Output Units

- ▶ Different output units are used depending on the task
- ▶ A linear output layer can be used for **regression**:

$$\mathbf{y} = \mathbf{W}^{(L)} \mathbf{h}^{(L-1)} + \mathbf{b}^{(L)}$$

- ▶ A sigmoid unit can be used for **binary classification**:

$$\mathbf{y} = \sigma(\mathbf{W}^{(L)} \mathbf{h}^{(L-1)} + \mathbf{b}^{(L)})$$

- ▶ A softmax layer can be used for **multiclass classification**:

$$\mathbf{z} = \mathbf{W}^{(L)} \mathbf{h}^{(L-1)} + \mathbf{b}^{(L)}$$

$$y_i = \frac{e^{z_i}}{\sum_{k=1}^K e^{z_k}}$$

NN and Function Composition

Let $f^{(\ell)}$ be the function implemented by the ℓ -layer, e.g.:

$$\mathbf{h}^{(1)} = f^{(1)}(\mathbf{x}) = \phi^{(1)}(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^1)$$


$$\mathbf{h}^{(2)} = f^{(2)}(\mathbf{h}^{(1)})$$

...

$$\mathbf{y} = f^{(L)}(\mathbf{h}^{(L-1)})$$

The NN computes the composite function:

$$\mathbf{y} = f^{(L)} \circ f^{(L-1)} \circ \dots \circ f^{(1)}$$

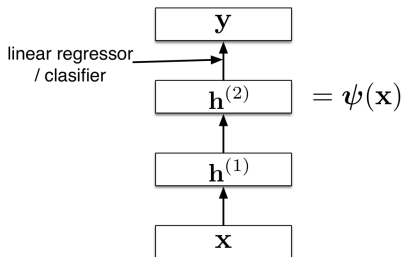


NN and Function Composition (2)

- ▶ The NN computes the composite function:

$$\mathbf{y} = f^{(L)} \circ f^{(L-1)} \circ \dots \circ f^{(1)}$$

- ▶ Output unit is often a linear function (regression) or a sigmoid (binary classification), similar to the Perceptron step function
- ▶ It is like applying a linear model to **features** $\psi(\mathbf{x})$ computed by the first $(L - 1)$ layers



From Linear to Nonlinear

Is it so important to use nonlinear activation functions?

Let's consider a model with a single hidden layer h , a **linear** activation function, and a linear output function:

$$h = W^{(1)}x + b^{(1)} \quad (10)$$

$$y = W^{(2)}h + b^{(2)} \quad (11)$$

We get:

$$y = W^{(2)}(W^{(1)}x + b^{(1)}) + b^{(2)} = \quad (12)$$


$$= W^{(2)}W^{(1)}x + W^{(2)}b^{(1)} + b^{(2)} = \quad (13)$$

$$= W'x + b' \quad (14)$$

With linear activations, you end up with a linear model, regardless of how many layers you use!

Example: XOR

Let's consider the following NN, with the Heaviside step function $H(z)$ for activation:



$$H(z) = \begin{cases} 1 & z \geq 0 \\ 0 & z < 0 \end{cases}$$



xor.ipynb

Expressivity

How “powerful” ANNs can be?

- ▶ We saw that linear models are easy to train, but they can only represent linear functions
- ▶ NNs overcome this limitation introducing nonlinearity
- ▶ **Question:** do we need a specialized model family for each nonlinear function to learn?

Universal Approximation Theorem

Theorem

A feedforward NN with a linear output layer, any “squashing” activation function, and enough hidden units can approximate any Borel measurable function from one finite-dimensional space to another with any desired nonzero amount of error.

- ▶ G. Cybenko, “Approximation by superpositions of a sigmoidal function” (1989)
- ▶ squashing: activation output is a finite interval (e.g., $[0, 1]$)
- ▶ Borel measurable function, TL;DR: any continuous function on a closed and bounded subset of \mathbb{R}^n
- ▶ Initially proven for the sigmoid activation, then proven for various functions, including ReLU

Universal Approximation Theorem (2)

But, don't be too enthusiastic!


- ▶ The theorem guarantees that a large enough NN can **represent** any function, but we are not given a training algorithm that can **learn** any function
- ▶ Training can fail for two reasons
 1. we completely fail to compute the parameters corresponding to the desired function
 2. as a result of **overfitting**, we learn the wrong function
- ▶ Moreover, the theorem does not say how **large** the hidden layer has to be!

Deep Neural Networks (DNN)

- ▶ In practice, better results can be achieved with **more layers** (i.e., **deeper** networks) rather than larger layers
- ▶ Traditional distinction between **shallow** and **deep** NNs
- ▶ No standard definition of “deep”
 - ▶ Popular one: “NNs with more than 1 hidden layer are deep”
 - ▶ Sometimes: “NNs with more than 2 hidden layers are deep”
- ▶ Just pick the definition you prefer...

Impact of Depth: Example

Multidigit number transcription from Street View address photographs

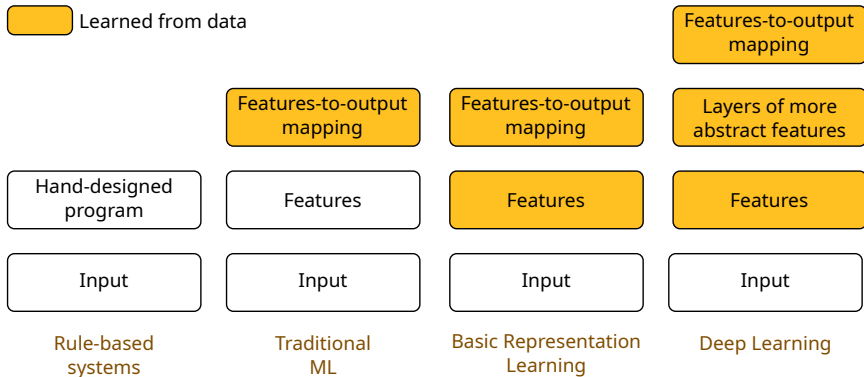


Goodfellow et al., "Multi-digit Number Recognition from Street View Imagery using Deep Convolutional Neural Networks", 2014

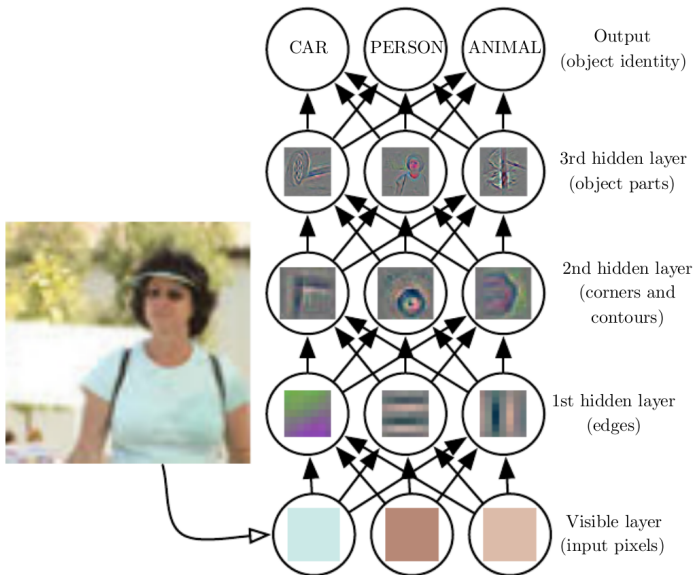
Deep Learning

- ▶ **Deep learning** (DL) is a collection of ML methods based on DNNs and **representation learning**, which can be adopted for (semi-)supervised, unsupervised and reinforcement learning tasks
- ▶ Not just a matter of having “more layers”
- ▶ Traditional ML: given a **representation** of the input data (i.e., features), learn a **mapping** from representation to output
- ▶ DL aims to learn both the representation and the mapping!
- ▶ Multiple layers → learning **multiple levels of composition**
 - ▶ discovering representations of the input expressed in terms of other, simpler representations

Comparison of AI Approaches



Example




Back to History

- ▶ ANNs had 3 waves of popularity
 - ▶ 1950s-1960s: Perceptron, **only linear models**
 - ▶ 1980s-1990s: MLPs, backpropagation, **difficulties to train large models**
 - ▶ 2006–today: deep learning
- ▶ Current renaissance began in 2006, when Hinton showed that a DNN could be efficiently trained to outperform a SVM-based solution on the MNIST benchmark.
 - ▶ *Hinton et al., “A Fast Learning Algorithm for Deep Belief Nets”*
- ▶ In general, DL “revolution” enabled by multiple factors:
 - ▶ more data (“Big Data”)
 - ▶ more and better hardware for parallel computing (e.g., GPGPUs)
 - ▶ new/improved software libraries

Increasing Dataset Sizes

- ▶ First DL applications date back to 1990s, but more as an art rather than an accessible technology (due to very difficult training)
- ▶ Nearly identical learning algorithms today reach human performance on complex tasks (though the models have undergone changes that simplify the training of very deep architectures)
- ▶ What else has changed? **We have much more data for training!**
 - ▶ ...and the computational resources to store and process those data


Example: Popular Datasets



Increasing Model Sizes

- ▶ Besides handling more data, we have the resources to run much **larger models** today
 - ▶ mainly faster CPUs, GPGPUs, better software infrastructures for parallel and distributed computing
 - ▶ ...and enough data to train them
- ▶ Unless new technologies enable significantly faster scaling, ANNs will reach the same number of neurons of the human brain not earlier than 2050

Example: Neural Network Size



- 1: Perceptron (1962) 4: Backpropagation NN (1986)
10: Deep Belief Network by Hinton (2006)
20: GoogLeNet (2014)

Increasing Impact

- ▶ The spectacular success of DL in several domains has further dramatically increased the interest of researchers and industries in the last decades
 - ▶ e.g, object recognition, speech recognition, image/video generation, ...
- ▶ Since 2015, DL has been also successfully and widely applied to reinforcement learning tasks

Recommended Readings

- ▶ Goodfellow et al., Chapter 6 (§6.1–§6.4)
- ▶ Dive into DL (d2l.ai): §5.1