

*Machine Learning*

# Recurrent Neural Networks

Gabriele Russo Russo    Francesco Lo Presti

Laurea Magistrale in Ingegneria Informatica – A.Y. 2023/24



**TOR VERGATA**  
UNIVERSITÀ DEGLI STUDI DI ROMA


Dipartimento di Ingegneria Civile e  
Ingegneria Informatica

# Learning with Sequences

- ▶ So far, we have focused primarily on fixed-length data
  - ▶ e.g., input vector  $x$  consisting of a fixed number of features
  - ▶ e.g., raw pixel values at each coordinate in an image
- ▶ Many tasks require dealing with **sequential data**
- ▶ Not only input sequences:
  - ▶ Input: time series prediction, video analysis, musical information retrieval
  - ▶ Output: image captioning, speech synthesis, music generation
  - ▶ Both: text translation

# Recurrent Neural Networks

- ▶ **Recurrent neural networks (RNNs)**: DL models that capture the dynamics of sequences via **recurrent connections**
  - ▶ cycles in the network of nodes
- ▶ RNNs are *unrolled* across time steps, with the same underlying parameters applied at each step



## Remark

- ▶ RNNs work well with sequential data, but they are not the only kind of DNN used for these tasks
- ▶ e.g., CNNs have been successfully applied to sequences (e.g., [WaveNet](#) to generate synthetic audio from raw audio)
- ▶ Recently, [Transformers](#) have been shown to outperform RNNs in many cases

# Hidden State

- ▶ Given an input sequence  $x_1, x_2, \dots, x_{t-1}$ , we are interested in

$$P(x_t | x_{t-1}, x_{t-2}, \dots, x_1) = ??$$

- ▶ Storing the full (or most recent) sequence of past observations would be unfeasible and would require an exponentially large number of parameters
- ▶ Therefore, we exploit a **hidden state**  $h_{t-1}$  to capture sequence information up to  $t - 1$

$$P(x_t | x_{t-1}, x_{t-2}, \dots, x_1) \approx P(x_t | h_{t-1})$$

$$h_t = f(x_t, h_{t-1})$$


# Recurrent Layers

- ▶ Recurrent (hidden) layers are characterized by hidden states
  - ▶ “hidden” refers to different concepts in this sentence
- ▶ Consider an input example at time  $t$  from the sequence:  $x_t$
- ▶ The output of the hidden layer is computed as

$$h_t = \phi(W_x x_t + W_h h_{t-1} + b)$$

- ▶ ...and this is the **hidden state** of the RNN
- ▶ **Recurrent**:  $h_t$  is defined in terms of  $h_{t-1}$
- ▶  $\phi()$  is usually tanh or sigmoid for RNNs

## Recurrent Layers (2)





Note: alternative RNN models exist (e.g., in the model we consider hidden state and output of the layer are the same, but they could be computed differently)

# Input and Output of RNNs

- ▶ **Sequence-to-sequence**: a sequence of inputs produces a sequence of outputs
  - ▶ e.g., time series forecasting: you feed data over the last N days, and output the series shifted by one day into the future
- ▶ **Sequence-to-vector**: feed a sequence of inputs and ignore all outputs except for the last one
  - ▶ e.g., given a social media post, output a sentiment score
- ▶ **Vector-to-sequence**: feed the RNN the same input vector over and over again at each time step and let it output a sequence
  - ▶ e.g., given an image, produce a caption



# Input and Output of RNNs (2)



# Training RNNs

- ▶ Recall forward + backward propagation for feedforward NNs
- ▶ Forward propagation relatively straightforward
- ▶ Applying backpropagation in RNNs is called **backpropagation through time** (Werbos, 1990)
- ▶ We expand (or unroll) the computational graph of an RNN one time step at a time
  - ▶ Unrolled RNN is essentially a feedforward NN, with the same parameters repeated throughout the unrolled network
  - ▶ We can apply the chain rule as usual
  - ▶ The gradient w.r.t. each parameter must be summed across all the parameter occurrences

# Backpropagation Through Time (BPTT)

Consider a simplified RNN model (identity act. function, no bias)

- ▶  $\mathbf{h}_t = \mathbf{W}_x \mathbf{x}_t + \mathbf{W}_h \mathbf{h}_{t-1}$
- ▶  $\mathbf{y}_t = \mathbf{W}_y \mathbf{h}_t$
- ▶ Loss function:  $\mathcal{L} = \frac{1}{T} \sum_{t=1}^T \ell(\mathbf{t}_t, \mathbf{y}_t)$

At any time step  $t$ , it is straightforward to compute:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{y}_t} = \frac{1}{T} \frac{\partial \ell(\mathbf{t}_t, \mathbf{y}_t)}{\partial \mathbf{y}_t} \quad (1)$$

# Backpropagation Through Time (2)

- ▶ The weights  $\mathbf{W}_y$  are used at **every time step** to compute the output given  $\mathbf{h}_t$
- ▶ Recall that  $\mathcal{L}$  depends on  $\mathbf{y}_1, \mathbf{y}_2, \dots$
- ▶ To compute the gradient w.r.t.  $\mathbf{W}_y$ :

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}_y} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial \mathbf{y}_t} \frac{\partial \mathbf{y}_t}{\partial \mathbf{W}_y} = \frac{1}{T} \sum_{t=1}^T \frac{\partial \ell(\mathbf{t}_t, \mathbf{y}_t)}{\partial \mathbf{y}_t} \mathbf{h}_t^\top \quad (2)$$

## Backpropagation Through Time (3)

- ▶ Next, we need the gradient of  $\mathcal{L}$  w.r.t.  $\mathbf{h}_t$
- ▶ For the final time step  $T$ , it is easy (the loss only depends on  $\mathbf{h}_T$  via  $\mathbf{y}_T$ )

$$\frac{\partial \mathcal{L}}{\partial \mathbf{h}_T} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}_T} \frac{\partial \mathbf{y}_T}{\partial \mathbf{h}_T} = \mathbf{W}_y^\top \frac{\partial \mathcal{L}}{\partial \mathbf{y}_T} \quad (3)$$

# Backpropagation Through Time (4)

For any time step  $t < T$ , things become trickier...

$$\frac{\partial \mathcal{L}}{\partial \mathbf{h}_t} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}_t} \frac{\partial \mathbf{y}_t}{\partial \mathbf{h}_t} + \frac{\partial \mathcal{L}}{\partial \mathbf{h}_{t+1}} \frac{\partial \mathbf{h}_{t+1}}{\partial \mathbf{h}_t} = \mathbf{W}_y^\top \frac{\partial \mathcal{L}}{\partial \mathbf{y}_t} + \mathbf{W}_h^\top \frac{\partial \mathcal{L}}{\partial \mathbf{h}_{t+1}} = \quad (4)$$

$$= \mathbf{W}_y^\top \frac{\partial \mathcal{L}}{\partial \mathbf{y}_t} + \mathbf{W}_h^\top \left( \mathbf{W}_y^\top \frac{\partial \mathcal{L}}{\partial \mathbf{y}_{t+1}} + \mathbf{W}_h^\top \frac{\partial \mathcal{L}}{\partial \mathbf{h}_{t+2}} \right) = \quad (5)$$

$$= \sum_{i=0}^{T-t} \left( \mathbf{W}_h^\top \right)^i \mathbf{W}_y^\top \frac{\partial \mathcal{L}}{\partial \mathbf{y}_{t+i}} \quad (6)$$

You can see that for long sequences we need to compute large powers of  $\mathbf{W}_h$  (eigenvalues  $< 1$  may vanish, and those  $> 1$  may diverge!)

# Backpropagation Through Time (5)


Finally, we can compute:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}_h} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{W}_h} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial \mathbf{h}_t} \mathbf{h}_{t-1}^\top \quad (7)$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}_x} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{W}_x} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial \mathbf{h}_t} \mathbf{x}_t^\top \quad (8)$$

# Example: Time Series Forecasting


- ▶ You have just been hired as a data scientist by Chicago's Transit Authority
- ▶ First task: build a model capable of forecasting the number of passengers that will ride on bus and rail the next day
- ▶ You have access to daily ridership data since 2001





# Example: Preliminary Observations

- ▶ We are facing a **multivariate time series**
- ▶ Looking at data, we note that a similar pattern is clearly repeated every week (**weekly seasonality**)
- ▶ Exploiting this observation, we could forecast tomorrow's ridership by just copying the values from a week earlier (**naive forecasting**): less than 10% mean error!



## Example: Using a RNN

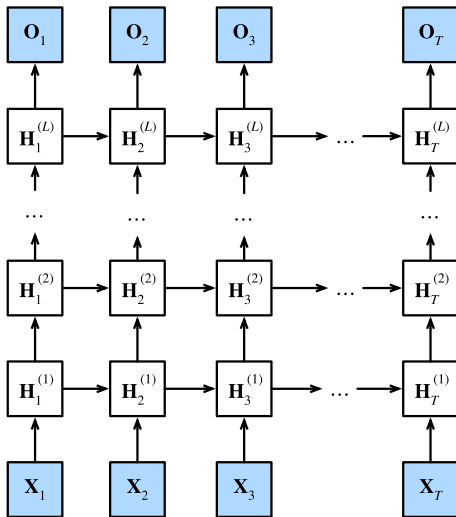
- ▶ We will use a RNN to (hopefully) outperform naive forecasting and a **SARIMA** model
  - ▶ SARIMA belongs to a family of widely used models for time series forecasting

 rnn\_timeseries.ipynb (part 1)

# Deep RNNs

- ▶ So far, we defined RNNs consisting of a sequence input, a single hidden RNN layer, and an output layer
- ▶ This NN is already **deep** in some sense: inputs from the first time step can influence outputs 100-1000s steps later
- ▶ We may also wish to retain the ability to express complex input-output relationships at a given time step
- ▶ We can stack RNN layers on top of each other

# Deep RNNs (2)



 rnn\_timeseries.ipynb (part 2)

# Handling Long Sequences

$$\frac{\partial \mathcal{L}}{\partial \mathbf{h}_t} = \sum_{i=0}^{T-t} \left( \mathbf{W}_h^\top \right)^i \mathbf{W}_y^\top \frac{\partial \mathcal{L}}{\partial \mathbf{y}_{t+i}}$$

- ▶ Long sequences cause large powers of  $\mathbf{W}_h$  to be computed (eigenvalues  $< 1$  may vanish, and those  $> 1$  may diverge!)
- ▶ The problem of **vanishing and exploding gradients** is one of the key challenges for RNN training
- ▶ Good parameter initialization, faster optimizers, dropout can help, but do not solve the problem
- ▶ Non-saturating activ. functions (e.g., ReLU) worsen the situation (tanh is a popular choice for RNNs)
  - ▶ If activation value is increased by weights update at the first step, it is further increased at every step!


# LSTM Cells

- ▶ Traversing RNNs, some information is lost at each time step
- ▶ After a while, state may contain no trace of the first inputs!
- ▶ Long short-term memory (LSTM): one of the first and most successful techniques to deal with vanishing gradients
- ▶ Published in 1997; become dominant model for sequence learning from 2011 until the rise of Transformer models in 2017
- ▶ Simple RNNs have long-term memory in the form of weights, which change slowly during training; RNNs also have short-term memory in the form of ephemeral activations, which pass from each node to successive nodes
- ▶ The LSTM model introduces an intermediate type of storage via memory cells, which replace recurrent nodes

# Memory Cell

- ▶ Each memory cell is equipped with an **internal state** and a number of **multiplicative gates**
- ▶ Gates determine whether
  1. input should impact the internal state (**input gate**)
  2. internal state should be reset (**forget gate**)
  3. internal state should impact cell's output (**output gate**)


# LSTM Gates



- ▶ 3 fully connected layers with sigmoid activation compute the values of the input, forget, and output gates
- ▶ Values of the gates range in  $(0, 1)$ 
  - ▶ e.g., forget gate determines *how much* of the current state should be kept



## LSTM Gates (2)



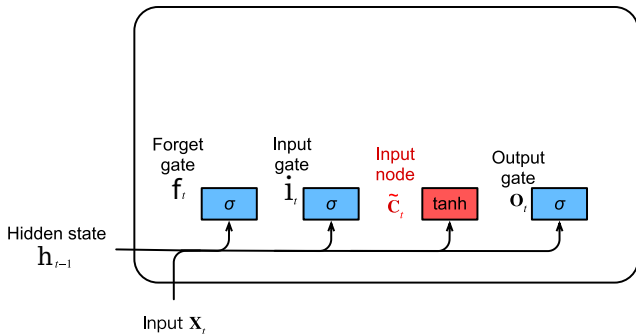
- ▶ Consider  $x \in \mathbb{R}^d$  and  $h$  hidden units
- ▶ Hidden state is  $h_{t-1} \in \mathbb{R}^h$

$$i_t = \sigma(x_t W_{xi} + h_{t-1} W_{hi} + b_i) \quad (9)$$

$$f_t = \sigma(x_t W_{xf} + h_{t-1} W_{hf} + b_f) \quad (10)$$

$$o_t = \sigma(x_t W_{xo} + h_{t-1} W_{ho} + b_o) \quad (11)$$

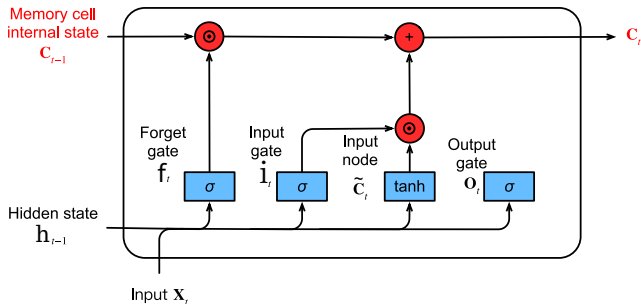
# Input Node



- $\tilde{c}_t \in (-1, 1)^h$  denotes the **cell input node**

$$\tilde{c}_t = \tanh(\mathbf{x}_t \mathbf{W}_{xc} + \mathbf{h}_{t-1} \mathbf{W}_{hc} + \mathbf{b}_c)$$

# Cell Internal State

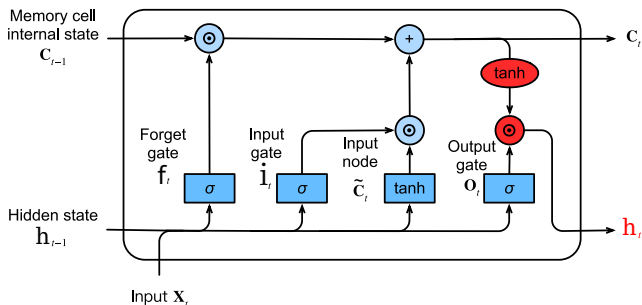


►  $c_t \in \mathbb{R}^h$  is the cell **internal state**

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

where  $\odot$  denotes the element-wise product

# Hidden State



- The **hidden state**  $h_t \in (-1, 1)^h$  is the output of the cell, as seen by next layer

$$h_t = o_t \odot \tanh(c_t)$$


## Example: Time Series Forecasting

- ▶ Last part of the notebook uses LSTM for the ridership forecasting
- ▶ Note that LSTM do not provide significant benefits in this simple example, but they generally work much better than simple RNNs in many tasks

 rnn\_timeseries.ipynb (part 3)

# WaveNet

- ▶ CNN proposed in 2016; excellent performance on 1D audio sequences (e.g., to generate synthetic voices)
- ▶ Idea: stacking 10 convolutional layers to obtain a large enough receptive field, able to capture long-term patterns at higher layers
- ▶ More details: Residual blocks and gated activations



# References

- ▶ D2L: 9.\*, 10.1, 10.3
- ▶ Hands-on ML: Chapter 15