

Machine Learning

Introduction to Reinforcement Learning

Gabriele Russo Russo Francesco Lo Presti

Laurea Magistrale in Ingegneria Informatica – A.Y. 2023/24




TOR VERGATA
UNIVERSITÀ DEGLI STUDI DI ROMA

Dipartimento di Ingegneria Civile e
Ingegneria Informatica

Reinforcement Learning

- ▶ Supervised learning
- ▶ Unsupervised learning
- ▶ **Reinforcement learning**
 - ▶ Branch of ML dealing with sequential decision-making

Reinforcement Learning



- ▶ Agent interacts with environment through **actions**
- ▶ Feedback in the form of **reward** (or **paid cost**)
- ▶ Goal: maximizing cumulated reward over the long run
- ▶ Trial-and-error experience (no complete knowledge of environment a priori)

Example: Tic-Tac-Toe

- ▶ **State**: representation of the board (3x3 matrix)
- ▶ **Actions**: available cells to mark
- ▶ **Reward**: 1 for a winning move, 0 otherwise

X	O	O
O	X	X
		X

Example: AlphaZero by DeepMind

- ▶ Software able to play Go, Chess and Shogi ¹
 - ▶ Board games with huge number of legal positions (i.e., state space)
- ▶ Trained via self-play and advanced deep RL techniques
- ▶ Superhuman level of play with 24-hour training
- ▶ First presented in 2017; in 2019 [MuZero](#), generalization to play Atari games and other board games without prior rule knowledge

¹<https://arxiv.org/abs/1712.01815>


Example: AlphaDev by DeepMind

- ▶ Announced in 2023²
- ▶ RL used to develop new C++ sorting algorithm, now accepted in the standard library
- ▶ 70% faster on short sequences (2-3 items), 1.7% faster on long sequences
- ▶ State: instructions generated so far and state of the CPU
- ▶ Actions: assembly instructions to add
- ▶ Reward: based on sorting correctness and efficiency


²<https://www.deepmind.com/blog/alphadev-discovers-faster-sorting-algorithms>

Other Examples

- ▶ Autonomous vehicles
- ▶ Robot control
- ▶ Trading
- ▶ Autonomous network and computer systems
- ▶ Videogames
- ▶ ...



Reinforcement Learning

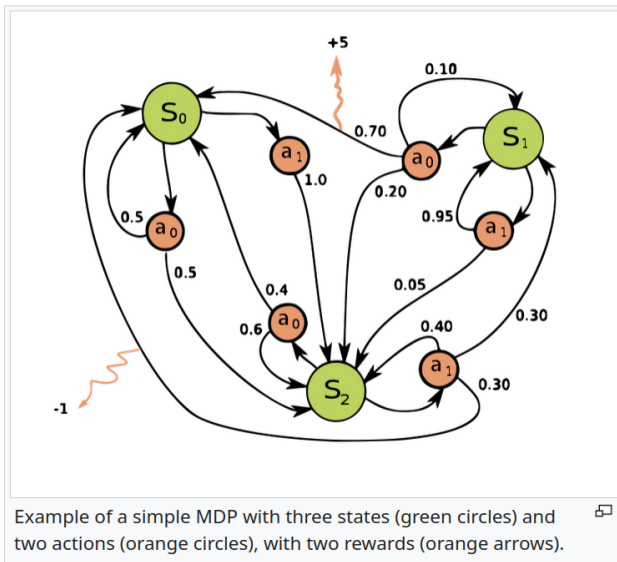


- ▶ Agent, environment, actions, state, rewards, ...
- ▶ Modeled depending on the specific task
 - ▶ e.g., autonomous car uses different state information compared to chess player
- ▶ Formally defined as a **Markov Decision Process (MDP)**
- ▶ A framework to model decision making in situations where outcomes are partly random

Markov Decision Process (MDP)

- ▶ Extension of discrete-time Markov chains
- ▶ At each time step t , the process is in some state s_t
- ▶ The decision maker (the agent) chooses an action a_t among those available in state s_t
 - ▶ e.g., robot observes current position and decides direction to move; some directions might be blocked by obstacles
- ▶ Following a_t , the process moves to (random) state s_{t+1}
 - ▶ e.g., autonomous drone chooses an action to reduce altitude; actual outcome may depend on (unpredictable) wind speed
- ▶ Agent receives a reward (or, equivalently, pays a cost)
 - ▶ e.g., robot may get a reward for reaching its final destination
 - ▶ e.g., chess player rewarded at the end of a match

Example



Example of a simple MDP with three states (green circles) and two actions (orange circles), with two rewards (orange arrows).

Markov Decision Process (2)

What defines an MDP?

- ▶ \mathcal{S} : a (finite) set of states
- ▶ \mathcal{A} : a (finite) set of actions
- ▶ p : state transition probabilities

$$p(s'|s, a) = P[S_{t+1} = s' | S_t = s, A_t = a]$$

- ▶ r : reward function (or, c : cost function)
 1. $r(s, a) = E[R_t | S_t = s, A_t = a]$
 2. $r(s, a, s') = E[R_t | S_t = s, A_t = a, S_{t+1} = s'] \longrightarrow$
 $r(s, a) = \sum_{s'} p(s'|s, a) r(s, a, s')$

Markov Property

“The future is independent of the past given the present”

Definition

A state S_t is **Markov** if and only if

$$P[S_{t+1} | S_1, \dots, S_t] = P[S_{t+1} | S_t]$$

- ▶ The state captures all relevant information from the history
- ▶ i.e., the state is a sufficient statistic of the future

Objective: Episodic Tasks

- ▶ Informally, we said that the agent aims to maximize the collected reward over time
- ▶ Let's consider an **episodic** task, where the agent-environment interaction naturally terminates at some final time step T
 - ▶ e.g., the end of a chess match
 - ▶ e.g., the time a robot reaches its destination or runs out of battery
- ▶ At time t , we aim to maximize the **expected return** G_t

$$G_t = R_t + R_{t+1} + \dots + R_T$$

Objective: Continuing Tasks

- ▶ In many cases the agent–environment interaction does not break naturally into identifiable episodes, but goes on continually without limit
 - ▶ e.g., an agent managing VM migration in a Cloud datacenter
 - ▶ e.g., the control system of RL-based traffic lights
- ▶ In this scenario, the goal of the agent is maximizing the **expected cumulative discounted** reward

$$G_t = R_t + \gamma R_{t+1} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k}$$

where $\gamma \in [0, 1)$ is the discount factor

Reward vs Cost

- ▶ You can either maximize the expected reward or minimize the expected cost

$$G_t = C_t + \gamma C_{t+1} + \dots = \sum_{k=0}^{\infty} \gamma^k C_{t+k}$$

- ▶ The two formulations are equivalent; you can easily switch between them by setting

$$r(s, a) = -c(s, a)$$

- ▶ In the following, we will mostly refer to costs; keep in mind this equivalence

Policy

Definition

A **policy** π is a distribution over actions given a state s

$$\pi(a|s) = p(A_t = a | S_t = s)$$

- ▶ A policy fully defines agent's behavior
- ▶ MDP policies depend on the current state only
- ▶ Special case: **deterministic policy**

$$\pi : \mathcal{S} \rightarrow \mathcal{A}$$

Example: Deterministic Policy

State Action	
s_1	a_1
s_2	a_1
s_3	a_2
s_4	a_1

Value Function

Value function is a prediction of future costs

- ▶ can be used to evaluate how good/bad states and/or actions are
- ▶ and therefore to select actions e.g.

State	a_1	a_2	a_3
s_1	10	5	3
s_2	8	6	4
s_3	6	5	6
s_4	5	4	6
s_5	4	3	7
s_6	1	5	9
s_7	0	9	15

s	$\pi(s)$
s_1	a_3
s_2	a_3
s_3	a_2
s_4	a_2
s_5	a_2
s_6	a_1
s_7	a_1

Value Functions

Action value function (or, Q function)

Expected cost starting from state s , taking action a and then following policy π

$$Q_{\pi}(s, a) = E_{\pi}[G_t | S_t = s, A_t = a]$$

State value function

Expected cost starting from state s and then following policy π

$$V_{\pi}(s) = E_{\pi}[G_t | S_t = s]$$

Action Value Functions

The **action value function** can be decomposed into two parts:

- ▶ immediate cost
- ▶ discounted costs from successor state S_{t+1}

$$\begin{aligned}Q_{\pi}(s, a) &= E_{\pi}[G_t | S_t = s, A_t = a] \\&= E_{\pi}[C_t + \gamma C_{t+1} + \gamma^2 C_{t+2} \dots | S_t = s, A_t = a] \\&= E_{\pi}[C_t + \gamma (C_{t+1} + \gamma C_{t+2} \dots) | S_t = s, A_t = a] \\&= E_{\pi}[C_t + \gamma G_{t+1} | S_t = s, A_t = a] \\&= c(s, a) + \gamma E_{\pi}[G_{t+1} | S_t = s, A_t = a]\end{aligned}$$

Bellman equation:

$$Q_{\pi}(s, a) = c(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) Q_{\pi}(s', \pi(s'))$$

State Value Functions

The **value function** can be similarly decomposed into two parts:

- ▶ immediate cost C_t
- ▶ discounted cost from successor state $V(S_{t+1})$

$$\begin{aligned}V_{\pi}(s) &= E_{\pi}[G_t | S_t = s] \\&= E_{\pi}[C_t + \gamma C_{t+1} + \gamma^2 C_{t+2} \dots | S_t = s] \\&= E_{\pi}[C_t + \gamma (C_{t+1} + \gamma C_{t+2} \dots) | S_t = s] \\&= E_{\pi}[C_t + \gamma G_{t+1} | S_t = s]\end{aligned}$$

Bellman equation:

$$V_{\pi}(s) = c(s, \pi(s)) + \gamma \sum_{s'} p(s' | s, \pi(s)) V_{\pi}(s')$$

Optimal Value Function

Optimal action value function

$Q^*(s; a)$ is the maximum action-value function over all policies

$$Q^*(s, a) = \max_{\pi} Q_{\pi}(s, a)$$

Optimal state value function

$V^*(s)$ is the minimum value function over all policies

$$V^*(s) = \max_{\pi} V_{\pi}(s)$$

Bellman Optimality Equations

$$Q_{\pi}(s, a) = c(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) Q_{\pi}(s', \pi(s'))$$



$$Q^*(s, a) = c(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) \min_{a'} Q^*(s', a')$$

$$V^*(s) = \min_a Q^*(s, a)$$


$$Q^*(s, a) = c(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V^*(s')$$

Optimal Policy

Given $Q^*(s, a)$ the optimal action when the system is in state s is:

$$\pi^*(s) = a^*(s) = \arg \min_{a \in \mathcal{A}} Q^*(s, a)$$

State	a_1	a_2	a_3
s_1	10	5	3
s_2	8	6	4
s_3	6	5	6
s_4	5	4	6
s_5	4	3	7
s_6	1	5	9
s_7	0	9	15



Optimal Action
a_3
a_3
a_2
a_2
a_2
a_1
a_1

How to compute V^* ?

- ▶ If we know the optimal value function, we have an optimal policy!
- ▶ **But...** how do we compute the optimal value function??

Value Iteration

Bellman Equation

$$Q^*(s, a) = c(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) \min_{a'} Q^*(s', a')$$

- ▶ Suppose we know the solution to subproblems $Q^*(s', a')$
- ▶ $Q^*(s, a)$ can be computed by one-step lookahead

$$Q^*(s, a) \leftarrow c(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) \min_{a'} Q^*(s', a')$$

- ▶ The idea is to apply these updates iteratively
- ▶ Proven to converge (see, Contraction Mapping Theorem in Sutton's book)

Value Iteration: Algorithm

Value Iteration

```
1  $i \leftarrow 0$ 
2  $Q_i(s, a) \leftarrow 0, \forall s \in \mathcal{S}, \forall a \in \mathcal{A}(s)$ 
3 repeat
4   forall  $s \in \mathcal{S}$  do
5     forall  $a \in \mathcal{A}(s)$  do
6        $Q_{i+1}(s, a) \leftarrow$ 
7          $c(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \min_{a' \in \mathcal{A}(s')} Q_i(s', a')$ 
8     end
9   end
10   $i \leftarrow i + 1$ 
11 until  $\max_{s,a} |Q_i(s, a) - Q_{i-1}(s, a)| < \epsilon$ 
12  $\pi^*(s) = \arg \min_a Q_i(s, a), \forall s \in \mathcal{S}$ 
```


Value Iteration: Alternative Algorithm

Value Iteration - Alternative

```
1  $i \leftarrow 0$ 
2  $Q_i(s, a) \leftarrow 0, \forall s \in \mathcal{S}, \forall a \in \mathcal{A}(s)$ 
3  $V_i(s) \leftarrow 0, \forall s \in \mathcal{S}$ 
4 repeat
5   forall  $s \in \mathcal{S}$  do
6     forall  $a \in \mathcal{A}(s)$  do
7        $Q_{i+1}(s, a) \leftarrow c(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_i(s)$ 
8     end
9      $V_{i+1}(s) = \min_{a' \in \mathcal{A}(s)} Q_{i+1}(s, a')$ 
10  end
11   $i \leftarrow i + 1$ 
12 until  $\max_{s,a} |Q_i(s, a) - Q_{i-1}(s, a)| < \epsilon$ 
13  $\pi^*(s) = \arg \min_a Q_i(s, a), \forall s \in \mathcal{S}$ 
```


Example: Maze

- ▶ Consider a $S \times S$ grid
- ▶ Episodes start with agent randomly located in a cell in the first column
- ▶ Goal: reaching target cell $(1, S)$
- ▶ Some cells are blocked
- ▶ Some cells are slippery: when entering, the agent has a probability p_{slip} of slipping one cell ahead along her current direction




Example: Maze (2)


- ▶ State: $s = (x, y)$
- ▶ Actions: $a \in \{(0, 1), (0, -1), (1, 0), (-1, 0)\}$
- ▶ Reward:
 - ▶ 0 for entering the goal cell
 - ▶ $-M$ for exiting the grid or crashing into a blocked cell ($M \gg 1$)
 - ▶ -1 otherwise

 maze.py (--agent mdp)


Maze: Optimal Policy



Maze: Optimal Value Function




Example: Cloud Auto-scaling



- ▶ We periodically make a decision about scaling in/out an app component (a thread, a VM, a container, ...)
- ▶ We are concerned with 3 objectives:
 - ▶ Monetary resource cost (or, resource usage in general)
 - ▶ Performance req. satisfaction (e.g., max response time)
 - ▶ Scaling overhead

Auto-scaling: MDP formulation



- ▶ State at time slot i : $s_i = (k_i, \lambda_i)$
 - ▶ k_i component parallelism
 - ▶ λ_i avg. arrival rate (of requests, jobs, data, ...)
- ▶ Action at time slot i : $a_i \in \{0, +1, -1\}$

MDP Model: Transition Probabilities

- ▶ State of the system $s = (k, \lambda)$
 - ▶ $1 \leq k \leq K^{max}$ Component parallelism
 - ▶ λ avg. input rate
 - ▶ λ is discretized, i.e., $\lambda_i \in \{0, \Delta\lambda, 2\Delta\lambda, (L-1)\Delta\lambda\}$
 - ▶ $\Delta\lambda$ quantization step size, L number of discrete values
- ▶ Available actions $\mathcal{A} = \{-1, 0, +1\}$
- ▶ Transition probabilities $p(s'|s, a) = p((k', \lambda')|(k, \lambda), a)$

$$\begin{aligned} p(s'|s, a) &= P[s_{t+1} = (k', \lambda') | s_t = (k, \lambda), a_t = a] = \\ &= \begin{cases} P[\lambda_{t+1} = \lambda' | \lambda_t = \lambda] & k' = k + a \\ 0 & \text{otherwise} \end{cases} = \\ &= \mathbb{1}_{\{k'=k+a\}} P[\lambda_{t+1} = \lambda' | \lambda_t = \lambda] \end{aligned}$$

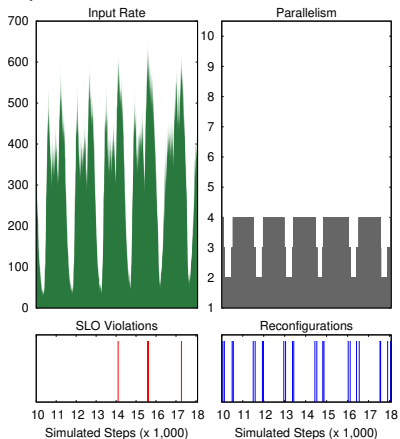
MDP Model: Cost Function

$$c(s, a, s') = \underbrace{w_{res} \frac{k + a}{K_{max}}}_{\text{Resource Cost}} + \underbrace{w_{perf} \mathbb{1}_{\{R(s, a, s') > R^{max}\}}}_{\text{Performance}} + \underbrace{w_{rcf} \mathbb{1}_{\{a \neq 0\}}}_{\text{Reconfig.}}$$


- ▶ $w_{res} + w_{perf} + w_{rcf} = 1, w_x \geq 0, x \in \{res, perf, rcf\}$
- ▶ $R(s, a, s')$: performance index, e.g, response time
- ▶ R^{max} : reference performance value
- ▶ We want to minimize $\sum_{t=0}^{\infty} \gamma^t c(s_t, a_t, s_{t+1}), \quad \gamma \in [0, 1)$

Trading-off Objectives

$$w_{perf}=0.6, w_{res} = w_{rcf}=0.2$$




$$w_{res}=0.6, w_{perf} = w_{rcf}=0.2$$



MDP Resolution

- ▶ We can use the Value Iteration algorithm to solve the MDP
 - ▶ i.e., finding the optimal policy
- ▶ Is this enough?
- ▶ Unfortunately, solving the MDP requires exact and complete knowledge of the underlying model
 - ▶ state transition probabilities
 - ▶ cost function
- ▶ In practice, we don't have such information!

Reinforcement Learning



- RL aims to learn the optimal policy through interaction and evaluative feedback

Model-free vs Model-based RL

- ▶ **Model-free** RL: no model of the environment is available or used; the optimal policy is learned through experience only
- ▶ **Model-based** RL: a (possibly partial) model of the environment is available and used to derive the optimal policy
 - ▶ a partial model can boost learning speed
 - ▶ RL may also be used in presence of a complete model instead of VI; e.g., with a large number of rarely visited states VI would unnecessarily run for a long time!
 - ▶ You may also try to learn the model online and use it to compute a policy

Value-based vs Policy-based RL

- ▶ **Value-based** RL: aims to learn the optimal value function through experience; the policy is derived from it
 - ▶ Simplest RL algorithms belong to this group
 - ▶ We will mainly focus on this group in the following
- ▶ **Policy-based** RL: aims to directly learn the optimal policy through experience; no explicit computation/learning of the value function
- ▶ **Hybrid** approaches: e.g., the Actor-Critic framework

Simple Value-based RL Algorithm

A simple RL algorithm

```
1  $t \leftarrow 0$   
2 Initialize  $Q$   
3 Loop  
4   |  $t \leftarrow t + 1$   
5 EndLoop
```

Q-learning

- ▶ Proposed by Chris Watkins in 1989
- ▶ One of the most known (and simplest) RL algorithms
- ▶ Proven to converge to the optimal policy under mild assumptions
 - ▶ ...after n steps, with $n \rightarrow \infty$

Q-learning: Action Selection

- ▶ How to choose an action at every time step?
- ▶ **Exploration vs Exploitation** dilemma
- ▶ **Exploitation**: using available knowledge to maximize reward
 - ▶ choose the “best” action, i.e., $a_t = \arg \max_a Q(s_t, a)$
- ▶ **Exploration**: discovering more information about the environment
 - ▶ choose other actions to learn more about the environment

Q-learning converges only if all state-action pairs are visited an infinite number of times as $t \rightarrow \infty$

- ▶ you can't **exploit** all the time
- ▶ you can't **explore** all the time

ϵ -Greedy Exploration

- ▶ Popular approach for the exploration-exploitation dilemma
- ▶ With probability $1 - \epsilon$ choose the greedy action
 $a^* = \arg \max_{a \in \mathcal{A}} Q(s, a)$
- ▶ With probability ϵ choose an action at random
- ▶ Improvement: ϵ -greedy with **decaying ϵ** (similar to decaying learning rate in SGD)

Softmax Action Selection

- ▶ Alternative to the ϵ -greedy strategy
- ▶ All actions assigned non-zero probability of being chosen
- ▶ Action $a \in \mathcal{A}$ is selected with probability

$$\pi(a|s) = \frac{\exp(Q(s, a) / \tau)}{\sum_{a' \in \mathcal{A}} \exp(Q(s, a') / \tau)}$$

- ▶ τ is the “temperature”
 - ▶ Small τ leads to greedy behavior
 - ▶ Large τ leads to random action selection
 - ▶ You usually start with a large temperature value and let it decay

Q-learning: Updating Q

With known model, we can compute Q iteratively using:

$$Q(s, a) \leftarrow c(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) \max_{a'} Q(s', a')$$

Q-learning uses *point estimates* on experience $\{s_t, a_t, c_t, s_{t+1}\}$:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \underbrace{\alpha_t}_{\text{Learning Rate}} \left[\underbrace{r_t + \gamma \max_{a' \in \mathcal{A}} Q(s_{t+1}, a')}_{\text{Target}} - Q(s_t, a_t) \right]$$

Q-learning: Algorithm


Q-learning


```
1  $t \rightarrow 0$ 
2 Initialize  $Q$  (e.g., zero-initialized)
3 Loop
4   choose  $a_t$  (e.g.,  $\epsilon$ -greedy or softmax selection)
5   observe next state  $s_{t+1}$  and reward  $r_t$ 
6    $Q(s_t, a_t) \leftarrow$   

    $Q(s_t, a_t) + \alpha_t [r_t + \gamma \max_{a' \in \mathcal{A}} Q(s_{t+1}, a') - Q(s_t, a_t)]$ 
7    $t \leftarrow t + 1$ 
8 EndLoop
```


Example: Maze

▶ `python maze.py --agent qlearning --episodes N
[-- plot_reward]`

 `maze.py`

 `qlearning.ipynb`

SARSA

- ▶ Q-learning is an **off-policy** algorithm
 - ▶ The algorithm uses the greedy policy to update Q , but likely chooses action according to another policy (e.g., ϵ -greedy)
- ▶ **SARSA**: **on-policy** algorithm similar to Q-learning
- ▶ The same policy is used to choose next action and to update Q

SARSA: Algorithm

SARSA

```
1  $t \rightarrow 0$ 
2 Initialize  $Q$  (e.g., zero-initialized)
3 choose  $a_t$  (e.g.,  $\epsilon$ -greedy or softmax selection)
4 Loop
5   observe next state  $s_{t+1}$  and reward  $r_t$ 
6   choose  $a_{t+1}$  (e.g.,  $\epsilon$ -greedy or softmax selection)
7    $Q(s_t, a_t) \leftarrow$   

    $Q(s_t, a_t) + \alpha_t [r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$ 
8    $t \leftarrow t + 1$ 
9 EndLoop
```

Dealing with Large State Spaces: Deep RL

Issues with Tabular RL

- ▶ So far, we have considered **tabular** representations of the value function

<i>State/Action</i>	a_1	a_2	...
s_1	$Q(s_1, a_1)$	$Q(s_1, a_2)$...
s_2	$Q(s_2, a_1)$	$Q(s_2, a_2)$...
...	...		
s_n	$Q(s_n, a_1)$	$Q(s_n, a_2)$...

- ▶ Not ideal as the state space grows...
- ▶ **Memory demand:** $\mathcal{O}(|\mathcal{S}||\mathcal{A}|)$
- ▶ **No generalization**
- ▶ How to handle **continuous state spaces?**

Value Function Approximation

Idea: using a **parametric approximation** of the value function

$$V_{\pi}(s) \approx \hat{V}(s, \mathbf{w}), \text{ or}$$

$$Q_{\pi}(s, a) \approx \hat{Q}(s, a, \mathbf{w})$$

- ▶ $\mathbf{w} \in \mathbb{R}^d$ is a vector of parameters
- ▶ We need to store \mathbf{w} instead of the Q table
 - ▶ Reduced memory demand if $d < |\mathcal{S}|$ ✓
- ▶ Potential generalization ✓
 - ▶ The experience gained in a state used to update \mathbf{w}
 - ▶ A single update possibly impacts the value of several states!
 - ▶ Can deal with continuous state spaces ✓

Value Function Approximation (2)

- ▶ How to choose a function \hat{Q} ?
- ▶ How to determine the value of w ?
- ▶ We search for a function and a vector w so as to approximate V (or Q) "well"
- ▶ First of all, what does "well" means?

Function Approximation: Objective

- ▶ A simple and natural choice is to minimize MSE:

$$J(\mathbf{w}) = \sum_{s \in \mathcal{S}} \mu(s) [V_{\pi}(s) - \hat{V}(s, \mathbf{w})]^2$$

- ▶ $\mu(s) \geq 0$ is a distribution over states
- ▶ $\mu(s)$ should reflect the importance or frequency of states

Optimizing Parameters

We can compute parameters \mathbf{w} through **gradient descent**

$$\begin{aligned}\mathbf{w}_{t+1} &= \mathbf{w}_t - \frac{1}{2}\alpha \nabla_{\mathbf{w}} J(\mathbf{w}_t) = \\ &= \mathbf{w}_t + \alpha \sum_{s \in \mathcal{S}} \mu(s) [V_{\pi}(s) - \hat{V}(s, \mathbf{w})] \nabla_{\mathbf{w}} \hat{V}(s, \mathbf{w})\end{aligned}$$

Two potential issues:

1. Summation over all states (may be expensive!)
2. We don't have the true values $V_{\pi}(s)$!

Optimizing Parameters: Issue 1

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \sum_{s \in \mathcal{S}} \mu(s) [V_\pi(s) - \hat{V}(s, \mathbf{w})] \nabla_{\mathbf{w}} \hat{V}(s, \mathbf{w})$$

Gradient computed over all states...

Stochastic gradient descent

one (or few) samples $(s_t, V_\pi(s_t))$ at each step

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [V_\pi(s_t) - \hat{V}(s_t, \mathbf{w})] \nabla_{\mathbf{w}} \hat{V}(s_t, \mathbf{w})$$

Optimizing Parameters: Issue 2

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \sum_{s \in \mathcal{S}} \mu(s) \left[V_{\pi}(s) - \hat{V}(s, \mathbf{w}) \right] \nabla_{\mathbf{w}} \hat{V}(s, \mathbf{w})$$

How to get exact values?

Stochastic semi-gradient descent:

we replace $V_{\pi}(s_t)$ with a noisy approximation U_t , based on estimated value func.

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \left[U_t - \hat{V}(s_t, \mathbf{w}) \right] \nabla_{\mathbf{w}} \hat{V}(s_t, \mathbf{w})$$

A possible approach (inspired by Q-learning):

$$U_t = r_t + \gamma \hat{V}(s_{t+1}, \mathbf{w}_t)$$

Linear Function Approximation

The simplest possible approximation model:

$$\hat{V}(s, \mathbf{w}) = \mathbf{w}^T \boldsymbol{\phi}(s) = \sum_{i=1}^d w_i \phi_i(s)$$

Weights $\mathbf{w} \in \mathbb{R}^d$

Features $\boldsymbol{\phi} : \mathcal{S} \rightarrow \mathbb{R}^d$

Update rule becomes very simple:

$$\nabla_{\mathbf{w}} \hat{V}(s, \mathbf{w}) = \boldsymbol{\phi}(s)$$
$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [U_t - \hat{V}(s_t, \mathbf{w}_t)] \boldsymbol{\phi}(s_t)$$

Linear Function Approximation (2)

We have equivalent formulas for Q :

$$\hat{Q}(s, a, \mathbf{w}) = \mathbf{w}^T \boldsymbol{\phi}(s, a) = \sum_{i=1}^d w_i \phi_i(s, a)$$

Weights $\mathbf{w} \in \mathbb{R}^d$

Features $\boldsymbol{\phi} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}^d$

$$\nabla_{\mathbf{w}} \hat{Q}(s, a, \mathbf{w}) = \boldsymbol{\phi}(s, a)$$
$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \left[U_t - \hat{Q}(s_t, a_t, \mathbf{w}_t) \right] \boldsymbol{\phi}(s_t, a_t)$$

??

Q-learning + Linear FA

Recall Q-learning update

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha_t [r_t + \gamma \max_{a' \in \mathcal{A}} Q(s_{t+1}, a') - Q(s_t, a_t)]$$


```
1  $t \rightarrow 0$ 
2 Initialize  $w$ 
3 Loop
4   choose action  $a_t$ 
5   gather experience  $\langle s_t, a_t, r_t, s_{t+1} \rangle$ 
6    $U_t \leftarrow r_t + \gamma \max_{a' \in \mathcal{A}} \hat{Q}(s_{t+1}, a', w_t)$ 
7    $w_{t+1} = w_t + \alpha [U_t - \hat{Q}(s_t, a_t, w_t)] \phi(s_t, a_t)$ 
8    $t \leftarrow t + 1$ 
9 EndLoop
```

(Linear) FA: Issues

- ▶ Linear FA+RL successfully applied on some tasks
- ▶ Nonlinear models (e.g., ANNs) have obtained significant results as well
 - ▶ e.g., [TD-Gammon](#) (1992)
- ▶ Efficacy of these approaches strongly depends on the [features](#) in use
 - ▶ how states (and actions) are represented
 - ▶ domain expertise necessary

Deep RL

- ▶ We have seen that the key advancement enabled by DNNs is the ability of **learning the features**
- ▶ Idea: exploiting this ability to learn suitable features for state and action representation



Deep Q Network

- ▶ First popular application of DNNs within RL in 2013
 - ▶ Mnih et al., “Playing Atari with Deep Reinforcement Learning”
<https://www.cs.toronto.edu/%7Evmnih/docs/dqn.pdf>
- ▶ Task: playing Atari 2600 games
- ▶ Two key innovations:
 - ▶ DNN to approximate Q (Deep Q Network)
 - ▶ Experience Replay buffer
- ▶ Learning algorithm adapted from Q-learning

Example: Atari games



Atari 2600 console (1977–1992)


Example: Atari games



Breakout: <https://www.youtube.com/watch?v=TmPfTpjtdgg>

Deep Q Network

- ▶ **Input:** state s (possibly preprocessed)
- ▶ **Output:** $\hat{Q}(s, a)$, for every action a



Training

- ▶ NN training usually based on (large) training set
 - ▶ collection of examples $(\mathbf{x}_i, \mathbf{y}_i)$
- ▶ To train a DQN we would need many examples $(s_i, [Q(s_i, a_1) \cdots Q(s_i, a_n)]^T)$
- ▶ **Problem:** we don't have true examples of $Q(s, a)$ to use!
 - ▶ agent only collects immediate rewards on-line
- ▶ We need to estimate Q on-line based on experience (as usual in RL)

Training (2)

Experience

$$\langle s_t, a_t, s_{t+1}, r_t \rangle$$

$$\langle s_{t-1}, a_{t-1}, s_t, r_{t-1} \rangle$$

$$\langle s_{t-2}, a_{t-2}, s_{t-1}, r_{t-2} \rangle$$

...

Training Sample

$$(s_t, a_t) \rightarrow r_t + \gamma \max_{a'} \hat{Q}(s_{t+1}, a', \mathbf{w})$$

$$(s_{t-1}, a_{t-1}) \rightarrow r_{t-1} + \gamma \max_{a'} \hat{Q}(s_t, a', \mathbf{w})$$

$$(s_{t-2}, a_{t-2}) \rightarrow r_{t-2} + \gamma \max_{a'} \hat{Q}(s_{t-1}, a', \mathbf{w})$$

- ▶ Naive idea: pick mini-batches of last b experience tuples and train the NN
 - ▶ i.e., at each iteration, train on most recent experience
- ▶ sequential observations likely correlated ✗
- ▶ less recent experience possibly forgotten ✗

Experience Replay

- ▶ Smarter approach: **experience replay** buffer
- ▶ Circular FIFO buffer with capacity $B > b$
- ▶ At each training iteration, b tuples drawn randomly from the buffer
- ▶ correlation between observations reduced/removed ✓
- ▶ if B is large, old observations are “seen” more than once ✓
 - ▶ improved data efficiency

Deep Q-learning (DQL)

```
1 Initialize  $\mathbf{w}$ 
2 Initialize empty buffer  $\mathcal{B}$ 
3  $i \leftarrow 0$ 
4 Loop
5   choose action  $a_i$ 
6   gather experience  $\langle s_i, a_i, r_i, s_{i+1} \rangle$  and add to  $\mathcal{B}$ 
7   sample minibatch of  $b$   $\langle s_j, a_j, r_j, s_{j+1} \rangle$  tuples from  $\mathcal{B}$ 
8    $y^{(j)} \leftarrow r_j + \gamma \max_{a'} \hat{Q}(s_{i+1}, a', \mathbf{w}), j = 1, \dots, b$ 
9    $\mathcal{L}^{(j)} = (y^{(j)} - \hat{Q}(s_j, a_j, \mathbf{w}))^2$  /* Loss */
10  update  $\mathbf{w}$  using, e.g., SGD on the minibatch
11   $i \leftarrow i + 1$ 
12 EndLoop
```

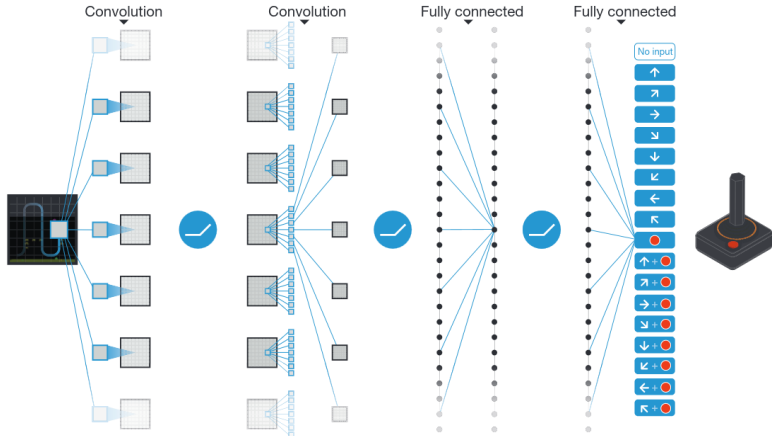

Example: Atari

- ▶ Frames are 210×160 pixel images with a 128 color palette
- ▶ Input dimensionality reduced via preprocessing
 - ▶ RGB to gray-scale conversion
 - ▶ down-sampling to 110×84
 - ▶ cropped to 84×84 to ease implementation
- ▶ State comprises last 4 frames
 - ▶ why?

Example: Atari

- ▶ NN input: $84 \times 84 \times 4$ image produced by preprocessing
- ▶ Conv. layer with 16 8×8 filters with ReLU
- ▶ Conv. layer with 32 4×4 filters with ReLU
- ▶ Fully-connected layer with 256 ReLU units
- ▶ Linear output layer with one unit for each valid action (from 4 to 18 in the considered games)
- ▶ Trained using RMSProp for a total of 50 million frames (around 38 days of game experience in total)
- ▶ Replay memory stores 1 million most recent frames

Example: Atari



Target Network

- ▶ DQN may suffer from instability during training, possibly preventing the algorithm to converge
- ▶ In traditional NN training, the training targets do not change over time
- ▶ In DRL, since we don't have ground-truth Q values, we use the approximated \hat{Q} in the update target value:

$$y^{(j)} \leftarrow r_j + \gamma \min_{a'} \hat{Q}(s_{i+1}, a', \mathbf{w})$$

- ▶but we keep changing \mathbf{w} at each iteration
- ▶ Let's use a second neural network to stabilize the targets

Deep Q-learning with Target Network

- 1 Initialize \mathbf{w} and $\mathbf{w}^- = \mathbf{w}$
- 2 Initialize empty buffer \mathcal{B}
- 3 $i \leftarrow 0$
- 4 **Loop**
 - 5 choose action a_i
 - 6 gather experience $\langle s_i, a_i, r_i, s_{i+1} \rangle$ and add to \mathcal{B}
 - 7 sample minibatch of b $\langle s_j, a_j, r_j, s_{j+1} \rangle$ tuples from \mathcal{B}
 - 8 $y^{(j)} \leftarrow r_j + \gamma \min_{a'} \hat{Q}(s_{j+1}, a', \mathbf{w}^-), j = 1, \dots, b$
 - 9 $\mathcal{L}^{(j)} = (y^{(j)} - \hat{Q}(s_j, a_j, \mathbf{w}))^2$
 - 10 update \mathbf{w} using, e.g., SGD on the minibatch
 - 11 every C steps: $\mathbf{w}^- \leftarrow \mathbf{w}$
 - 12 $i \leftarrow i + 1$
- 13 **EndLoop**

Remark

- ▶ DQN can seamlessly work with **continuous** state spaces
- ▶ Action space must be finite

Example: CartPole with DQN

- ▶ Environment provided by [OpenAI Gym](#)
 - ▶ Large collection of ready-to-use environments
- ▶ DQN implemented using [TF-Agents](#)
 - ▶ RL library part of Tensorflow ecosystem
- ▶ https://www.tensorflow.org/agents/tutorials/1_dqn_tutorial?hl=en

Policy-based RL

- ▶ So far, we have considered **value-based** RL algorithms
 - ▶ Learn the value function; get a policy from it
- ▶ Now we turn our attention to **policy-based** RL (or, **policy gradient methods**)
 - ▶ Directly learn a **policy**
 - ▶ Algorithms may still learn the value function, but it is not used to derive the policy

Policy Gradient Methods

- ▶ Algorithms learn a **parameterized policy**

$$\pi(a|s, \theta) = P(A_t = a | S_t = s, \theta_t = \theta)$$

$\theta \in \mathbb{R}^m$ is the vector of policy parameters

- ▶ $\pi(a|s, \theta)$ can be any function, as long as it is differentiable w.r.t. parameters θ

Note

To avoid ambiguity, we will keep using $\mathbf{w} \in \mathbb{R}^d$ to denote the vector of parameters used to approximate the value function, if necessary (e.g., $V(s, \mathbf{w})$)

Policy Gradient Methods (2)

- ▶ Suppose that $J(\theta)$ is a performance measure of the policy resulting from parameters θ (the higher the better)
- ▶ To maximize performance, we can update θ by **gradient ascent**:

$$\theta_{t+1} = \theta_t + \alpha \nabla J(\theta_t)$$

- ▶ The expression “policy gradient” refers to all the methods based on the idea introduced above

Policy Approximation: Why?

- ▶ Often (but not always), the policy is an easier function to approximate compared to the value function
- ▶ Policy parameterization lets action probabilities change smoothly as a function of the learned parameters, while they can change dramatically for a small change in the action values (if a different action gets the highest value)
 - ▶ stronger convergence guarantees are available
- ▶ Stochastic policies can be learned
- ▶ Continuous action spaces are supported

Policy Approximation via Action Preferences

- ▶ Let's suppose that the action space is discrete (and not too large)
 - ▶ We will discuss later other scenarios
- ▶ A natural choice for policy approximation is **softmax in action preferences**:

$$\pi(a|s, \theta) = \frac{e^{h(s,a,\theta)}}{\sum_{a'} e^{h(s,a',\theta)}}$$

- ▶ $h(s, a, \theta)$ is a parameterized numerical preference value for every state-action pair

Policy Approximation via Action Preferences

- ▶ Note that we don't need any specific strategy to determine the preference values $h(s, a, \theta)$
- ▶ They are just a convenient way to parameterize the policy $\pi(a|s, \theta)$
- ▶ For instance, we could use a DNN with a softmax output layer to approximate the policy
 - ▶ Hidden layers compute the action preferences
 - ▶ Output layer produces the policy probabilities

Action Preferences vs. Action Values

- ▶ Action values $Q(s, a)$ may differ by a small amount
 - ▶ Softmax based on Q may struggle to approach a deterministic policy (unless a very small temperature coefficient is used)
- ▶ Action preferences instead do not need to converge to specific values (e.g., the optimal value function), but rather to the best values for the policy to learn
 - ▶ if a deterministic policy is optimal, preference for the optimal action will be as high as possible than the other actions
 - ▶ if a stochastic policy is optimal, more than one action will have a high preference value (e.g., card games with incomplete information) ...
 - ▶ ...and we can learn *arbitrary* probabilities for actions

Policy Gradient in Episodic Tasks

- ▶ Let's consider an **episodic** task starting in state s_0
- ▶ In this case, performance of the policy can be evaluated as

$$J(\theta) = V_{\pi_\theta}(s_0)$$

- ▶ **How to update θ to improve performance?**
- ▶ Performance depends both on (1) action selection and (2) the distribution of states occurring in the episode
- ▶ Both depend on the parameters!
- ▶ (2) is particularly difficult as it also depends on the environment

Policy Gradient Theorem

Policy Gradient Theorem

$$\nabla_{\theta} J(\theta) \propto \sum_s \mu(s) \sum_a Q_{\pi}(s, a) \nabla_{\theta} \pi(a|s, \theta)$$

- ▶ Proportionality constant is the average length of an episode
 - ▶ in gradient ascent, constant absorbed by step size α ✓
- ▶ we don't need the derivative of $\mu(s)$ ✓

Policy Gradient Theorem: Proof

To simplify notation, we leave it implicit that π is a function of θ , and that gradients are w.r.t. θ

$$\begin{aligned}\nabla V_{\pi}(s) &= \nabla \left[\sum_a \pi(a|s) Q_{\pi}(s, a) \right] = \\ &= \sum_a \nabla [\pi(a|s) Q_{\pi}(s, a)] = \\ &= \sum_a [\nabla \pi(a|s) Q_{\pi}(s, a) + \pi(a|s) \nabla Q_{\pi}(s, a)] = \\ &= \sum_a \left[\nabla \pi(a|s) Q_{\pi}(s, a) + \pi(a|s) \nabla \sum_{s', r'} p(s', r'|s, a) (r + V_{\pi}(s')) \right] =\end{aligned}$$

Proof (2)

$$= \sum_a \left[\nabla \pi(a|s) Q_\pi(s, a) + \pi(a|s) \nabla \sum_{s', r'} p(s', r'|s, a) (r + V_\pi(s')) \right] =$$

1) Reward does not depend on θ (gradient is 0)

2) $\sum_{s'} \sum_{r'} p(s', r'|s, a) V_\pi(s') = \sum_{s'} p(s'|s, a) V_\pi(s')$

$$= \sum_a \left[\nabla \pi(a|s) Q_\pi(s, a) + \pi(a|s) \sum_{s'} p(s'|s, a) \nabla V_\pi(s') \right] =$$

Note: we are computing $\nabla V_\pi(s)$ and now we have a recursive term $\nabla V_\pi(s')$! Let's unroll the recursion...

Proof (3)

$$= \sum_a \left[\nabla \pi(a|s) Q_\pi(s, a) + \pi(a|s) \sum_{s'} p(s'|s, a) \cdot \sum_{a'} \left(\nabla \pi(a'|s') Q_\pi(s', a') + \pi(a'|s') \sum_{s''} p(s''|s', a') \nabla V_\pi(s'') \right) \right] =$$

After repeated unrolling ...

$$= \sum_{x \in \mathcal{S}} \sum_{k=0}^{\infty} P(s \rightarrow x, k, \pi) \sum_a [\nabla \pi(a|x) Q_\pi(x, a)]$$

where $P(s \rightarrow x, k, \pi)$ is the probability of transitioning from s to x in k steps under policy π .

Proof (4)

We can now write an expression for the gradient of J

$$\begin{aligned}\nabla J(\theta) &= \nabla V_{\pi}(s_0) = \sum_s \sum_{k=0}^{\infty} P(s_0 \rightarrow s, k, \pi) \sum_a [\nabla \pi(a|s) Q_{\pi}(s, a)] = \\ &= \sum_s \eta(s) \sum_a [\nabla \pi(a|s) Q_{\pi}(s, a)] =\end{aligned}$$

$\eta(s)$: avg. number of steps spent in s within an episode

$$= \sum_{s'} \eta(s') \sum_s \frac{\eta(s)}{\sum_{s'} \eta(s')} \sum_a [\nabla \pi(a|s) Q_{\pi}(s, a)] =$$

Proof (5)

$$\begin{aligned} &= \sum_{s'} \eta(s') \sum_s \frac{\eta(s)}{\sum_{s'} \eta(s')} \sum_a [\nabla \pi(a|s) Q_\pi(s, a)] = \\ &= \sum_{s'} \eta(s') \sum_s \mu(s) \sum_a [\nabla \pi(a|s) Q_\pi(s, a)] \\ \nabla J(\theta) &\propto \sum_s \mu(s) \sum_a [\nabla \pi(a|s) Q_\pi(s, a)] \end{aligned}$$

REINFORCE

- ▶ $\mu(s)$ is the on-policy distribution of states under π
 - ▶ if π is followed, states will occur in that proportion

$$\begin{aligned}\nabla_{\theta} J(\theta) &\propto \sum_s \mu(s) \sum_a Q_{\pi}(s, a) \nabla_{\theta} \pi(a|s, \theta) = \\ &= E_{\pi} \left[\sum_a Q_{\pi}(s_t, a) \nabla_{\theta} \pi(a|s_t, \theta) \right]\end{aligned}$$

- ▶ as we did to approximate the value function, we can perform a **stochastic gradient** ascent on occurring states s_t

REINFORCE (2)

- ▶ We replaced a sum over states with an expectation under π
- ▶ We want to do the same with the sum over actions
- ▶ First, we need to weigh actions by $\pi(a|s, \theta)$

$$\begin{aligned}\nabla_{\theta} J(\theta) &\propto E_{\pi} \left[\sum_a Q_{\pi}(s_t, a) \nabla_{\theta} \pi(a|s_t, \theta) \right] = \\ &= E_{\pi} \left[\sum_a \pi(a|s_t, \theta) Q_{\pi}(s_t, a) \frac{\nabla_{\theta} \pi(a|s_t, \theta)}{\pi(a|s_t, \theta)} \right] = \\ &= E_{\pi} \left[Q_{\pi}(s_t, a_t) \frac{\nabla_{\theta} \pi(a_t|s_t, \theta)}{\pi(a_t|s_t, \theta)} \right] = \\ &= E_{\pi} \left[G_t \frac{\nabla_{\theta} \pi(a_t|s_t, \theta)}{\pi(a_t|s_t, \theta)} \right]\end{aligned}$$

REINFORCE (3)

$$\nabla J(\theta) \propto E_{\pi} \left[G_t \frac{\nabla_{\theta} \pi(a_t | s_t, \theta)}{\pi(a_t | s_t, \theta)} \right]$$

- ▶ We can sample G_t for each time step, and we got an expression proportional to the gradient
- ▶ We can use stochastic gradient ascent to update parameters
- ▶ The resulting algorithm is **REINFORCE** (1992)

$$\theta_{t+1} \leftarrow \theta_t + \alpha G_t \nabla_{\theta} \ln \pi(a_t | s_t, \theta)$$

$$\frac{\nabla_{\theta} \pi(a_t | s_t, \theta)}{\pi(a_t | s_t, \theta)} = \nabla_{\theta} \ln \pi(a_t | s_t, \theta)$$

REINFORCE: Algorithm

We also include a discount factor γ

```
1 Initialize  $\theta$  (e.g., to 0)
2 Loop
3   generate episode  $s_0, a_0, r_1, s_1, \dots, r_T$  following  $\pi$ 
4   for  $t=0,1,\dots,T$  do
5      $G_t \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} r_k$ 
6      $\theta \leftarrow \theta + \alpha \gamma^t G_t \nabla_{\theta} \ln \pi(a_t | s_t, \theta)$ 
7   end
8 EndLoop
```

Policy Gradient: Another Perspective

- ▶ Consider a NN used for multi-class classification
- ▶ Softmax final layer outputs a probability y_c for all classes c
- ▶ Gradient of cross-entropy loss used for training

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} \left[\sum_c \bar{y}_c \ln y_c \right]$$

- ▶ Intuitively, training will increase prob. y_c for the class c labeled as correct (ground truth)
 - ▶ likelihood maximization
- ▶ Replacing classes with “actions”, the NN outputs $\pi(a|s, \theta)$
- ▶ REINFORCE update is proportional to $\nabla_{\theta} \ln \pi(a_t|s_t, \theta)$
 - ▶ without a ground truth, prob. is increased/decreased based on return

REINFORCE with Baseline

- ▶ A slight generalization of REINFORCE involves the use of a **baseline** $b(s)$
- ▶ We compare the value of each action to $b(s)$
- ▶ The policy gradient theorem remains true

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a (Q(s, a) - b(s)) \nabla \pi(a|s, \theta)$$

- ▶ Baseline can be useful to reduce the variance of the update and speed learning

$$\theta_{t+1} = \theta_t + \alpha (G_t - b(s_t)) \nabla \ln \pi(a_t|s_t, \theta)$$

Actor-Critic

- ▶ Idea to avoid high variance of returns used by REINFORCE
- ▶ Using the **one-step return** $G_{t:t+1}$ instead

$$G_{t:t+1} = r_t + \gamma V(s_{t+1})$$

- ▶ We call **critic** the role of the value function used in this way
- ▶ We call the resulting approach **actor-critic**

$$\begin{aligned}\theta_{t+1} &= \theta_t + \alpha (G_{t:t+1} - \hat{V}(s_t, \mathbf{w})) \nabla \ln \pi(a_t | s_t, \theta) = \\ &= \theta_t + \alpha (r_{t+1} + \gamma \hat{V}(s_{t+1}, \mathbf{w}) - \hat{V}(s_t, \mathbf{w})) \nabla \ln \pi(a_t | s_t, \theta)\end{aligned}$$

Actor-Critic (2)

```
1 Initialize  $\theta$  and  $w$  (e.g., to 0)
2 Loop
3   Initialize  $s$  as first state of the episode
4    $l \leftarrow 1$ 
5   while  $s$  not terminal do
6     choose action  $a$  according to  $\pi(\cdot|s, \theta)$ 
7     observe  $s'$  and  $r$ 
8      $\delta \leftarrow r + \gamma \hat{V}(s', w) - \hat{V}(s, w)$ 
9      $w \leftarrow w + \alpha^w \delta \nabla \hat{V}(s, w)$ 
10     $\theta \leftarrow \theta + l \alpha^\theta \delta \nabla \ln \pi(a|s, \theta)$ 
11     $s \leftarrow s'$ 
12  end
13 EndLoop
```

The Continuing Case

- ▶ Policy gradient theorem holds for continuing tasks as well
- ▶ The performance measure $J(\theta)$ must be changed to the [average reward](#)
- ▶ Proof and updated Actor-Critic alg. in Sutton-Barto, 13.6

Continuous or Large Action Spaces

- ▶ Policy gradient methods can be useful in presence of very large or continuous action spaces
 - ▶ Computing the learned probability for every action can be expensive/unfeasible
- ▶ The solution: learn the parameters of a probability distribution, instead of the probability of choosing each action
- ▶ Example: policy defined as the normal probability density

$$\pi(a|s, \theta) = \frac{1}{\sigma(s, \theta)\sqrt{2\pi}} \exp\left(-\frac{(a - \mu(s, \theta))^2}{2\sigma(s, \theta)^2}\right)$$

Example

CartPole with Actor-Critic in TensorFlow:

https://www.tensorflow.org/tutorials/reinforcement_learning/actor_critic?hl=en

Advanced Policy Methods

- ▶ **Deterministic Policy Gradient (DPG)**: similar to the theorem proven above, but for deterministic policies
 - ▶ Silver et al. (2014), “Deterministic Policy Gradient Algorithm”,
<http://proceedings.mlr.press/v32/silver14.pdf>
- ▶ **Deep Deterministic Policy Gradient (DDPG)**: DPG with DNNs
 - ▶ Lillicrap et al. (2016), “Continuous control with deep reinforcement learning”,
<https://arxiv.org/abs/1509.02971>
- ▶ ...
- ▶ **Proximal Policy Optimization (PPO)**: state-of-the-art algorithm
 - ▶ Schulman et al., “Proximal Policy Optimization Algorithms”,
<https://arxiv.org/pdf/1707.06347.pdf>

Advanced RL Topics

- ▶ Multi-agent RL
- ▶ Hierarchical RL
- ▶ RL + Heuristic Tree Search
- ▶ Transfer RL
- ▶ ...