

Hardware-based Attacks and Countermeasures

Alessandro Pellegrini
a.pellegrini@ing.uniroma2.it

Setting up Reading Primitives

Time and Security

- The time taken to perform some operation on data might reveal information on the data itself

```
int strcmp(char *t, char *s)
{
    for( ; *t == *s ; s++, t++)
        if(*t == '\0')
            return 0;
    return *t - *s;
}
```

- Comparing “ABCD” with “ABxDE” will return after three iterations

Timing the Cache

- When reading data, the CPU actually reads the data from the cache
- Cache protocols try to serve a memory request as fast as possible
- If we can control the state of the cache, we can exploit the time required to serve a memory request to leak data
- Thanks to SMT, we can leak data from L1 caches
- Since lower level of caches are shared, we can also leak data from lower level caches
- The fundamental aspect relates to *how* we can control the state of the cache

Discovering Code Paths

Montgomery Ladder

Input: Point P , scalar n , k bits

Output: Point nP

$R_0 \leftarrow \mathcal{O}$

$R_1 \leftarrow P$

for i from k to 0 **do**

if $n_i = 0$ **then**

$R_1 \leftarrow R_0 + R_1$

$R_0 \leftarrow 2R_0$

else

$R_0 \leftarrow R_0 + R_1$

$R_1 \leftarrow 2R_1$

end

end

} cache line A

} cache line B

} cache line C

} cache line D

- Scalar multiplication is common in cryptosystems
- Nonces should remain secret to prevent reconstructing keys
- We cannot time this algorithm: the number of operations is even
- We can anyhow determine *which branch* of the if statement is taken *for each bit*

Side-channel Attacks

- Side channel: *memory which allows to read other memory content or detect access patterns*
- There are various ways to mount a side channel:
 - Prime + Probe
 - Flush + Reload
 - Flush + Flush
 - Evict + Time
 - Evict + Reload
 - Prime + Abort
- The key idea is the same: bring the cache into a known state and observe side effects generated by other processes
- In this way, we can leak data bypassing OS-based process isolation

Side-channel Attacks

- Basic construction of an attack follows through several phases
 1. *Pre-attack*: the target is acquired (single cache line, or cache set) and timing thresholds are established if needed
 2. *Active attack*:
 - a) Initialization: bring the channel in a known state
 - b) Wait for the victim to make an access to memory
 - c) Analyze the access, by observing side effects left by the victim
 - d) Repeat until all the data are leaked
- Aspects to take into account:
 - Caches are cached either virtually or physically
 - Caches are shared differently depending on the level
 - Interleave of execution between different processes

Evict + Time

- This approach uses the targeted eviction of lines, together with overall execution time measurement
- The attacker first causes the victim to run, preloading its working set, and establishing a baseline execution time
- The attacker then evicts a line of interest, and runs the victim again
- A variation in execution time indicates that the line of interest was accessed

esempio c: lavoriamo con cache L1 perchè facciamo flush di indirizzo logico, che lo ha la cache L1. Inoltre attaccante e attaccato lavorano su stesso processore e stessa cpu. Funziona bene con shared library (così fu rotto AES). Nell'esempio il target è deterministico, se così non fosse, sarebbe più complicato. Inoltre, nel loop, carichiamo il target che vogliamo, ma se così non fosse?

Flush + Reload

- Relies on the existence of shared virtual memory (such as shared libraries or page deduplication), and the ability to flush by virtual address
- The attacker first flushes a shared line of interest (by using dedicated instructions or by eviction through contention)
- Once the victim has executed, the attacker then reloads the evicted line by touching it, measuring the time taken
- A fast reload indicates that the victim touched this line (reloading it), while a slow reload indicates that it didn't

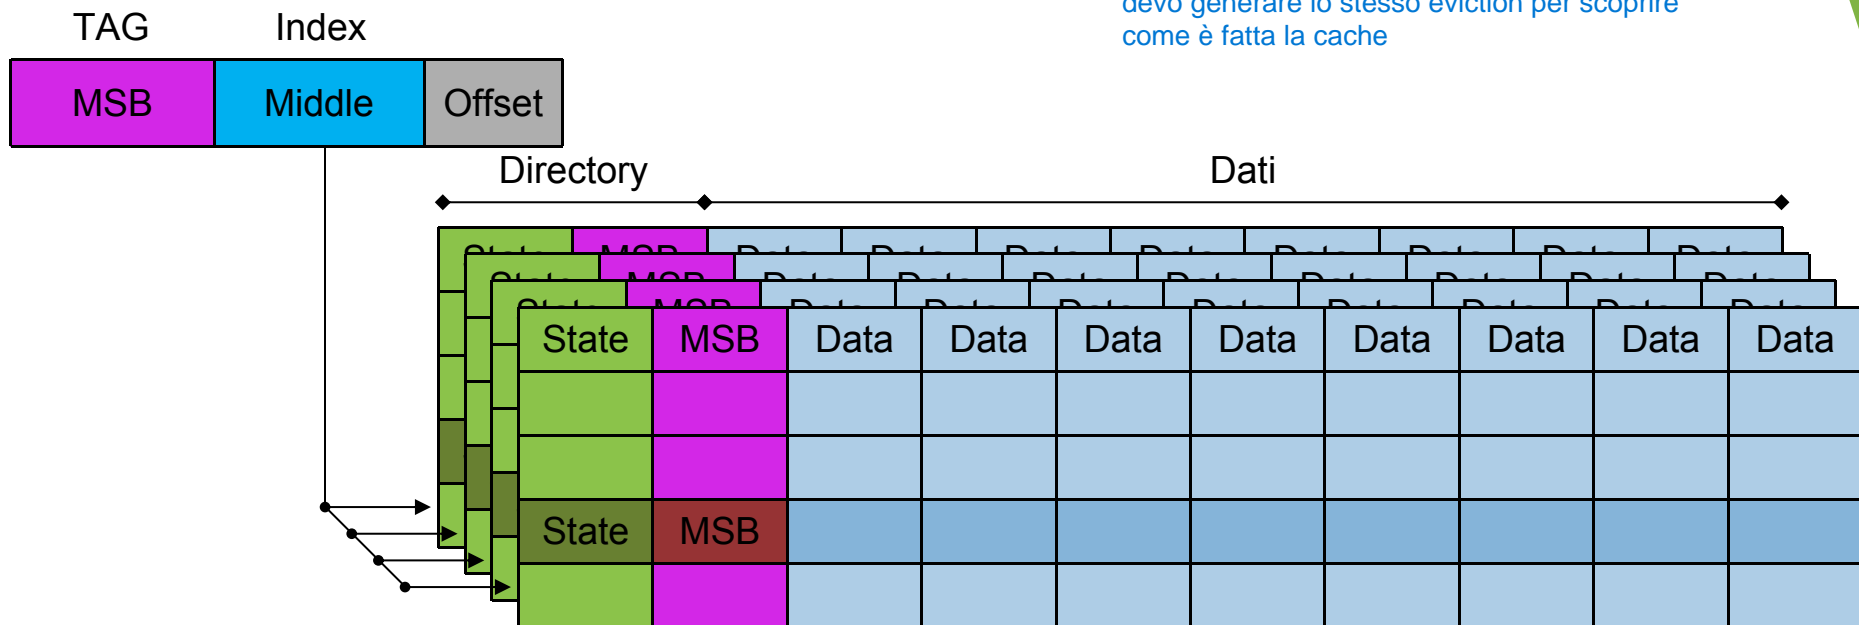
Prime + Probe

- The oldest and largest family of cache attacks, and the most general
- The original attack targeted L1, but it can be extended to L3 (without the need to rely on SMT)
 - L3 Prime + Probe can also work across VMs
- This technique detects the eviction of the attacker's working set by the victim:
 - The attacker first *primes* the cache by filling one or more sets with its own lines
 - Once the victim has executed, the attacker *probes* by timing accesses to its previously-loaded lines, to see if any were evicted
 - If so, the victim must have touched an address that maps to the same set

Prime + Probe

- A successful attack is not that simple (anymore)
 - The attacker requires sets of **colliding memory addresses (eviction sets)**
 - In conventional caches, the mapping of memory to cache is static

devo generare io stesso eviction per scoprire
come è fatta la cache



Prime + Probe: A Hardware Countermeasure

- Modern caches define the mapping of memory-to-cache at runtime
- This mapping changes over time
- The attacker has therefore a reduced time window in which to create the eviction set:
 - Choose randomly N memory addresses
 - Load them into the cache by accessing them
 - There is a possibility of self-collision
 - These self-colliding addresses should be removed from the set
- At this point, the classical attack can take place

Prime + Abort

- A Hardware Transaction is **aborted if there is a conflict in the write set**, *independently* of whether it happened in a transactional context
- The target is a single cache line (attacker must know a precise address to target)
- Simplest naïve attack:
 - Start a transaction
 - Access the target transactionally
 - Wait for an abort
 - **On abort, we know the address was accessed by another process**
- TSX works as an effective *hardware callback*
 - No need for a “timing phase”

Prime + Abort

- Prime + Abort on L1 cache:
 - A cacheline which has been *written* during the transaction (i.e., a cacheline in the transaction's write set) is evicted from the L1 cache
 - The transactional access entails a write in memory
 - We can therefore monitor evictions from L1
 - We can only spy on threads running on the same core, thanks to SMT
- Prime + Abort on L3 cache:
 - A cacheline which has been *read* during the transaction (i.e., a cacheline in the transaction's "read set") is evicted from the L3 cache
 - The transactional access entails reading from memory

Flush + Flush

- It is essentially a variant of Flush + Reload
- This attack is based on the fact that flushing a line can have different timings as well
 - `clflush` takes a different time to execute depending on whether a line is available in cache or not
- Stealthy: it's more difficult to detect flush della linea è ciò che fanno le app, ad esempio per i dataset. Quindi è difficile da capire chi fa cosa
- Faster: performance measures tell that Flush + Flush is faster than other cache attacks

Meeting Out-of-order Pipelines

- A speculative processor can let an attacker exploit micro-architectural effects by using *phantom instructions* even though the pipeline is flushed
- Phantom instructions leave effects in the micro-architectural state for a transient period
- These effects can be observed: *Transient Execution Attacks*
- This is the rationale behind *Spectre/Meltdown* attacks
 - It affects Intel, AMD and ARM processors

metto in pipeline istruzioni non valide, le quali generano effetti sulla cpu

Meltdown Primer

trovare l'address è un altro tipo di problema, che non vediamo qui



```
uint8_t *probe_array = new uint8_t[256 * 4096];  
// Make sure probe_array is not cached  
uint8_t kernel_memory = *(uint8_t*)(kernel_address);  
uint64_t final_kernel_memory = kernel_memory * 4096;  
uint8_t dummy = probe_array[final_kernel_memory];  
// handle (eventual) SEGFAULT  
// determine which of 256 slots in probe_array is cached
```

possible byte values

Intel OOO Execution

Avoid prefetching
& different lines

- Simply measuring the latency to access `probe_array` at the end of this execution path allows to know the value of the kernel memory byte

Fooling the Branch Prediction Unit

- The BPU will “learn” from recent past, independently of how complex is its organization
- Branch prediction units can be *poisoned*:
 - An attacker repeatedly executes a certain pattern of code which makes the BPU learn a certain outcome
 - Once the BPU becomes stable, the code alters its behavior
 - The BPU will continue to make the same prediction
 - The OOO pipeline might allow unauthorized code pass through the pipeline
 - This will be detected eventually, but micro-architectural side effects will still be around

Spectre Primer v1



not available in cache allora devo indovinare, quindi ritardo il momento in cui in pipeline si farà il check.

```
if(x < array1_size) {  
    y = array2[array1[x] * 4096];  
}
```

speculate through (vendor specific)

byte of interest

address available in cache

finally out of bounds

- It is then possible to inspect the cache state of `array2` to see what was the speculatively accessed value of `array1[x]`.

Si può avere segfault, ma sappiamo che possiamo catturarla e gestirla.

Spectre Primer v2

individuo delle shared library (quindi codice esistente), per identificare un gadget, utile per fare qualcos'altro diverso dal suo scopo originale.

- The attacker chooses a *gadget* from the victim's address space
 - This gadget can be anywhere, also in eBPF programs!
 - It must be in the victim's executable address space
- The **Branch Target Buffer** is **poisoned** to mispredict an indirect branch instruction to the address of the gadget, resulting in its speculative execution
- There is **no vulnerability in victim's code here**
 - The code is correct, it's only exploited to run according to the attacker's plan
- There are several megabytes of shared libraries where to look a gadget for

Spectre Primer v2



- A possible gadget (taken from Windows 10):

```
adc edi,dword ptr [ebx+edx+13BE13BDh]  
adc dl,byte ptr [edi]
```

con questo valore nelle [...].
potrei arrivare in zone kernel
- The attacker can:
 - set `edi` to the base address of a probe array
 - this can be a memory region in a shared library
 - $ebx = m - 0x13BE13BD - edx$
- The first instruction in the gadget will read 32 bits from m and add it into `edi`
- The second instruction fetches the index m in the probe array into the cache

cosa posso fare con questo attacco?

Mitigating Side Channel Attacks

Possible Mitigations and Detection

- Timing attacks are difficult to mitigate
 - The behaviour of the caches should be modified
 - This would require a constant-time cache, with no difference between hit/miss time
 - This would be a fundamental performance drop, making the caching architecture mostly pointless

Mitigations: the hard way

- Remove or restrict access to high-resolution timers such as `rdtsc`
 - unlikely: necessary to benchmark various hardware properties
- Allow certain memory to be marked as uncacheable
 - hardware challenge!
- Use AES-NI instructions in Intel chips to compute AES
 - but what about other encryption algorithms?
- Scatter-gather: secret data should not be allowed to influence memory access at coarser-than-cache-line granularity. oscurare non è mai la soluzione
- Disable TSX (as they did!) Transactional Memory, cioè Transactional Synchronization Extensions
- Disable OOO pipelines (performance-critical!)

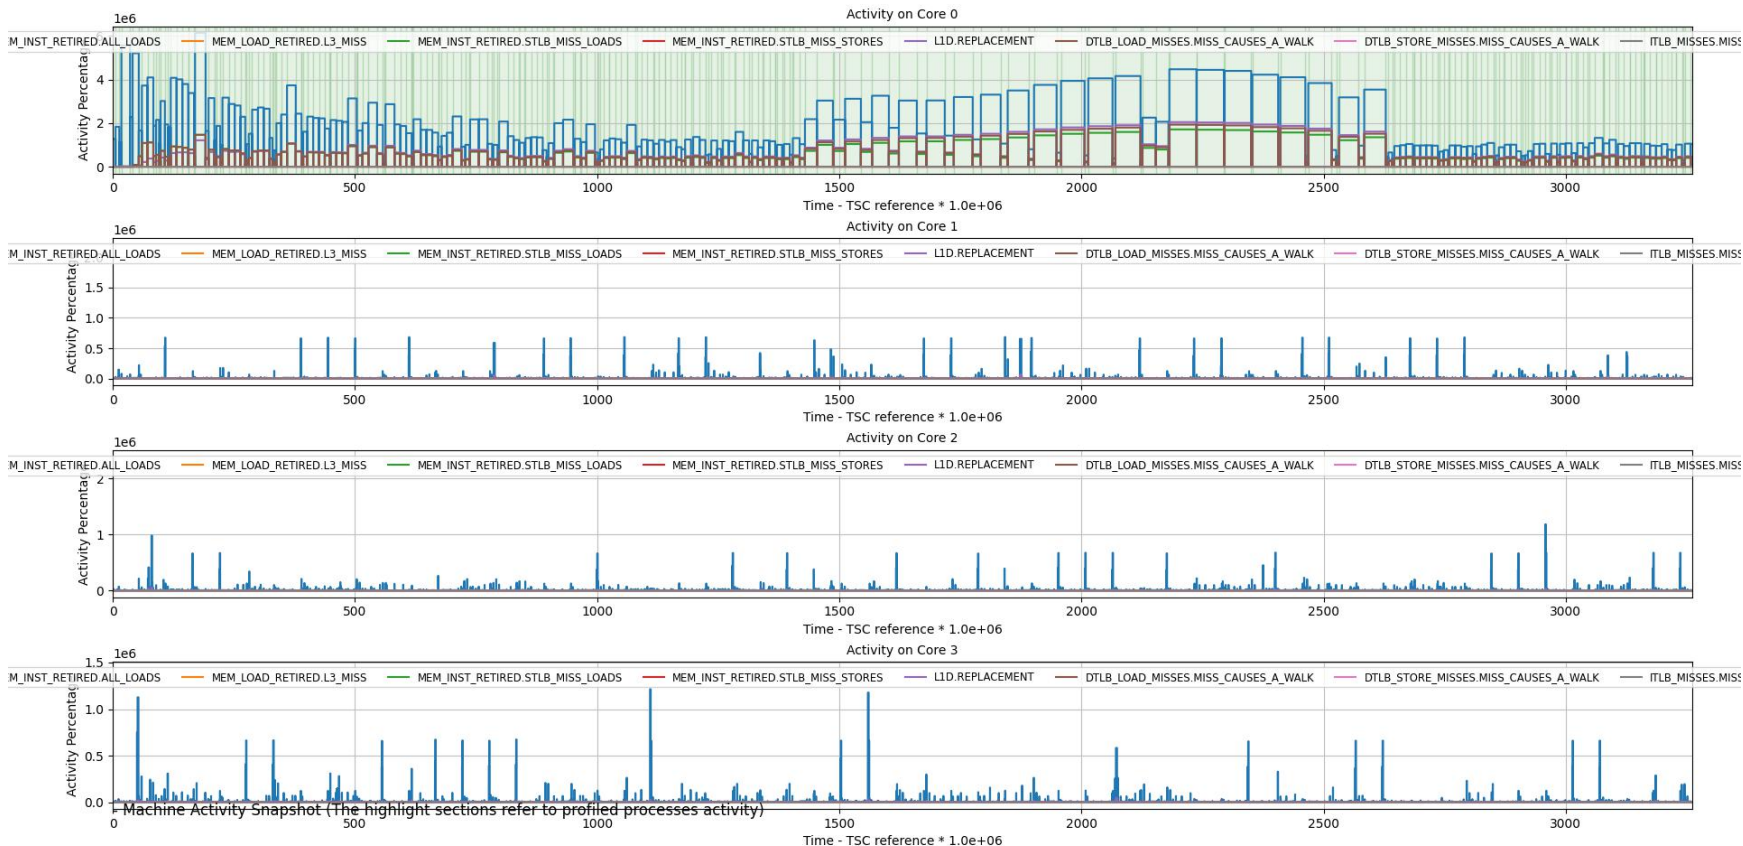
Mitigations to L1 attacks

- Cross-process L1 attacks are only possible when two separate threads are sharing the same L1
- L1 are private per-core
- Therefore, L1 attacks are only viable if relying on Simultaneous Multi-Threading
- The solution is therefore to turn off SMT

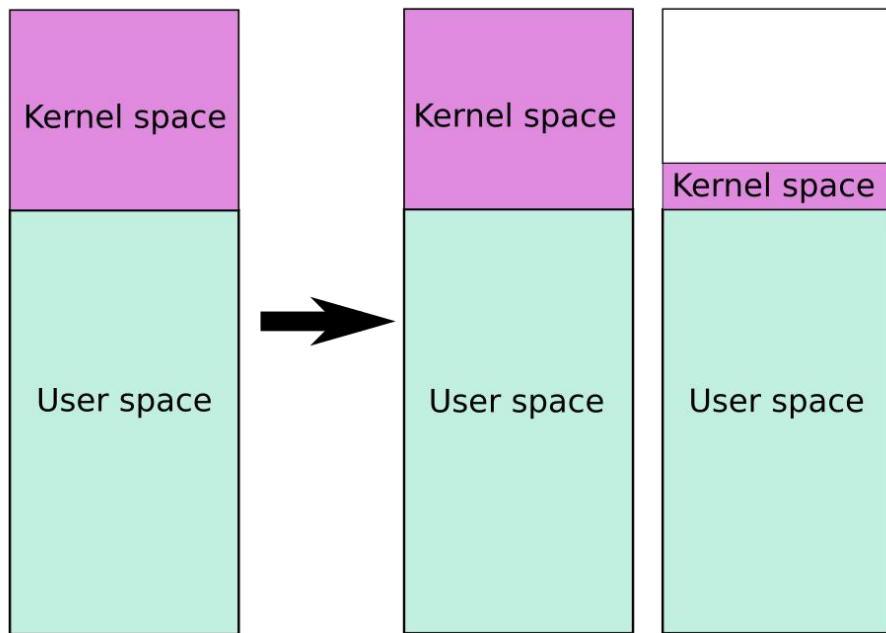
Mitigations: the detection way

- Setting up a side channel creates noise on the caching hierarchy
 - This is especially true for the *prepare* phase of some attacks
- Modern CPUs are equipped with **Hardware Performance Counters**
- They allow to track micro-architectural level events and count them
- They can be used to *observe* the behaviour of applications running in the system, and *selectively activate* countermeasures or mitigations
- Unfortunately, performance counters are not stable across vendors and different generations of CPUs

Mitigations: the detection way



Meltdown Mitigation: Kernel Page Table Isolation



non vedo locazione
della parte più grande
del kernel quando sono user.

User mode
Kernel mode
traditional

Kernel mode

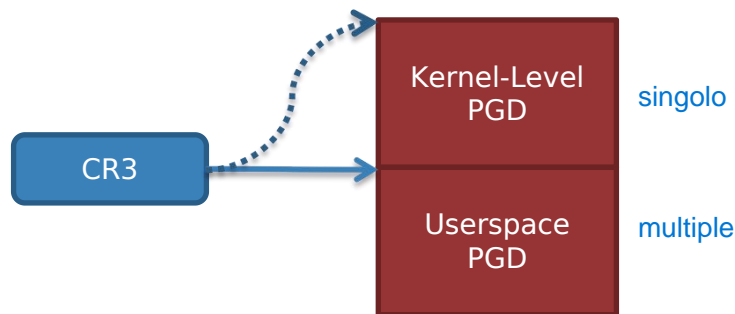
User mode

cpu_entry_area

```
struct cpu_entry_area {  
    char gdt[PAGE_SIZE];  
  
    struct entry_stack_page entry_stack_page;  
  
    struct tss_struct tss;  
  
    char exception_stacks[...];  
};  
  
DECLARE_PER_CPU(struct cpu_entry_area *, cpu_entry_area);
```

Double Page General Directory

- The first level of the page table is composed of a buffer of 8 KBs (two actual pages) 4 Kb è la dimensione minima vista dal s.o.
- One page is used to map the kernel-level memory view
- The other one is used to map the userspace memory view



CR3 is updated when transitioning to and from kernel mode

Switch CR3

/arch/x86/entry/entry_64.S:

SYM_CODE_START(entry_SYSCALL_64)

...

SWITCH_TO_KERNEL_CR3 scratch_reg=%rsp

...

SWITCH_TO_USER_CR3_STACK scratch_reg=%rdi

/arch/x86/entry/calling.h:

.macro SWITCH_TO_KERNEL_CR3 scratch_reg:req

mov %cr3, \scratch_reg leggo contenuto registro

andq \$(~PTI_USER_PGTABLE_AND_PCID_MASK), \scratch_reg faccio flip 1 bit

mov \scratch_reg, %cr3 aggiorno cr3

.endm

Retpoline

- The retpoline (*return trampoline*) is a software construct to prevent branch-target injection forzo la cpu ad eseguire un percorso, anche sotto speculazione.
 - Indirect branches are isolated from speculative execution
 - It is built around return instructions
 - Speculative execution will spin indefinitely
- This is a compiler-assisted technique, which prevents several indirect branch instructions to be emitted in a binary, and replaces them with thunks
 - A thunk is a subroutine used to delay a calculation until the result is actually needed
- Examples of avoided x86 instructions:
 - `jmp *%rax`
 - `call *%rax`

se cambio rax, è facile avere side-effect.
Tali operazioni sono molto diffuse.

devo sostituirle

Retpoline thunks

quando faccio call set_up_target, salvo return address sul top dello stack (sarebbe capture_spec).

set_up_target faccio una mov, aggiornando il return address, che normalmente sarebbe l'istruzione successiva di capture_spec, mettendoci invece r11.

Quindi, se c'è speculazione, sarà cattura da capture_spec, che però non crea side effect.

```
jmp %r11      jmp retpoline_r11_trampoline;
call %r11     call retpoline_r11_trampoline;
```

```
retpoline_r11_trampoline:
```

```
    call set_up_target
```

```
capture_spec: ←
```

```
    pause
```

```
    jmp capture_spec
```

```
set_up_target:
```

```
    mov %r11, (%rsp) ←
```

```
    ret
```

Target known at compile time: will not trigger speculative target resolution

Speculation is captured here, which is a loop which will be eventually flushed

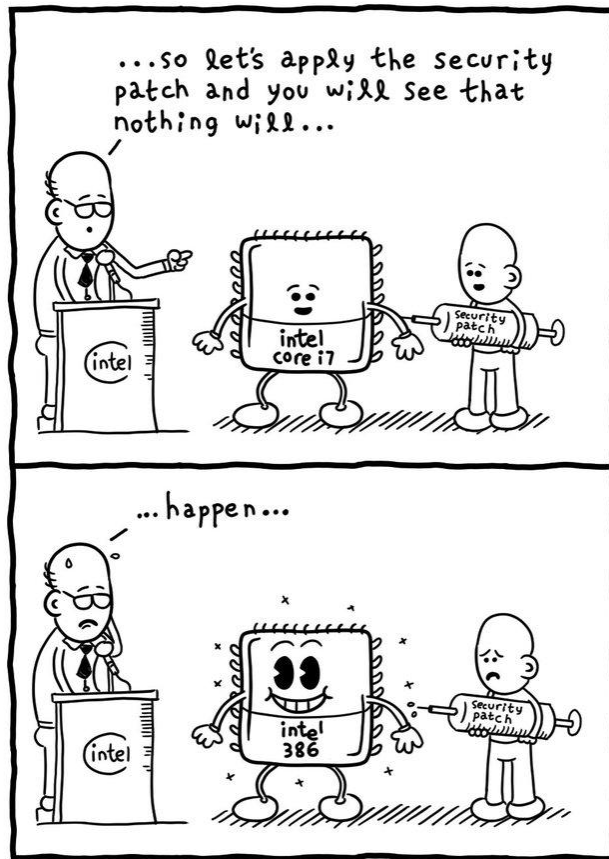
Modify the target of the return instruction on stack, leaving the Return Stack in the CPU unaffected

- In the code above, there is no point that can be controlled by an attacker
- There is one shared trampoline for each register in the final binary!

Prevent Branch Poisoning

- Additional facilities by the ISA
- Indirect Branch Restricted Speculation (IBRS)
 - The processor enters in a special IBRS mode
 - BPU is not affected by predictions outside of IBRS
- Single Thread Indirect Branch Prediction (STIBP)
 - Restricts branch prediction sharing between software executing on hyperthreads of the same core
- Indirect Branch Predictor Barrier (IBPB)
 - Prevents software running before setting the barrier from affecting branch prediction by software running after the barrier
 - This is done by flushing the state of the Branch Target Buffer

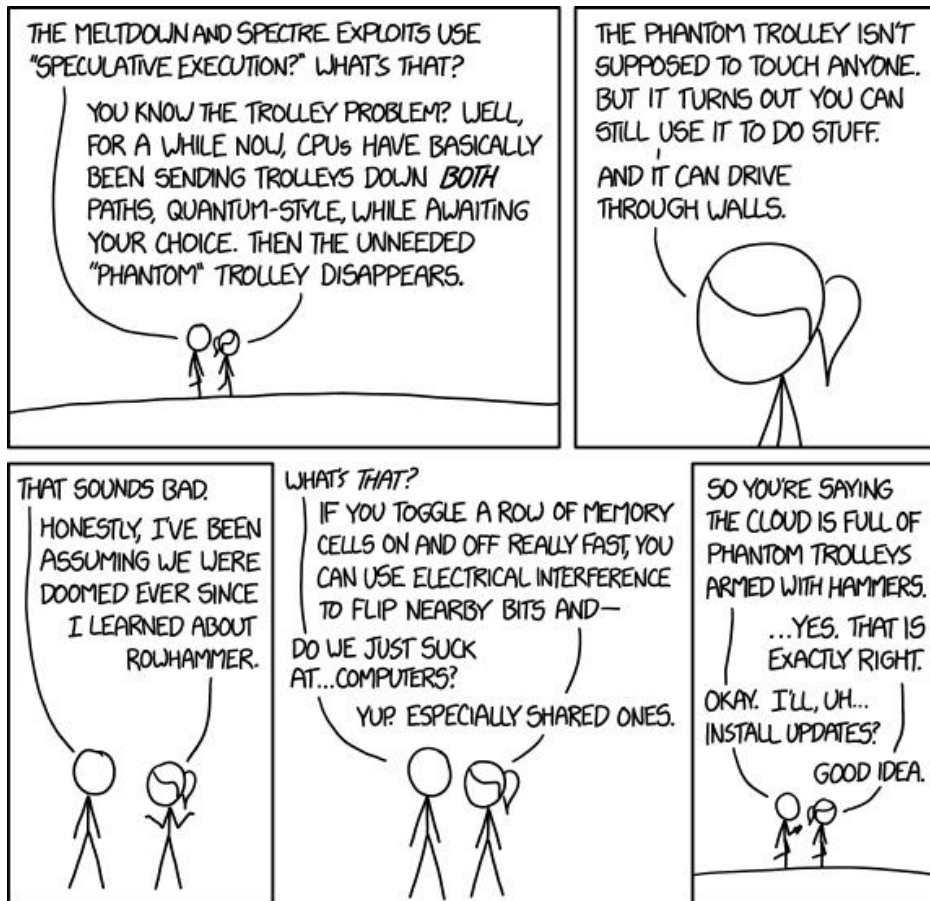
How to make your CPU more secure



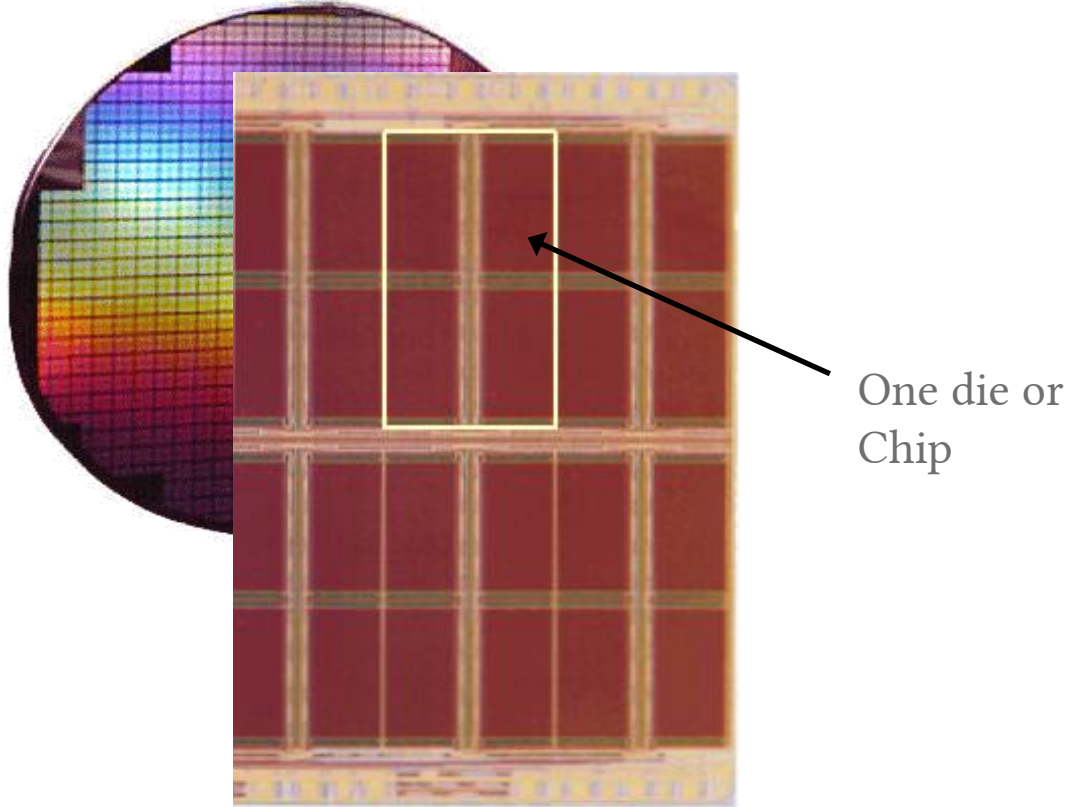
Daniel Stori {turnoff.us}

fino ad ora abbiamo operato in lettura,
vediamo in scrittura

Setting up a Writing Primitive

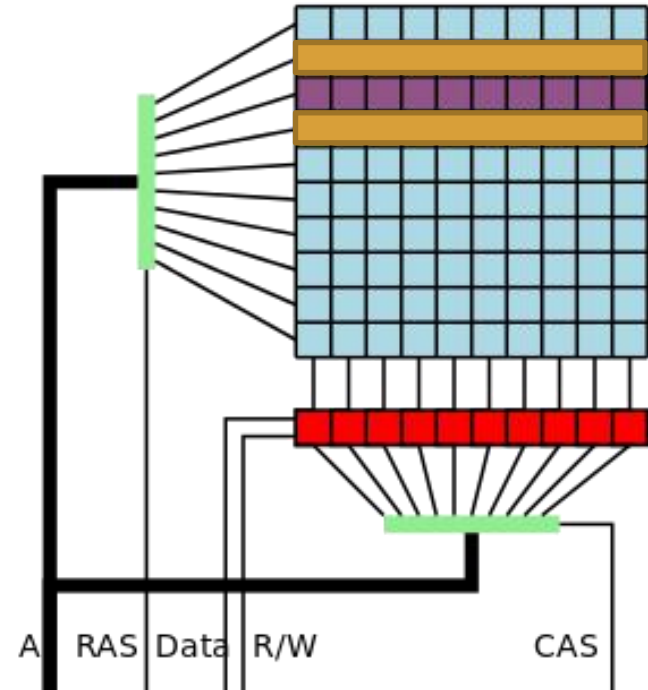


Anatomy of a Memory Chip



RowHammer

- Frequent row activations cause voltage fluctuations on the associated row selection lines
- These voltage fluctuations create an electromagnetic field (Lorentz force)
- This field can generate electric effects in nearby capacitors (adjacent, in most cases)
- Higher-than-natural discharge rates or charged capacitors
 - bit flips!



RowHammer POC

code1a:

```
mov (X), %eax // read from address X
mov (Y), %ebx // read from address Y
clflush (X)    // flush cache for address X
clflush (Y)    // flush cache for address Y
mfence
jmp code1a
```

ho creato campo elettromagnetico
su queste due linee, che carica alcuni bit

flusho subito, e poi riaccedo,
sto caricando fino a che non
ci sarà un flip.

Mitigations?

- Error correction codes in memory
 - might not detect multiple bit flips per memory word
- Reduce the 64 ms refresh interval
 - higher power consumption
 - increased processing overhead
- Detection: repeatedly accessing memory creates noise on cache
 - HPCs can be used to detect this noise as in side-channel attacks
- Pseudo Target Row Refresh (pTRR)
 - track row activation and refresh victim rows
 - implemented in the memory controller
 - “security through obscurity”

lavora con oscurità, alla fine lavora come un trigger.
Se vede qualcosa di sospetto, rinfresha le vittime.

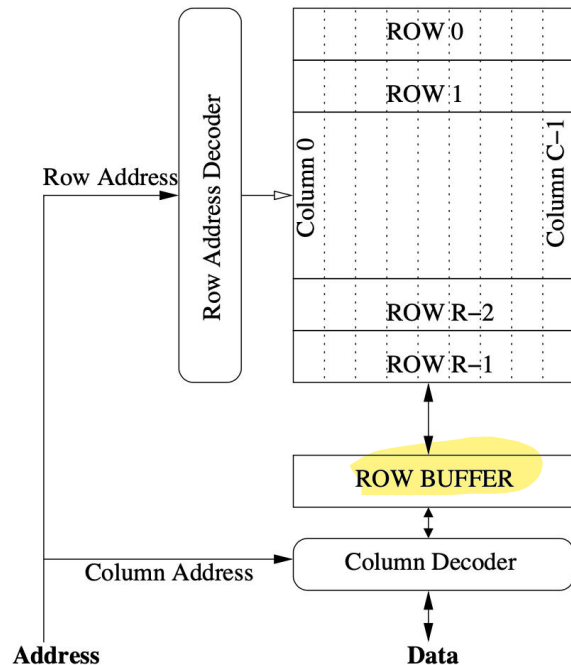
Memory Performance Attacks

Denial of Memory Service in Multi-
Core Systems

DRAM Memory Scheduler

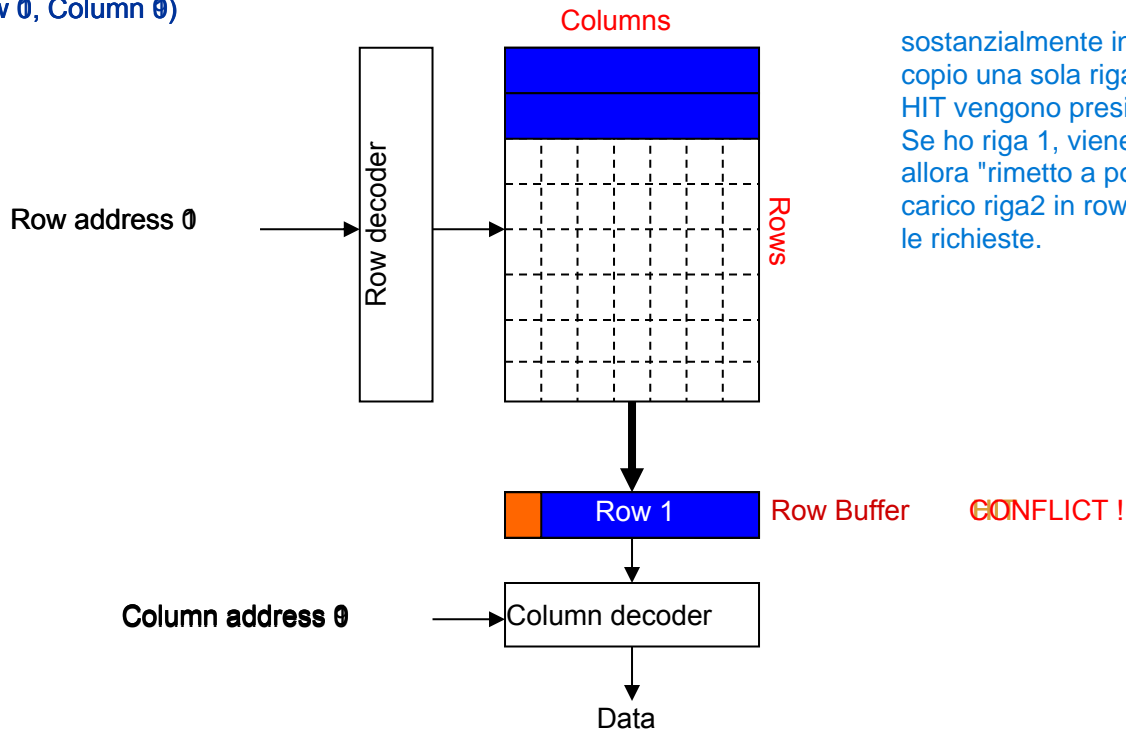
- DRAM schedulers have been designed for single core chips!
- To improve speed, each memory bank has a *single* row buffer
- Memory access is performed only through this buffer
- Request hits or misses the current content of the row buffer

gli accessi in memoria passato per row buffer



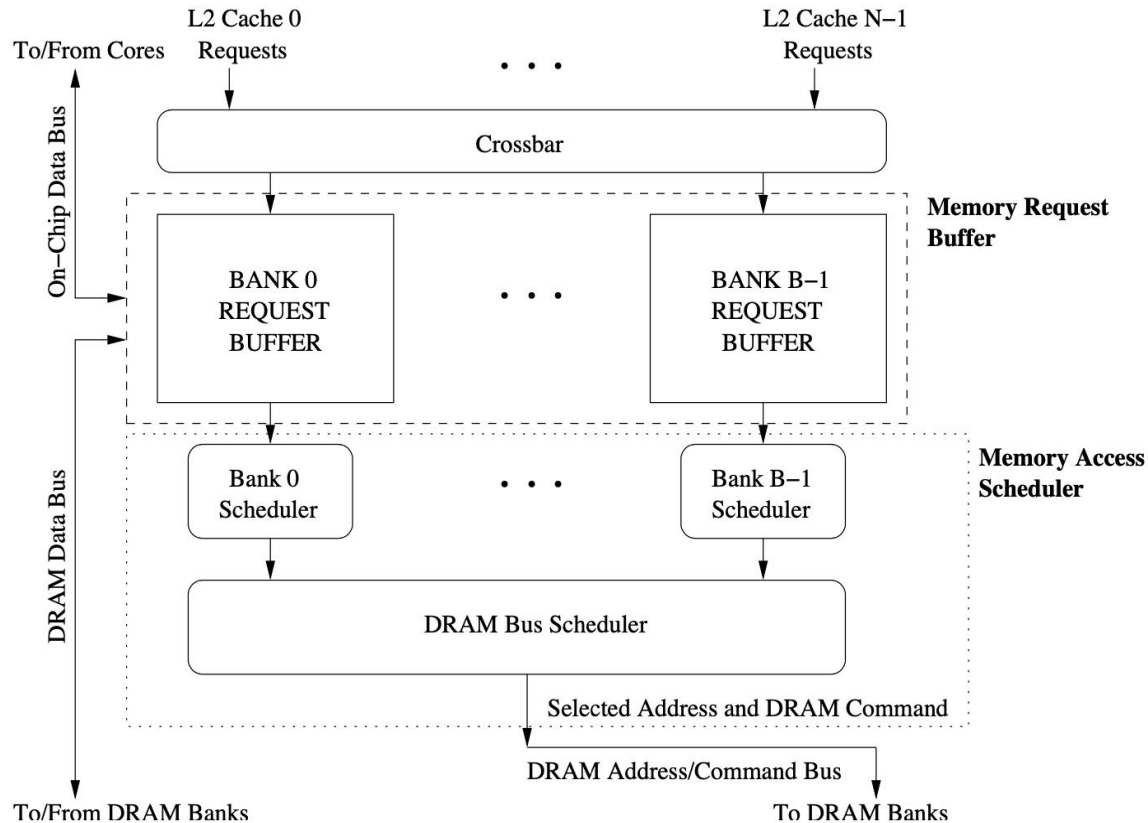
DRAM Bank Operations

Access Address
(Row 0, Column 0)



sostanzialmente in row buffer
copio una sola riga, gli accessi
HIT vengono presi da row buffer.
Se ho riga 1, viene chiesta riga 2,
allora "rimetto a posto" riga 1,
carico riga2 in row buffer, e servirò con lui
le richieste.

Multi-bank DRAM Memory Systems



DRAM Memory Access Scheduling: FR-FCFS

Bank scheduler:


- Row-hit-first
 - Low row-buffer locality threads have low priority
- Oldest-within-bank-first
 - Prioritizes threads that generate requests fast

Bus scheduler:


- Oldest-across-banks-first (among all requests proposed by individual bank scheduler)

An Aggressive, High Row-Buffer Locality Thread

richieste
da thread diversi



Request Buffer



un problema è che io
potrei avere in row buffer riga 0,
poi c'è richiesta riga1, e poi
richiesta riga0.
avrebbe senso far passare
riga0 perchè già disponibile,
però così, ripetendolo, potrei
evitare al thread che chiede
riga1 di usare la ram.

NON ESISTE FIX che
non modifica
totalmente l'architettura

It hogs DRAM!