

29/03/2022

Performance Modeling of Computer Systems and Networks

Prof. Vittoria de Nitto Personè

Analytical results
KP further results

Università degli studi di Roma Tor Vergata
Department of Civil Engineering and Computer Science Engineering

Copyright © Vittoria de Nitto Personè, 2021
<https://creativecommons.org/licenses/by-nc-nd/4.0/>



1

P 31 slide,

Analytical models
M/G/1

The waiting time and the remaining service time

new job che arriva

N_Q

t_c

S_{rem}

combinazione

$T_Q = t_c \oplus S_{rem}$

T_Q in funzione di:

$E(S_{rem}) = \frac{\lambda}{2} E(S^2)$ distribuzione

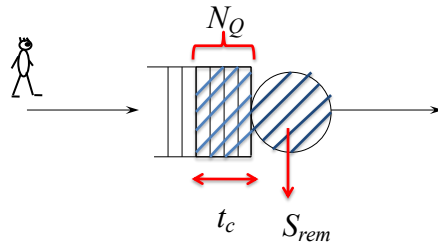
media del quadrato del servizio

Prof. Vittoria de Nitto Personè

2

2

The waiting time and the remaining service time



$$E(S_{rem}) = \frac{\lambda}{2} E(S^2)$$

$$\rho = \lambda E(S)$$

nel caso exponential

ho:

$$E(S^2) = 2E(S)^2 \longrightarrow E(S_{rem}) = \frac{\lambda}{2} E(S)^2$$

$$E(S_{rem}) = \rho E(S)$$

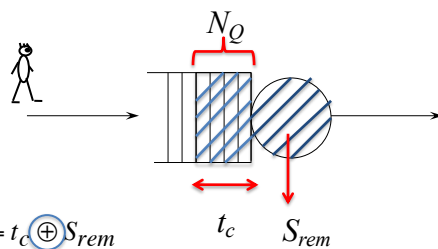
prob che ci sia qualcuno, e' esponenziale.

Prof. Vittoria de Nitto Personè

3

3

The waiting time and the remaining service time



$$T_Q = t_c \oplus S_{rem}$$

dalla
tabella

$$E(T_Q) = \frac{\rho E(S)}{1 - \rho} = \frac{E(S_{rem})}{1 - \rho} = \frac{1}{1 - \rho} E(S_{rem})$$

exponential

$$E(S_{rem}) = \rho E(S)$$

↳ fattore, e' un numero puro, non tempo!

Prof. Vittoria de Nitto Personè

4

4

caso generale

$$\begin{aligned}
 E(T_Q) &= \frac{\rho}{1-\rho} \frac{C^2+1}{2} E(S) = \\
 &= \frac{\rho}{2(1-\rho)} \left[\frac{\sigma^2(S)}{E(S)^2} + 1 \right] E(S) = \\
 &= \frac{\rho}{2(1-\rho)} \left[\frac{E(S^2) - E(S)^2}{E(S)^2} + 1 \right] E(S) = \\
 &= \frac{\lambda E(S)}{2(1-\rho)} \left[\frac{E(S^2)}{E(S)^2} - 1 + 1 \right] E(S) = \\
 &= \frac{\lambda}{2(1-\rho)} \left[\frac{E(S^2)}{E(S)^2} \right] E(S)^2 = \frac{\frac{\lambda}{2} E(S^2)}{1-\rho}
 \end{aligned}$$

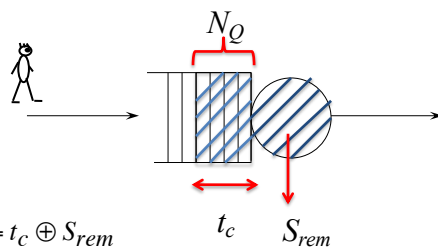
Prof. Vittoria de Nitto Personè

5

5

Analytical models
M/G/1

The waiting time and the remaining service time



$$\frac{1}{1-\rho}$$

represents the mean time to complete the jobs in the queue at the arrival instant

• Rappresenta il fattore legato a t_c MA NON è un tempo

$$\rho \cdot E(S) = E(S_{rem})$$

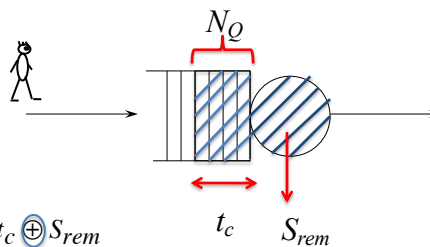
is the mean time to complete the job in service at the arrival instant

Prof. Vittoria de Nitto Personè

6

6

The waiting time and the remaining service time



$T_Q = t_c \oplus S_{rem}$

$\frac{1}{1-\rho}$

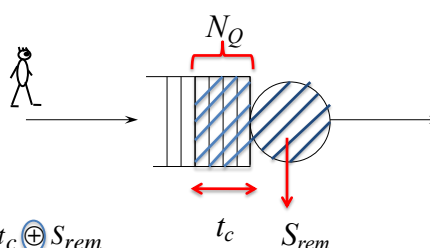
$P \cdot E(S) = E(S_{rem})$

the two terms are compounded through the multiplication:
if one is 0, the result is 0
if one is 1, it does not give contribution

Prof. Vittoria de Nitto Personè

7

The waiting time and the remaining service time



$T_Q = t_c \oplus S_{rem}$

$\frac{1}{1-\rho}$

$E(S_{rem}) = 0$ \rightarrow Non c'è coda, new job subito servito \rightarrow The queue cannot be \rightarrow The incoming job does not wait, but it is immediately served

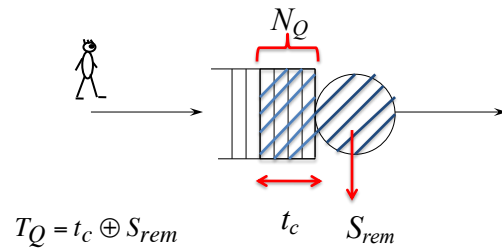
$E(S_{rem})$ $\frac{1}{1-\rho} = 1$ \rightarrow The component does not give contribution, the incoming job finds just the job in service with remaining service so small that it does not wait

coda con piccola che il new job trova solo job in esecuzione che ha "quasi finito"

Prof. Vittoria de Nitto Personè

8

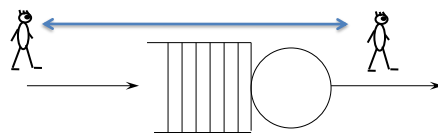
The waiting time and the remaining service time



$$E(T_Q) = \frac{E(S_{rem})}{1-\rho} = \frac{\frac{\lambda}{2} E(S^2)}{1-\rho}$$

tempo servizio rimanente
M/G/1
colonna

The response time



$$E(T_S) = E(T_Q) + E(S) = \frac{\frac{\lambda}{2} E(S^2)}{1-\rho} + E(S)$$

M/G/1

$$E(T_S) = \frac{\rho E(S)}{1-\rho} + E(S) = \frac{E(S)}{1-\rho}$$

M/M/1

P 32 slide

Analytical models
scheduling

Non-preemptive abstract scheduling

(ho risultati più forti di quelli della KP)

Def. 1
A policy is *preemptive* if a job may be stopped part way through its execution and then resumed at a later point in time from the same point where it was stopped. A policy is *non-preemptive* if jobs are always run-to-completion.

Preemptive se fermo Job e dopo lo riprendo da dove si era fermato

Def. 2
A *work-conserving* scheduling policy is one which always performs work on some job when there is a job in the system.

Theorem 1 (Conway, Maxwell, Miller¹).
All non-preemptive service orders that do not make use of job sizes have the same distribution on the number of jobs in the system.

NON solo le medie!

$E(N_S) \quad E(T_S) \quad E(N_Q) \quad E(T_Q)$

¹Richard Conway, William Maxwell, and Louis Miller, Theory of Scheduling Addison-Wesley Publishing Company, Inc., 1967. Chapter 8

Prof. Vittoria de Nitto Personè

11

11

Analytical models
scheduling

Non-preemptive abstract scheduling

$$E(T_Q) = \frac{\frac{\lambda}{2} E(S^2)}{1 - \rho}$$

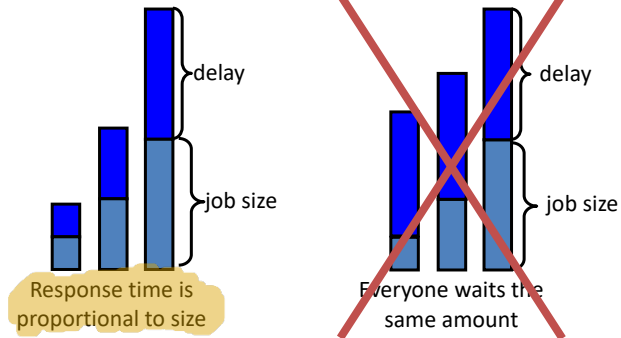
which is very high when $E(S^2)$ is high (anche se ρ piccolo!)

Prof. Vittoria de Nitto Personè

12

12

WHAT IS FAIR?



13

Let us consider the mean time in system for a job of size x

dipende da taglia 'x'

$$E(T_S(x)) = E(x + T_Q(x)) = x + E(T_Q) = x + \frac{\frac{\lambda}{2} E(S^2)}{1 - \rho}$$

non dipende da size (in $\frac{\lambda}{2} \frac{E(S^2)}{1 - \rho}$ non c'è x)non è $E(S)$, ma ciò che chiedo!

14

Analytical models
scheduling

WHAT IS FAIR?

Response time is proportional to size

Everyone waits the same amount

Define the **slowdown** for a job of size x as
condizionato

$$E(sd(x)) = \frac{E(T_S(x))}{x}$$

↑
Misura fairness

15

Analytical models
scheduling

Slowdown for jobs of size x

Def. The mean slowdown for jobs of size x is the observed mean response time in respect of their size, that is

$$E(sd(x)) = \frac{E(T_S(x))}{x} = \frac{[\text{tempo risposta size}]}{[\text{quanto chiede}]}$$

metto λP e divido per x

$$1 < E(sd(x)) = 1 + \frac{\lambda E(S^2)}{x(1-\rho)}$$

Se $\rho = 1$ non attendo nulla! (chiedo e spendo x)

inverso a x

Sono in sist non-preemptive, (es FIFO), con T_a e $E(S)$, normalmente ha senso pensare che job più piccoli siano svantaggiati nell'attendere big jobs.

Note that small jobs have a higher expected slowdown than do big jobs.

es: se P_1 e P_2 aspettano $T_q = 30s$, $x_1 = 5s$ e $x_2 = 35s$, ovviamente P_1 è svantaggiato (Tipo andare alle poste, zone di fila) per operazione veloce

come succede a quelli grandi? non pochi ma ci sono, si portano davanti i job piccoli, i grandi vanno in starvation?

Prof. Vittoria de Nitto Personè

16

Analytical models
scheduling

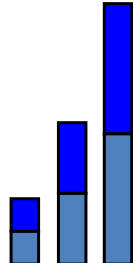
Slowdown vs Response time

Response Time tends to be representative of the performance of just a few jobs — the bigger ones
tends to emphasize the performance of the really big jobs, since they count the most in the mean, since their response time tends to be the greatest (emphasized for heavy-tail distr.)

Slowdown tends to be representative of the performance of most jobs — because it is dominated by the performance of the large number of small jobs.

we would like to make $E(T_S(x))$ smaller for small x

(se alle poste op veloce, vorrei aspettare poco!)



17

Prof. Vittoria de Nitto Personè

17

33 slide, 508 Performance

Analytical models
scheduling

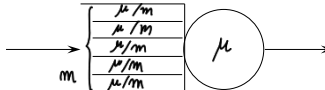
Processor sharing

we would like to make $E(T_S)$ smaller for small x

How do we do this if we DON'T know job sizes?

two reasons historically why CPU scheduling is (approximately) processor-sharing

1. in a multi-resource system (including a CPU, disk, memory, etc.) it is useful to have many jobs running simultaneously (rather than just one job at a time) because jobs requiring different resources can be overlapped to increase throughput. (se \forall job do quanto di tempo, little job finiscono, big jobs no.)
2. PS is a good way to get small jobs out fast, given that we don't know the size of the jobs.



18

Prof. Vittoria de Nitto Personè

18

Analytical models
PS scheduling

Processor sharing tempi risposta medi

should be better than FIFO with respect to $E(T_S)$, because PS gets small jobs out faster, and PS should be a lot better than FIFO with respect to $E(sd)$!

Process. Sharing *ammortizza variabilità* esponenziale FIFO

distribuzione $Pr\{N_S = n\}^{M/G/1/PS} = \rho^n (1 - \rho) = Pr\{N_S = n\}^{M/M/1/FIFO}$

popolazione globale $E(N_S)^{M/G/1/PS} = \frac{\rho}{1 - \rho} = E(N_S)^{M/M/1/FIFO}$

tempo risposta $E(T_S)^{M/G/1/PS} = \frac{E(S)}{1 - \rho} = E(T_S)^{M/M/1/FIFO}$

PS is better than FIFO when $C^2 > 1$

the M/G/1/PS queue is insensitive to the variability of the service time distribution, G (variabilità non compare nelle formule)

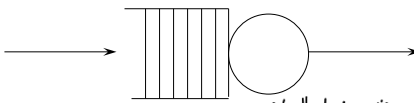
Prof. Vittoria de Nitto Personè

19

Quanto appena detto è vero sempre, o esistono casi particolari?

2 arrivi simultanei che richiedono 1 s di servizio

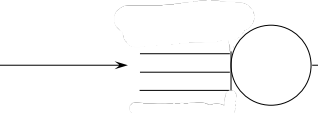
Fifo



→ un job attende 1 e viene servito per un altro, l'altro subito servito

$E(T_S) = (1+2)/2 = 1.5 \text{ s}$ **$E(T_S) ?$**

PS



Adesso J_1 e J_2 accedono insieme, ma capacità server è la metà: entrambi hanno $\frac{1}{2}$ capacità a testa per secondo.

$E(T_S) = 2 \text{ s}$ $E(T_S) = 2 (0.5 \cdot \frac{1}{s} + 0.5 \cdot \frac{1}{s}) = 2$

PS può essere peggio su alcune sequenze di job!

Prof. Vittoria de Nitto Personè

20

Processor sharing

$$E(T_S(x))^{M/G/1/PS} = \frac{x}{1-\rho}$$

$$E(sd(x))^{M/G/1/PS} = \frac{1}{1-\rho}$$

all jobs have same slowdown: PS as "FAIR" scheduling, *non c'è la size*, *PS è con prelazione*, *equo, rallenta tutti ugualmente*

$$F_i F_o: E(sd(x))^{M/G/1/abstract} = 1 + \frac{\lambda E(S^2)}{x(1-\rho)}$$

21

all the preemptive non-size-based scheduling policies produce the same mean slowdown for all job sizes

$$E(sd(x))^{M/G/1/preempt-non-size-based} = \frac{1}{1-\rho}$$

We would like to get lower slowdowns for the smaller jobs

But how can we give preference to the smaller jobs if we don't know job size?

we do know a job's age (CPU used so far), and age is an indication of remaining CPU demand

If the job size distribution has DFR (e.g. Pareto distribution) then the greater the job's age, the greater its expected remaining demand

→ give preference to jobs with low age (younger jobs) and this will have the effect of giving preference to jobs which we expect to be small

(heavy tail: leggere par. 20.5!)

Performance's Book

Prof. Vittoria de Nitto Personè

In pratica: se job nuovo scommetto rimanga poco (age bassa). Se age alta scommetto vi rimanga tanto (Decreasing Failure Rate)

22

35 slide, 516 Par f

Analytical models
scheduling

Scheduling in Unix:

Foreground-Background scheduling (= Multi Level Process Sharing)

1 gestione PS

gruppo prioritario

gruppo non prioritario

2

age a

- gruppo 1, dopo 'age a', passa a gruppo 2.
- chi chiede poco se ne va prima.
- gruppo 2 servito se gruppo 1 vuoto.

L'ordine è basato su ciò che chiedono e sul tempo rimanente.

Coda 1 principale, se vuota passa alla 2. Un job passa da coda 1 a coda 2 superando un **age threshold**. (se job più vecchio di tale parametro, va in coda 2). Appena arriva job in coda 1, servo lui dandogli precedenza. Mi sono sull'**age** (NOTA), poiché non conosco la SIZE.

Prof. Vittoria de Nitto Personè

23