



# Performance Modeling of Computer Systems and Networks

Prof. Vittoria de Nitto Personè

Discrete-Event Simulation  
examples

Università degli studi di Roma Tor Vergata  
Department of Civil Engineering and Computer Science Engineering

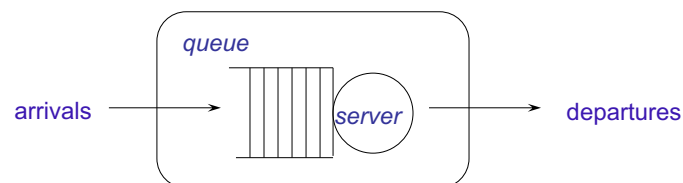
Copyright © Vittoria de Nitto Personè, 2021  
<https://creativecommons.org/licenses/by-nc-nd/4.0/>



1

Discrete-Event Simulation  
case study ssq

## Single Server Queue



Arrival times:  $a_i$

~~15 47 71 111 123 152 166 226 310 320~~

**Pseudo-random generators**

Service times:  $s_i$

~~43 36 34 30 38 40 31 29 36 30~~

Prof. Vittoria de Nitto Personè

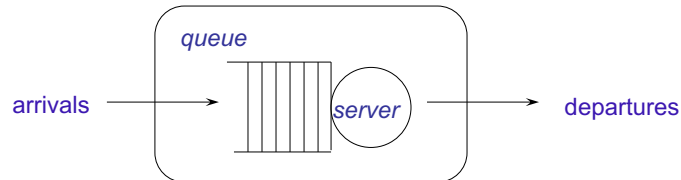
2

2

P. 97 discrete

Discrete-Event Simulation

## Single Server Queue



- assume **service times** are between 1.0 and 2.0 minutes

- The distribution within this range is unknown
- Without further knowledge, we assume no time is more likely than any other

media varianza  
 → Uniform(1.0, 2.0)

Prof. Vittoria de Nitto Personè

3

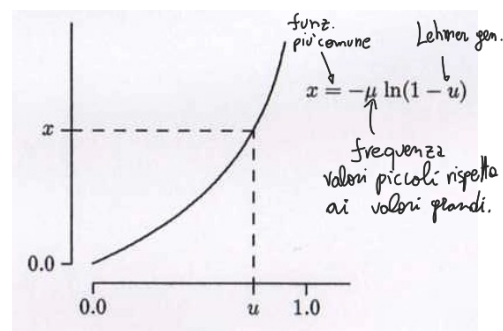
3

Discrete-Event Simulation

## Exponential distribution

- In general, it is unreasonable to assume that all possible values are equally likely (uniform ha il valore stessa prob, nella realtà non sempre è realistico)
- Frequently, small values are more likely than large values
- We need a non-linear transformation that maps  $0.0 \rightarrow 1.0$  to  $0.0 \rightarrow \infty$

this is the most frequently used function  
 $\mu > 0$  is a parameter that "control" the frequency of large values in respect of the small ones



Prof. Vittoria de Nitto Personè

4

4

- the transformation is monotone increasing, one-to-one

$$\begin{aligned}
 0 < u < 1 &\Leftrightarrow 0 < (1 - u) < 1 \\
 &\Leftrightarrow -\infty < \ln(1 - u) < 0 \\
 &\Leftrightarrow 0 < -\mu \ln(1 - u) < \infty \\
 &\Leftrightarrow 0 < x < \infty
 \end{aligned}$$

```

double Exponential(double μ)    /* use μ > 0.0 */
{
    return (-μ * log(1.0 - Random()));
}

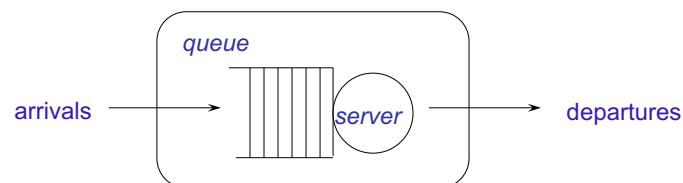
```

$\hookrightarrow Q_n$

- the parameter  $\mu$  specifies the sample mean

abbiamo visto i tempi di interarrivo, usando TRACCE,  
ora facciamo senza!

### Single Server Queue



Arrival times:  $a_i$

- use the exponential function for the interarrival times

$$a_i = a_{i-1} + \text{Exponential}(\mu); i = 1, 2, 3, \dots, n$$

Service times:  $s_i$

*Uniform(1.0, 2.0)*

- program `ssq2` is an extension of `ssq1` <sup>usava le tracce</sup>
  - arrival times are drawn from `Exponential(2.0)`
  - service times are drawn from `Uniform(1.0, 2.0)`

## trace-driven simulation

```
#include <stdio.h>

#define FILENAME "ssq1.dat" /* input data file */
#define START 0.0

double GetArrival(FILE *fp) /* read an arrival time */
{
    double a;
    fscanf(fp, "%lf", &a);
    return (a);}

double GetService(FILE *fp) /* read a service time */
{
    double s;
    fscanf(fp, "%lf\n", &s);
    return (s);}
```

## ssq2.c distribution-driven simulation

```

#include <stdio.h>
#include <math.h>
#include "rng.h"
#define LAST      10000L    /* number of jobs processed */
#define START     0.0

double Exponential(double m)          /* -----*
{return (-m * log(1.0 - Random())); }      m > 0.0
                                           -----*/

double Uniform(double a, double b)    /* -----*
{return (a + (b - a) * Random()); }      a < b
                                           * -----*/

double GetArrival(void)
{static double arrival = START;
 arrival += Exponential(2.0); ← genero tempo interarrivo
 return (arrival);}

double GetService(void)
{return (Uniform(1.0, 2.0));}

```

Prof. Vittoria de Nitto Personè

9

9

## Discrete-Event Simulation case study ssq

- the program generates all first-order statistics

$$\bar{r}, \bar{w}, \bar{d}, \bar{s}, \bar{l}, \bar{q}, \bar{x}$$

Prof. Vittoria de Nitto Personè

10

10

trace-driven simulation = distribution-driven simulation

```
int main(void)
{ FILE *fp; /* input data file */
  long index = 0; /* job index */
  double arrival = START; /* arrival time */
  double delay; /* delay in queue */
  double service; /* service time */
  double wait; /* delay + service */
  double departure = START; /* departure time */
  struct { /* sum of ... */
    double delay; /* delay times */
    double wait; /* wait times */
    double service; /* service times */
    double interarrival; /* interarrival times */
  } sum = {0.0, 0.0, 0.0};
```

Prof. Vittoria de Nitto Personè

11

11

trace-driven simulation

```
fp = fopen(FILENAME, "r");
if (fp == NULL) {
  fprintf(stderr, "Cannot open input file %s\n", FILENAME);
  return (1);
}
while (!feof(fp)) {
```

distribution-driven simulation

```
PutSeed(123456789);
while (index < LAST) {
  ultima job che vogliamo simulare
```

Prof. Vittoria de Nitto Personè

12

12

```

                                trace-driven: fp
while (index < LAST) {
    index++;
    arrival      = GetArrival();
    if (arrival < departure)
        delay = departure - arrival; /* delay in queue */
    else delay = 0.0;                /* no delay */
    service = GetService();
    wait = delay + service;
    departure = arrival + wait; /* time of departure */
    sum.delay += delay;
    sum.wait += wait;
    sum.service += service;
}
sum.interarrival = arrival - START;

```

Prof. Vittoria de Nitto Personè

13

13

```

printf("\nfor %ld jobs\n", index);
printf("  average interarrival time = %6.2f\n", sum.interarrival / index);
printf("  average wait ..... = %6.2f\n", sum.wait / index);
printf("  average delay ..... = %6.2f\n", sum.delay / index);
printf("  average service time .... = %6.2f\n", sum.service / index);
printf("  average # in the node ... = %6.2f\n", sum.wait / departure);
printf("  average # in the queue .. = %6.2f\n", sum.delay / departure);
printf("  utilization ..... = %6.2f\n", sum.service / departure);
return (0);

```

Prof. Vittoria de Nitto Personè

14

14

# ESEMPIO SINGLE SERVER QUEUE



- coda ∞
- server singolo
- $\lambda = 0.5$  j/s esponenziale
- $E(S) = 1.5$  s, Unif(a=1, b=2)

## SVOLGIMENTO

Con Little  $E(T_q) = \frac{\rho}{1-\rho} \left( \frac{C^2+1}{2} \right) \cdot E(S)$ ; media:  $\bar{x} = E(\bar{S}) = \frac{a+b}{2} = 1.5$  s; VARianza  $\sigma^2(x) = E(S^2) = \frac{(b-a)^2}{12} = \frac{1}{12}$

• calcolo  $\rho = \lambda E(S) = 0.75$  e  $E(T_q) = 2.33$  s (se usi  $C^2 = \frac{1}{12} \cdot \frac{1}{(1.5)^2} = 0.037037$ ) tempo attesa medio in coda.  $E(T_s) = E(T_q) + E(S) = 2.33 + 1.5 = 3.83$  s.

• Uso LITTLE per i risultati:  $E(N_q) = 1.1667$ ;  $E(N_s) = 1.9167 = 1.1667 + \rho$

NB:  $\frac{\rho}{1-\rho} \left( \frac{C^2+1}{2} \right) E(S) = \frac{\lambda}{\lambda - \rho} \left( \frac{C^2+1}{2} \right) E(S) = \frac{C^2 \sigma^2(S)}{E(S)^2} \frac{\lambda}{2\rho(1-\rho)} [E(S)^2 + 1] E(S) = \frac{\lambda \cdot E(S^2)}{2\rho(1-\rho)}$  DIVERSI!

- $E(N_s)$  ci dice che ha in media quasi 2 job nel server
- server occupato per il 75%
- un job che chiede  $E(S) = 1.5$  attende  $E(T_s) = 3.8 \approx$  doppio!

Discrete-Event Simulation  
case study ssq

Example 1

- The "theoretical" averages using *Exponential(2.0)* (rate 0.5 j/s) arrivals and *Uniform(1.0, 2.0)* (rate 0.67) service times are

$\bar{r}$	$\bar{w}$	$\bar{d}$	$\bar{s}$	$\bar{l}$	$\bar{q}$	$\bar{x}$
2.00	3.83	2.33	1.50	1.92	1.17	0.75

exact analytical results,  
No simulation!

$\bar{r}$ : server risponde ogni  
 $\bar{w}$ : server  
 $\bar{d}$ : Job chiede ogni 1.5  
 $\bar{s}$ : in media, 2 Job  
 $\bar{l}$ : in media, 2 Job  
 $\bar{q}$ : in media, 2 Job  
 $\bar{x}$ : in media, 2 Job

- Although the server is busy 75% of the time, on average there are approximately 2 jobs in the service node
- A job can expect to spend more time in the queue than in service
- To achieve these averages, many jobs must pass through node

(simulando posso vedere il comportamento per molti job, osservando quando le statistiche tendono a questi valori asintotici)

Prof. Vittoria de Nitto Personè

15

Discrete-Event Simulation  
case study ssq

- The accumulated average wait was printed every 20 job

Number of jobs, n

Initial seed  
 ○ - 12345  
 ◇ - 54321  
 \* - 2121212  
 (sempi diversi)

- The convergence of  $w$  is slow, erratic, and dependent on the initial seed

Prof. Vittoria de Nitto Personè

16



(quando diventa  
stazionario)

- the program can be used to study the **steady-state** behavior
  - Will the statistics converge independent of the initial seed?
  - How many jobs does it take to achieve steady-state behavior?

(e il grafico di prima, con 5/5 simulazioni diverse)

- the program can be used to study the **transient** behavior (fase iniziale curve)
  - Fix the number of jobs processed and replicate the program with the initial state fixed (2000 Job nel grafico, tutte le simulazioni partono da stessa config.)
  - Each replication uses a different initial rng seed (cambio semi)

Prof. Vittoria de Nitto Personè

17

17

## Steady-state analysis

	$\bar{r}$	$\bar{w}$	$\bar{d}$	$\bar{s}$	$\bar{l}$	$\bar{q}$	$\bar{x}$	
theoretical	2.00	3.83	2.33	1.50	1.92	1.17	0.75	
n=10	2.85	1.74	0.39	1.35	0.59	0.13	0.45	
n=100	2.06	3.16	1.67	1.48	1.50	0.80	0.71	
n=1000	2.03	3.44	1.94	1.50	1.69	0.96	0.74	
n=10000	2.02	3.86	2.36	1.50	1.91	1.17	0.74	
n=100000	2.00	3.85	2.35	1.50	1.92	1.17	0.75	
n=1000000	2.00	3.81	2.31	1.50	1.90	1.15	0.75	seed=123456789
n=1000000	2.00	3.84	2.34	1.50	1.92	1.17	0.75	seed=1

Initial transiente										
POCH, comportamento molto variabile	{	n=10	2.13	2.36	0.75	1.62	1.02	0.32	0.69	seed=1
		n=10	1.48	2.37	0.87	1.50	1.24	0.46	0.79	seed=987654321
		n=10	1.49	1.89	0.49	1.40	1.12	0.29	0.83	seed=21212121

## Transient analysis

Prof. Vittoria de Nitto Personè

18

18

## Simulating an initial steady state

Inizio stazionarietà

$\bar{d} = 2.33$  (delay, attesa in coda)

departure=3, il sistema parte in stato stabile (departure=3), e NON da VUOTO

$a_1 = a_0 + \overset{\text{media}}{\text{expo}(2)} = 0 + 0.8 = 0.8$

$0.8 < 3$  (c'è attesa)

$d_1 = 3 - 0.8 = 2.2 \approx$  attesa media in coda

$s_1 = \overset{\text{servizio}}{\text{Uniform}(1,2)} = 1.3$

$w_1 = 2.2 + 1.3 = 3.5$  (tempo risposta: attesa + servizio)

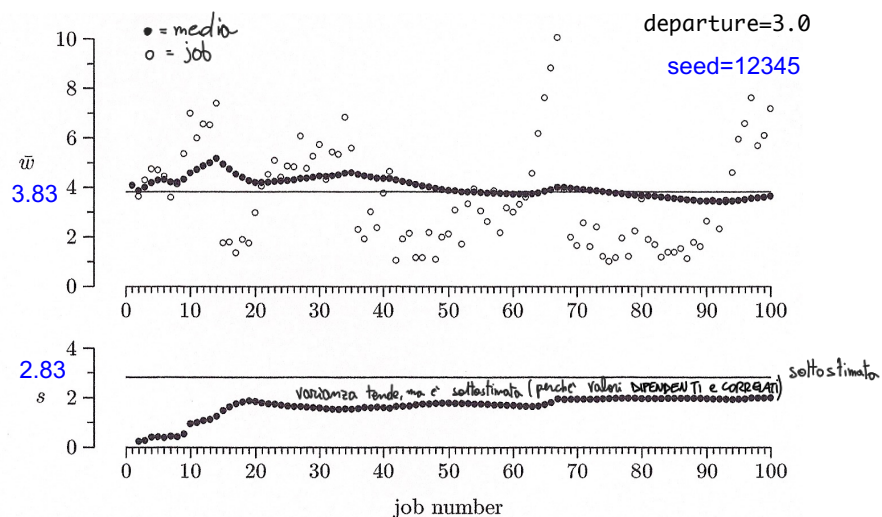
$c_1 = 0.8 + 3.5 = 4.3$  (quando job esce dal centro)

Prof. Vittoria de Nitto Personè

19

19

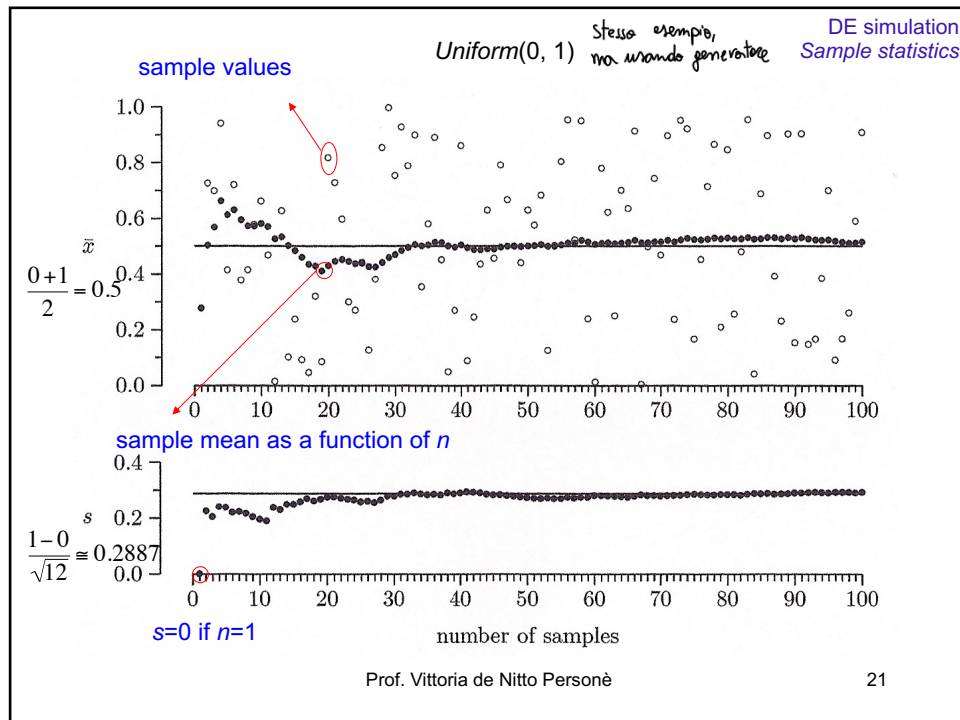
DE simulation  
Sample statistics



Prof. Vittoria de Nitto Personè

20

20



21

DE simulation  
Sample statistics

## Serial correlation

- **Independence:** each  $x_i$  value does not depend on any other point
- Time-sequenced DES output is typically not independent
- E.g.: wait times of consecutive jobs have positive *serial correlation*
- Example: Consider output from ssq2
  - Exponential(2) interarrivals, Uniform(1,2) service
- wait times  $w_1, w_2, \dots, w_{100}$ , have high positive serial correlation
  - The correlation produces a bias in the standard deviation

Prof. Vittoria de Nitto Personè

22

22

## Example 2

- assume that jobs arrive at random with a steady-state arrival rate of 0.5 jobs per minute
- assume that Job service times are "composite" with two components:
  - the number of service tasks is  $1 + \text{Geometric}(0.9)$   $\rightarrow$  median =  $\frac{1}{0.9}$
  - the time (in minutes) per task is  $\text{Uniform}(0.1, 0.2)$   $\rightarrow$  median =  $\frac{0.1+0.2}{2} = 0.15$

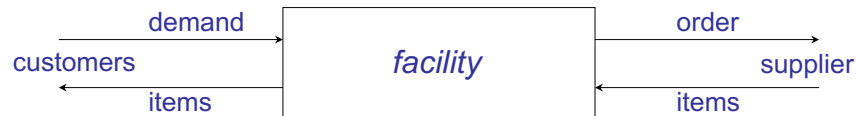
```
double GetService(void)
{
    long k;
    double sum = 0.0;
    long tasks = 1 + Geometric(0.9); ←
    for (k = 0; k < tasks; k++)
        sum += Uniform(0.1, 0.2);
    return (sum);
}
```

- The theoretical steady-state statistics for this model are

$\bar{r}$	$\bar{w}$	$\bar{d}$	$\bar{s}$	$\bar{l}$	$\bar{q}$	$\bar{x}$	
2.00	5.77	4.27	1.50	2.89	2.14	0.75	exact analytical results, No simulation!
" "	(3.83	2.33	" "	1.92	1.17)	" "	prodotti da ssq 2

- The arrival rate, service rate, and utilization are identical to the previous case (example 1)
- The other four statistics are significantly larger
- performance measures are sensitive to the choice of service time distribution

## A simple inventory system (revisione)



- Distributes items from current inventory to customers
- Customer demand is discrete
- Simple: one type of item
- Inventory review is periodic
- Items are ordered, if necessary, only at review times
- $(s, S)$  are the min,max inventory levels,  $0 \leq s < S$

fisso  
variabile per vedere  
i benefit

Prof. Vittoria de Nitto Personè

25

25

P. 02

## Simply Inventory System

- Program sis2 has randomly generated demands using an *Equilikely(a, b)* random variate (non realistico) → passo alla geometria (anche lei non perfetta)
- Using random data, we can study transient and steady-state behaviors

#include &lt;stdio.h&gt;

sis1.c

#define FILENAME "sis1.dat"

#define MINIMUM 20

#define MAXIMUM 80

#define STOP 100

#define sqr(x) ((x) \* (x))

~~long GetDemand(FILE \*fp)~~~~{~~~~long d;~~~~fscanf(fp, "%ld\n", &d);~~~~return (d);}~~

Prof. Vittoria de Nitto Personè

26

26

sis2.c

distribution driven

```

#include <stdio.h>
#include "rng.h"

#define MINIMUM 20
#define MAXIMUM 80
#define STOP 100 /* 100 weeks = about 2 years*/
#define sqr(x) ((x) * (x))

long Equilikely(long a, long b)
{ return (a + (long) ((b - a + 1) * Random()));}

long GetDemand(void)
{
    return (Equilikely(10, 50)); }

```

27

```

int main(void)
{
    long index      = 0;
    long inventory = MAXIMUM;
    long demand;
    long order;
    struct {
        double setup;
        double holding; /*inventory hold (+) */
        double shortage; /*inventory short (-) */
        double order;
        double demand;
    } sum = { 0.0, 0.0, 0.0, 0.0, 0.0 };

    PutSeed(123456789);

```

Prof. Vittoria de Nitto Personè

28

28

```

while (index < STOP) {
    index++;
    if (inventory < MINIMUM) {
        order = MAXIMUM - inventory;
        sum.setup++;
        sum.order += order;
    }
    else order = 0;
    inventory += order; /* there is no delivery lag */ demand =
    GetDemand();
    sum.demand += demand;
    if (inventory > demand)
        sum.holding += (inventory - 0.5 * demand);
    else {
        sum.holding += sqr(inventory) / (2.0 * demand);
        sum.shortage += sqr(demand - inventory) / (2.0 * demand);
    }
    inventory -= demand;
}

```

Prof. Vittoria de Nitto Personè

29

29

```

if (inventory < MAXIMUM) {
    order = MAXIMUM - inventory;
    sum.setup++;
    sum.order += order;
    inventory += order;
}

...

```

Prof. Vittoria de Nitto Personè

30

30

for 100 time intervals with an average demand of 27.68  
and policy parameters  $(s, S) = (20, 80)$

average order ..... = 27.68  
setup frequency ..... = 0.36  
average holding level .... = 44.81  
average shortage level ... = 0.14

sls2.dat, con  $n=100$  e  $(s, S) = (20, 80)$ , restituisce:

$$\bar{o} = \bar{d} = 27,68 \quad \bar{m} = 0,36 \quad \bar{l}^+ = 44,81 \quad \bar{l}^- = 0,14$$

Prof. Vittoria de Nitto Personè

31

31

```
int main(void)
{ long seed;
  long index      = 0;
  long inventory = MAXIMUM;
  long demand;
  long order;
  struct {
    double setup;
    double holding; /*inventory held (+) */
    double shortage; /*  inventory short (-)  */
    double order;
    double demand;
  } sum = { 0.0, 0.0, 0.0, 0.0, 0.0 };

  PutSeed(-1); //clock sistema
  GetSeed(&seed);
  printf("\nwith an initial seed of %ld", seed);
```

Prof. Vittoria de Nitto Personè

32

32



for 100 time intervals with an average demand of 27.68  
and policy parameters (s, S) = (20, 80)

average order ..... = 27.68  
setup frequency ..... = 0.36  
average holding level .... = 44.81  
average shortage level ... = 0.14

with an initial seed of 1333437895 (*seed* = -1)  
for 100 time intervals with an average demand of 31.00  
and policy parameters (s, S) = (20, 80)

average order ..... = 31.00  
setup frequency ..... = 0.40  
average holding level .... = 43.39  
average shortage level ... = 0.37

Prof. Vittoria de Nitto Personè

33

33

Discrete-Event Simulation  
case study sis

### Simply Inventory System

equilicely( $a=10, b=50$ )  $\rightarrow$  average demand =  $\frac{a+b}{2} = 30$

if (a, b) = (10, 50) and (s, S) = (20, 80), then the approximate  
steady-state statistics are

$\bar{d}$	$\bar{o}$	$\bar{u}$	$\bar{l}^+$	$\bar{l}^-$	$\left( \begin{array}{l} \text{valore teorico} \\ \text{prodotto da sis2} \end{array} \right)$
30.00	30.00	0.39	42.86	0.26	

(trace-driven (con la traccia)

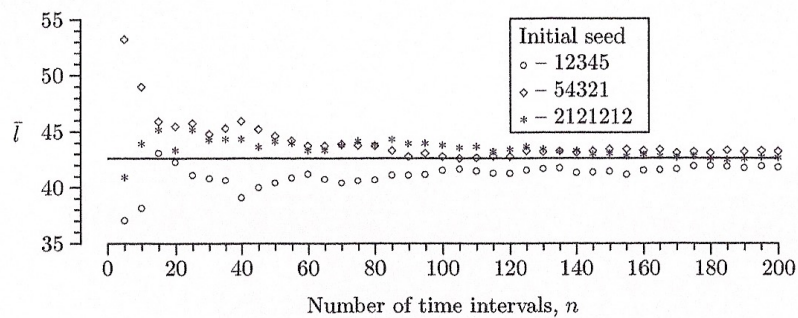
$\bar{o} = \bar{d} = 29.29$        $\bar{u} = 0.39$        $\bar{l}^+ = 42.40$        $\bar{l}^- = 0.25$  )

Prof. Vittoria de Nitto Personè

34

34

The average inventory level  $\bar{I} = \bar{I}^+ - \bar{I}^-$   
approaches steady state after several hundred time intervals



Convergence is slow, erratic, and dependent on the initial seed

Prof. Vittoria de Nitto Personè

35

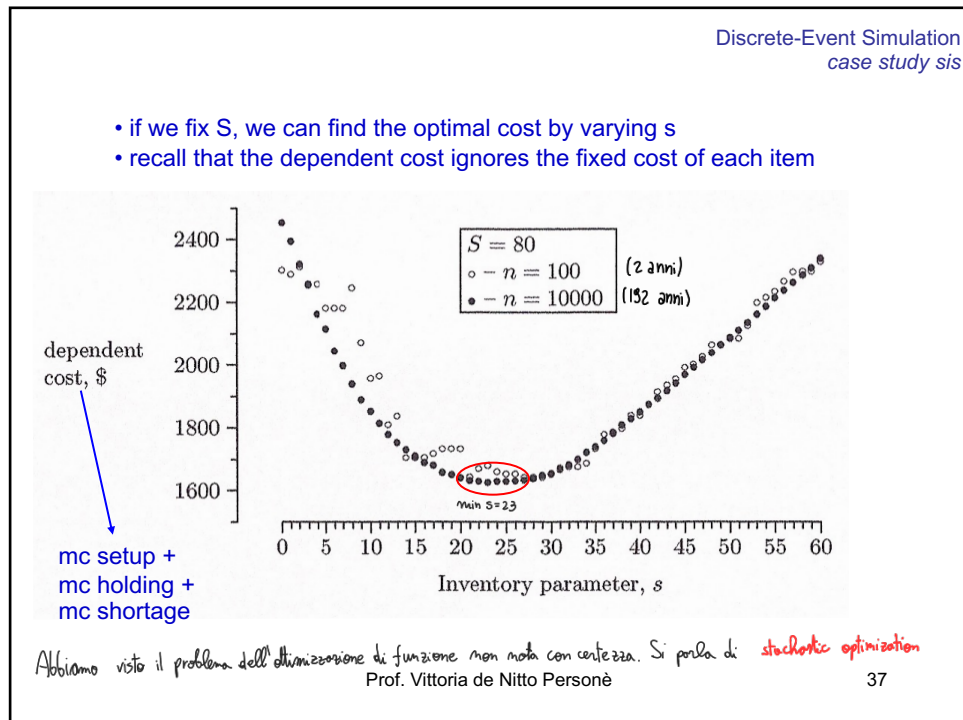
35

- using a fixed initial seed guarantees the exact same demand sequence  
(in the example 12345), *devo cambiare solo lo s, non tocco altri valori!*
- any changes to the system are caused solely by the change of  $s$
- a steady state study of this system is unreasonable:
  - all parameters would have to remain fixed for many years
  - when  $n=100$ , we simulate approximately 2 years
  - when  $n=10000$ , we simulate approximately 192 years

Prof. Vittoria de Nitto Personè

36

36



37

P.104 discrete

Discrete-Event Simulation

### Statistical Considerations (sect. 3.1.3)

- *variance reduction:*  
(intuitive approach) use the same random numbers  
(We kept the same initial seed and changed only  $s$ )

NOTE:  
transient behavior will always have some inherent uncertainty

(poiché prodotto da una discrete event simulation dipende dalla sequenza di variabili random generate durante l'esecuzione del programma)

- *Robust Estimation:*  
when a data point is not sensitive to small changes in assumptions
  - values of  $s$  close to 23 have essentially the same cost
  - Would the cost be more sensitive to changes in  $S$  or other assumed values?

• se  $S$  varia in un intorno di 80, o average demand cambia in un intorno di 30, o introduco delivery lags, la stima  $s=23$  è robusta  $\Leftrightarrow s$  rimane nell'intorno di 23 anche con queste variazioni. In generale ROBUST ESTIMATORS sono insensibili al modello assunto.

Prof. Vittoria de Nitto Personè

38

38

# Exercises

- derive analytical results on p.12 (*example 1 slide*)
- study program ssq2.c; run it and compare output with the results on p.12
- Exercises: 3.1.1, 3.1.2, 3.1.4, 3.1.5, 3.1.6