

Performance Modeling of Computer Systems and Networks

Prof. Vittoria de Nitto Personè

Multi-stream application examples


Università degli studi di Roma Tor Vergata
Department of Civil Engineering and Computer Science Engineering

Copyright © Vittoria de Nitto Personè, 2021
<https://creativecommons.org/licenses/by-nc-nd/4.0/>



1

Multi-stream RNG
Case study ssq



1. one stream per stochastic component

2. uncoupled processes → no-overlapped streams

Prof. Vittoria de Nitto Personè

2

2

Quando uso il multistream (lo userò sempre) non uso la PutSeed, bensì la PlantSeeds, perchè devo 'piantarli' tutti.

Multi-stream RNG
Case study ssq

ssq2 revisited

- Use rngs functions for arrivals and services

```
double GetArrival(void) {
    static double arrival = START;
    SelectStream(0);
    arrival += Exponential(2.0);
    return (arrival);
}

double GetService(void) {
    SelectStream(2);
    return (Uniform(1.0, 2.0));
}
```

- include "rngs.h" and use PlantSeeds(12345)
(in place of PutSeed(12345))

Prof. Vittoria de Nitto Personè
3

3

arrival and service processes are *uncoupled*

stream 0 for arrivals, stream 1 for services

for 10025 jobs

average interarrival time =	1.99
average wait	3.92
average delay	2.41
average service time	1.50
average # in the node ...	1.96
average # in the queue ..	1.21
utilization	0.75

stream 0 for arrivals, stream 2 for services (or e.g. stream 128 to get more separation)

for 10025 jobs

average interarrival time =	1.99
average wait	3.86
average delay	2.36
average service time	1.50
average # in the node ...	1.93
average # in the queue ..	1.18
utilization	0.75

Theoretical values

\bar{r}	\bar{w}	\bar{d}	\bar{s}	\bar{l}	\bar{q}	\bar{x}
2.00	3.83	2.33	1.50	1.92	1.17	0.75

(in riferimento all'esercizio con Unif(1,2))

Prof. Vittoria de Nitto Personè
4

4

Le cose cambiano, stessi arrivi, distribuzioni etc, ma sto prendendo due pezzi di numeri random diversi!
Ci sono piccole oscillazioni, ma sono normali!

Uncoupling Stochastic Processes

perchè non prenderne una $Unif(0,3)$?
ha un'altra varianza ancora, sarebbe un "altro caso".

Consider changing the service process to

$Uniform(0.0, 1.5) + Uniform(0.0, 1.5)$

La media è sempre la stessa (1.5), variabilità qui più ampia.
(perchè è come fosse un'unica uniforme da 0 a 3, sommando)

- Without uncoupling, arrival process sequence would change!
- With uncoupling, the service process "sees" exactly the same arrival sequence
- Important variance reduction technique

Se voglio vedere l'effetto della variabilità, devo toccare solo lei,
e devo tenere i flussi diversi. Non posso cambiare più cose insieme!

Prof. Vittoria de Nitto Personè

5

5

Theoretical values

\bar{r}	\bar{w}	\bar{d}	\bar{s}	\bar{l}	\bar{q}	\bar{x}
2.00	3.83	2.33	1.50	1.92	1.17	0.75

stream 0 for arrivals, stream 1 for services

for 10025 jobs

average interarrival time = 1.99
average wait = 4.29
average delay = 2.78
average service time = 1.51
average # in the node ... = 2.15
average # in the queue .. = 1.40
utilization = 0.76

nuovi risultati, alcuni sono cambiati di
pochissimo (servizio medio),
altri maggiormente (avg wait).

risultati con uniform(1.0, 2.0)

for 10025 jobs
average interarrival time = 1.99
average wait = 3.92
average delay = 2.41
average service time = 1.50
average # in the node ... = 1.96
average # in the queue .. = 1.21
utilization = 0.75

Prof. Vittoria de Nitto Personè

6

6

ssq with Multiple Job Types

- Consider multiple job types, each with its own arrival and service process
- Two job types:
 - arrivi più frequenti service con stessa media, cambia varianza, in classe 1 maggiore.
–Class 0: *Exponential*(4.0) interarrivals, *Uniform*(1.0, 3.0) service
 - Classe 1: *Exponential*(6.0) interarrivals, *Uniform*(0.0, 4.0) service
- Use rngs to allocate a different stream to each stochastic process

Prof. Vittoria de Nitto Personè

7

7

errore frequente: devo sequenziare istanti di arrivo di due classi. A seconda dell'istante di arrivo che stiamo 'trattando' dobbiamo simulare il successivo. Un errore sarebbe generare TUTTI INSIEME GLI ARRIVI in uno stesso momento.

E' un errore quindi generare subito e tutti insieme istanti di arrivo e tempi di servizio.

Arrival process

```
double GetArrival(int *j)          /* j corrisponde to job type */
{
    const double mean[2] = {4.0, 6.0};
    static double arrival[2] = {START, START};
    static int init = 1;
    double temp;
    if (init) {
        SelectStream(0);
        arrival[0] += Exponential(mean[0]);
        SelectStream(1);
        arrival[1] += Exponential(mean[1]);
        init = 0;
    }

    if (arrival[0] <= arrival[1]) {
        *j = 0;
        temp = arrival[0];
        SelectStream(*j);
        arrival[*j] += Exponential(mean[*j]);
        return (temp);
    }
    else {
        *j = 1;
        temp = arrival[1];
        SelectStream(*j);
        arrival[*j] += Exponential(mean[*j]);
        return (temp);
    }
}
```

li genero via via, infatti ora parto sono dai due START.

trovo quale è "venuto" prima

The first arrival is of class 0!

se sto consumando arrival 0, devo generare il prossimo arrivo.

- streams 0 and 1 are used for interarrival times of class 0 and class 1 jobs respectively

Prof. Vittoria de Nitto Personè

8

8

Ad ogni istante di tempo della simulazione, ho array con i prossimi istanti di arrivo. Come abbiamo detto prima, non li genero tutti, ma solo quando ne sto consumando uno!

Service process

```
double GetService(int j)    j = 0 or 1, è la classe.
{
    const double min[2] = {1.0, 0.0};
    const double max[2] = {3.0, 4.0};
    SelectStream(j + 2); mi sposto sullo stream 2 o 3, a seconda di j
    return (Uniform(min[j], max[j]));
}                        Unif[1,3] or Unif[0,4]
```

- j corrisponde to the job type (0 or 1)
- streams 2 and 3 are used for service times of class 0 and 1 respectively
- All four simulated stochastic processes are uncoupled!
- Any process could be changed without altering the random sequence of others!

Prof. Vittoria de Nitto Personè

9

9

Consistency checks

- The “teorethical” steady-state statistics are

\bar{r}	\bar{w}	\bar{d}	\bar{s}	\bar{l}	\bar{q}	\bar{x}	
2.40	7.92	5.92	2.00	3.30	2.47	0.83	exact analytical results, No simulation!

- obvious consistency checks: $\bar{w} = \bar{d} + \bar{s}$ $\bar{l} = \bar{q} + \bar{x}$
- other consistency checks:
 - both job types have avg service time of 2.0 $\rightarrow \bar{s} = 2.00$
 - arrival rate should be $1/4 + 1/6 = 5/12 \rightarrow \bar{r} = 12/5 = 2.40$
 - \bar{x} should be ratio of arrival to service rates

$$\frac{5/12}{1/2} = 5/6 \approx 0.83$$

Prof. Vittoria de Nitto Personè

10

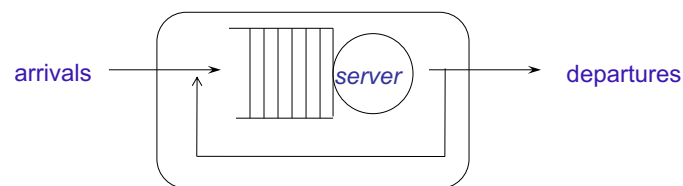
10

Exercises

- Exercises: 3.2.3, 3.2.4, 3.2.7

un job che arriva al sistema potrebbe richiedere 'servizi' ulteriori. Discrete-Event Simulation
Il completamento del servizio non è detto che corrisponda alla partenza. ssq with feedback
Qui una partenza corrisponde al completamento di un numero di servizi richiesti.
C'è quindi completamento singolo e totale.

Single Server Queue with feedback



- If the service a job receives is incomplete or unsatisfactory, the job feeds back
- Completion of service and departure now have different meanings

Non so a priori quante volte richiede feedback un job. E' probabilistico.


p.116 discrete

DE simulation
ssq with feedback

Model Considerations

"ni" quindi $1/n_i$ è tempo di servizio per UN SOLO passaggio.

ν service rate



bilanciamento flusso

$$\lambda = \lambda' (1 - \beta)$$

- When feedback occurs the job joins the queue consistent with the queue discipline
- The decision to depart or feed back is random with feedback probability β

Come calcolare l'utilizzazione? perchè non è più semplicemente λ/ν $\rho = \frac{\lambda'}{\nu}$

Prof. Vittoria de Nitto Personè

13

13

Discrete-Event Simulation
ssq with feedback

Model Considerations

- Feedback is independent of past history
- In theory, a job may feed back arbitrarily many times
- Jobs from outside the system are merged with jobs from the feedback process
- The steady-state request-for-service rate is larger than λ by the positive additive factor $\beta \bar{x} \nu$ probabilità feedback * avg job in servizio * rate servizio
- Note that \bar{x} increases with feedback but $1/\nu$ is the average service time per request

Prof. Vittoria de Nitto Personè

14

14

Flow Balance and Saturation

- Jobs flow into the service node at the average rate of λ
- To remain flow balanced jobs must flow out of the service node at the same average rate
- The average rate at which jobs flow out of the service node is

$$\bar{x}(1-\beta)v$$


- Flow balance is achieved when $\lambda = \bar{x}(1-\beta)v$
- Saturation occurs when $\bar{x}=1$ or as $\beta \rightarrow 1 - \lambda/v$

Prof. Vittoria de Nitto Personè

15

15

Coda con feedback



$$\rho = \lambda' / v$$

$$\rho = \lambda / (1-\beta)v$$

Saturation:

$$\rho \rightarrow 1$$

$$\lambda / (1-\beta)v \rightarrow 1$$

$$\beta \rightarrow 1 - \lambda / v$$

Flow Balance: $\lambda = \lambda' (1-\beta)$

$$\lambda' = \lambda / (1-\beta)$$

Prof. Vittoria de Nitto Personè

16

16

- Feedback is independent of past history
- In theory, a job may feed back arbitrarily many times

snippet di codice per vedere se per un job c'è o non c'è feedback

```
int GetFeedback(double beta)    /* 0.0 <= β < 1.0 */
{
    SelectStream(2);
    if (Random() < beta)
        return (1);            /* feedback*/
    else
        return (0);            /* no feedback */
}
```

17

Statistical considerations

Nella coda si mischiano job arrivanti dall'esterno a job ritornati causa feedback.

- Index $i=1, 2, 3, \dots$ counts jobs that enter the service node
 - fed-back jobs are not recounted (job $i=3$ che rientra grazie al feedback, è sempre job 3, non lo rinomino)
- Using this indexing, all job-averaged statistics remain valid
- We must update delay times, wait times and service times for each feedback
- Jobs from outside the system are merged with jobs from the feedback process
- The steady-state request-for-service rate is larger than λ by the positive additive factor $\beta \bar{x} \nu$ = beta * che ci sia qualcuno in servizio * tasso servizio
 - tempo servizio medio cresce al crescere di chi riceve feedback
- Note that \bar{s} increases with feedback but $1/\nu$ is the average service time per request

18

Ci serve struttura dati per ordinare i job con feedback rispetto a quelli che entrano normalmente.

Algorithm and Data Structure Considerations

job index	1	2
Arrival/feedback	1	3
service	9	3
departure	10	13

(istante completamento Ci)

job 2 arriva quando job1 in servizio, inizia a lavorare a $t = 10$, quando job 1 esce.
job 2 richiede feedback.

Prof. Vittoria de Nitto Personè

19

19

i job aventi departure rossi richiedono feedback.

Algorithm and Data Structure Considerations

job index	1	2	3	4	5	.	6
Arrival/feedback	1	3	4	7	10		14
service	9	3	2	4	7		6
departure	10	13	15	19	26		37

devo vedere sequenza dei tempi per dire dove collocare il '13', perchè subisce feedback.

Prof. Vittoria de Nitto Personè

20

20

Algorithm and Data Structure Considerations

job index	1	2	3	4	5	.	6	.	7	8	.	9	...
Arrival/feedback	1	3	4	7	10	13	14	15	19	24	26	30	...
service	9	3	2	4	7	5	6	3	4	6	3	7	...
departure	10	13	15	19	26	31	37	40	44	50	53	60	...

At the computational level, some algorithm and data structure is necessary


il job 2 va tra i job 5 e 6, poichè il suo $\text{departure} == \text{arrival} = 13$, e lo metto tra arrival 10 e arrival 14

Theoretical values

\bar{r}	\bar{w}	\bar{d}	\bar{s}	\bar{l}	\bar{q}	\bar{x}
2.00	3.83	2.33	1.50	1.92	1.17	0.75

Program ssq2 was modified to incorporate immediate feedback

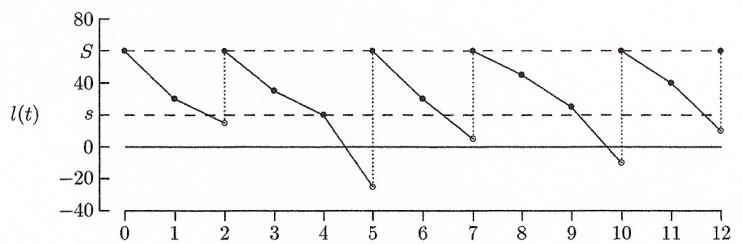
- interarrivals = *Exponential*(2.0)
- service times = *Uniform*(1.0, 2.0)



- It appears saturation is achieved as $\beta \rightarrow 0.25$

Inventory system with delivery lag

- *delivery lag* (dl) occurs when orders are not delivered immediately
- Lag is assumed to be random and independent of order size
- Without lag, inventory jumps occur only at inventory review times




Prof. Vittoria de Nitto Personè

23

23

Inventory system with delivery lag

- With delivery lag, inventory jumps occur at arbitrary times



- The last order is assumed to have no lag
- We assume that orders are delivered before the next inventory review
- With this assumption, there is no change to the specification model


il livello di scorte può scendere ancora (cerchi rossi) perchè non soddisfo subito la richiesta, e quindi quando a inizio settimana ricevo gli ordini non ne avrò in quantità massima, poichè da quando ho richiesto la quantità, nel mentre ne ho perse altre.

Prof. Vittoria de Nitto Personè

24

24

- If $l_{i-1} \geq s$ the equations for \bar{l}_i^+ and \bar{l}_i^- remain correct
- If $l_{i-1} < s$, the time-averaged holding and shortage intervals must be modified
 - the dl for interval i is $0 < \delta_i < 1$



Prof. Vittoria de Nitto Personè

25

25

Consistency Checks

- It is fundamentally important to verify extended models with the parent model (before the extension)
 - Set system parameters to special values
- Set $\beta = 0$ for the ssq with feedback
 - Verify that all statistics agree with parent
- Using the library `rngs` facilitates this kind of comparison
- It is a good practice to check for intuitive “small-perturbation” consistency
 - Use a small, but non-zero β and check that appropriate statistics are slightly larger

Prof. Vittoria de Nitto Personè

26

26 Un buon controllo è vedere se nel nuovo codice non ho introdotto errori nuovi, quindi se ho modificato il codice in vista di parametri come delivery lag, mi aspetterei che questo nuovo codice, se testato nel caso originale (no delivery lag), fornisca i risultati originali prodotti dal vecchio programma.
E' come dire: complico il modello introducendo un rumore, ma se tale rumore è 0, mi aspetto risultati uguali al vecchio modello che non considera il rumore.


- For the InvSys with delivery lag, $\delta_i = 0.0$ iff no order during i^{th} interval, $0 < \delta_i < 1.0$ otherwise
 - The InvSys is lag-free iff $\delta_i = 0.0$ for all i
 - If (S, s) are fixed then, even with small dl:
 - $\bar{o}, \bar{d}, \bar{u}$ are the same regardless of delivery lag
 - Compared to the lag-free system, \bar{I}_i^+ will decrease (i costi)
 - Compared to the lag-free system, \bar{I}_i^- will increase or remain unchanged (penalità data dal fatto che non soddisfo la richiesta cresce)
- con delivery lag, shortage maggiore, e spese holding minori (c'è meno merce)

Prof. Vittoria de Nitto Personè

27

27

- Delivery lags are independent $Uniform(0.0, 1.0)$ random variates

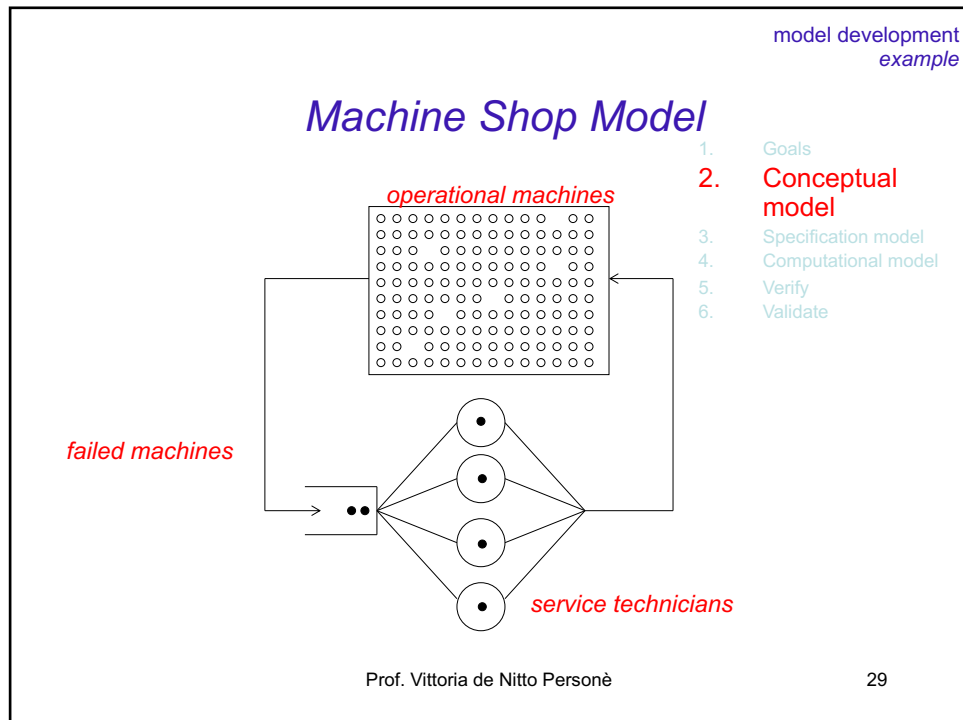


- delivery lag causes \bar{I}_i^+ to decrease and \bar{I}_i^- to increase or remain the same
- with $C_{\text{hold}} = \$25$ and $C_{\text{short}} = \$700$, cause shift up and to the left

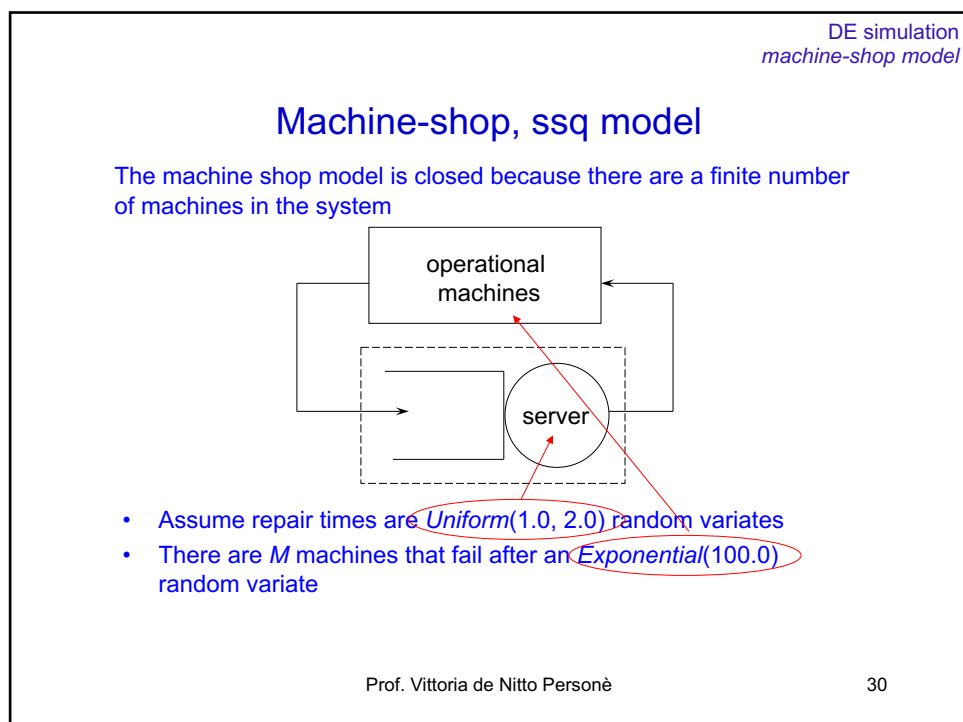
Prof. Vittoria de Nitto Personè

28

28



29



30

Program ssms

- program `ssms` simulates a Single Server Machine-Shop
- the library `rngs` is used to uncouple the random process
- the failure process is defined by the array `failures`
 - a $O(M)$ search is used to find the next failure (M è numero macchine)
 - alternate data structures can be used to increase computational efficiency

Prof. Vittoria de Nitto Personè

31

31

```
double GetFailure(void)
{
    SelectStream(0);
    return (Exponential(100.0));
}

double NextFailure(double failure[], int *m)
{
    int i = 0;
    double t = failure[0];

    *m = i;
    for (i = 1; i < M; i++)
        if (failure[i] < t) {
            t = failure[i];
            *m = i;
        }
    return (t);
}

double GetService(void)    qui si usa uno stream diverso ovviamente!
{
    SelectStream(1);
    return (Uniform(1.0, 2.0));
}
```

Prof. Vittoria de Nitto Personè

32

32


```

                                inalterato
int main(void)
{
    long    index    = 0;        /* job (machine failure) index */
    double  arrival   = START;   /* time of arrival (failure) */
    double  delay;              /* delay in repair queue */
    double  service;           /* service (repair) time */
    double  wait;              /* delay + service */
    double  departure = START;   /* time of service completion */
    int     m;                 /* machine index 0,1,...(M-1) */
    double  failure[M];         /* list of next failure times */
    struct {                    /* sum of ... */
        double wait;           /* wait times */
        double delay;          /* delay times */
        double service;        /* service times */
        double interarrival;    /* interarrival times */
    } sum = {0.0, 0.0, 0.0};
}

```

Prof. Vittoria de Nitto Personè

33

33

```

PlantSeeds(123456789);

for (m = 0; m < M; m++)          /* initial failures */
    failure[m] = START + GetFailure();

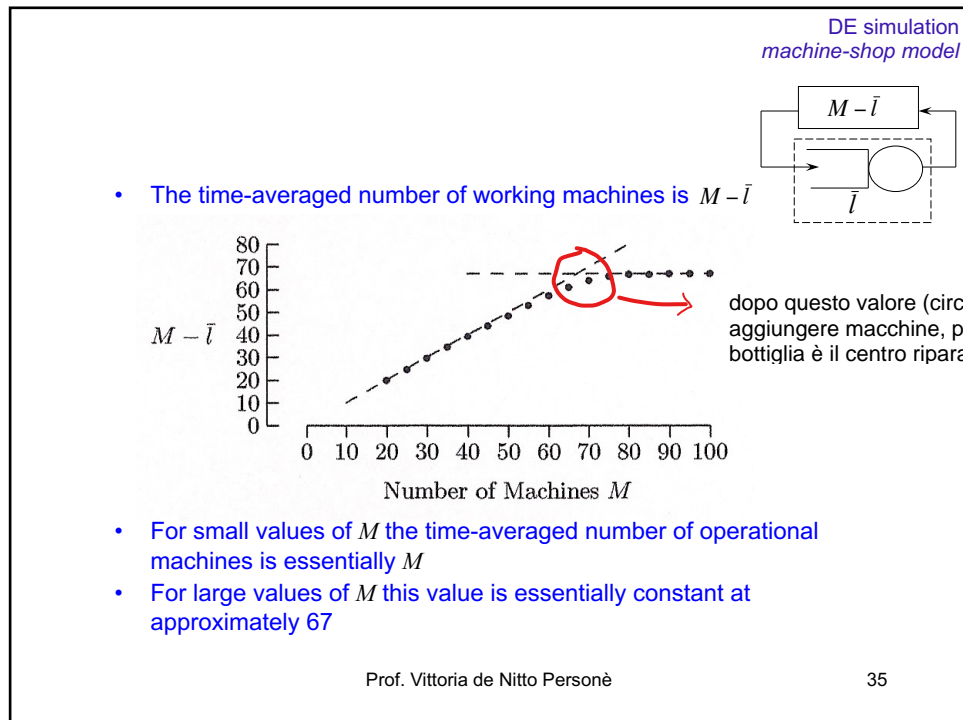
while (index < LAST) {
    index++;
    arrival   = NextFailure(failure, &m);
    if (arrival < departure)
        delay   = departure - arrival;
    else
        delay   = 0.0;
    service    = GetService();
    wait       = delay + service;
    departure  = arrival + wait;          /* completion of service */
    failure[m] = departure + GetFailure(); /* next failure, machine m */
    sum.wait   += wait;
    sum.delay  += delay;
    sum.service += service;
}
sum.interarrival = arrival - START;

```

Prof. Vittoria de Nitto Personè

34

34



35

DE simulation
machine-shop model

Exercises

- Exercises: 3.3.2, 3.3.3, 3.3.4, 3.3.7

Prof. Vittoria de Nitto Personè

36

36