Performance Modeling
of Computer Systems and Networks

Prof. Vittoria de Nitto Personè

Interval Estimation

Università degli studi di Roma Tor Vergata
Department of Civil Engineering and Computer Science Engineering

1

model development

Algorithm 1.1: how to develop a model

1. Goals and objectives
2. *Conceptual* model (cm)
3. Convert cm into a *specification* model (sm)
4 Convert sm into a *computational* model (cptm)
5 Verify
6 Validate

2

Algorithm 1.2: using the resulting model

7. Design simulations experiments
   - What parameters should be varied?
   - perhaps many combinatoric possibilities
8. Make production runs
   - Record initial conditions, input parameters
   - Record statistical output
9. Analyze the output
   - Random components → statistical analysis
     (means, standard deviations, percentiles, histograms etc.)
10. Make decisions
    - The step9 results drive the decisions → actions
    - Simulation should be able to correctly predict the outcome of these actions (→ further refinements)
11. Document the results
    - summarize the gained insights in specific observations and conjectures useful for subsequent similar system models

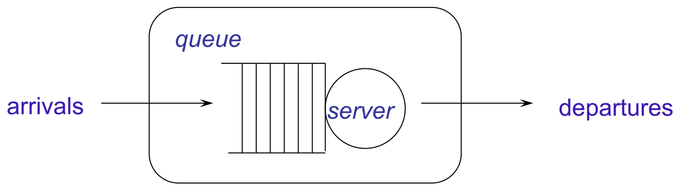

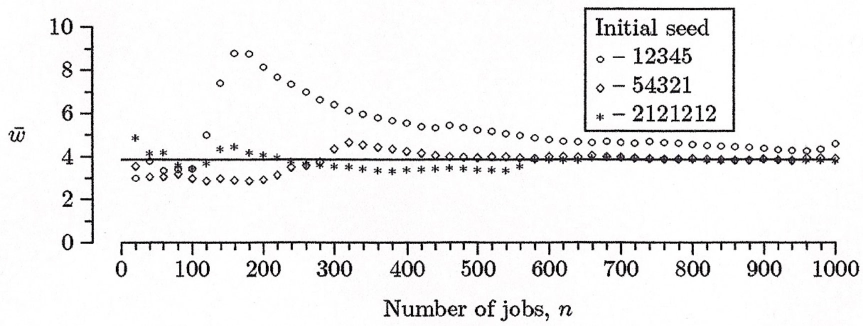

• The accumulated average wait was printed every 20 job

Consider a sample $x_1, x_2, \ldots, x_n$ (continuous or discrete) with

$$\bar{x} = \frac{1}{n}\sum_{i=1}^{n} x_i \qquad s^2 = \frac{1}{n}\sum_{i=1}^{n}(x_i - \bar{x})^2$$

job average

Consider a piecewise constant sample path

$$x(t) = \begin{cases} x_1 & t_0 < t \le t_1 \\ x_2 & t_1 < t \le t_2 \\ . & . \\ . & . \\ x_n & t_{n-1} < t \le t_n \end{cases}$$

processi stocastici
che variano nel tempo

$$\bar{x} = \frac{1}{\tau}\int_0^\tau x(t)dt = \frac{1}{t_n}\sum_{i=1}^{n} x_i \delta_i \qquad s^2 = \frac{1}{\tau}\int_0^\tau (x(t) - \bar{x})^2 dt = \frac{1}{t_n}\sum_{i=1}^{n}(x_i - \bar{x})^2 \delta_i$$

Central limit theorem

variabili aleatorie random
If $X_1, X_2, \ldots, X_n$ is an iid sequence of random variables (RVs) with

- common mean μ
- common standard deviation σ

and if $\bar{X}$ is the (sample) mean of these RVs $\bar{X} = \frac{1}{n}\sum_{i=1}^{n} X_i$
then $\bar{X}$ approaches a *Normal*(μ, $\sigma / \sqrt{n}$)
as $n \to \infty$

S la dimensione del campione, la distribuzione è distribuita
come normale di stessa media mu e deviazione std sigma/sqrt(n).
Ovvero la media di un campione molto grande ha questo comportamento fissato.

L'idea è:
costruisco n campioni di certa lunghezza uguale per tutti, per ogni campione calcolo media e varianza campionaria.
Tutti questi campioni li diamo in pasto al programma che genera gli istogrammi caso discreto e caso continuo.
Tale programma ci ritorna media, deviazione std, etc dell'istogramma.

Sample Mean Distribution

- Choose one of the random variate generators in `rvgs` to generate a sequence of random variable samples with fixed sample size $n > 1$
- with the n-point samples indexed $j=1, 2, \ldots,$ the corresponding sample mean $\bar{x}$ and sample standard deviation s can be calculated using Welford's algorithm

$$\underbrace{x_1, x_2, \ldots, x_n}_{\bar{x}_1, s_1}, \underbrace{x_{n+1}, x_{n+2}, \ldots, x_{2n}}_{\bar{x}_2, s_2}, \underbrace{x_{2n+1}, x_{2n+2}, \ldots, x_{3n}}_{\bar{x}_3, s_3}, x_{3n+1}$$

- A continuous-data histogram can be created using program `cdh`

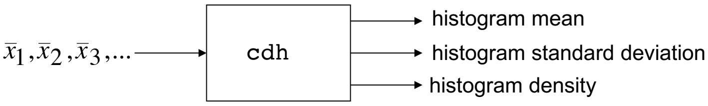

7

Indipendentemente dalla dimensione del campione (non quanti campioni sono), la media è 'mu', la dev.std è sigma/sqrt(n).
Per n grande replico anche la forma della distribuzione, cioè se confronto densità teorica con questo risultato li trovo simili.

Properties of Sample Mean Histogram

If we denote with μ and σ the theoretical mean and standard deviation respectively of the random variates

- independent of n
  - the histogram mean is approximately μ
  - the histogram standard deviation is approximately $\sigma / \sqrt{n}$
- if n is sufficiently large,
  - the histogram density approximates the *Normal*$(\mu, \sigma / \sqrt{n})$ pdf

8

SKIP FINO A SLIDE 20.

10000 *n*-point *Exponential*(μ) samples

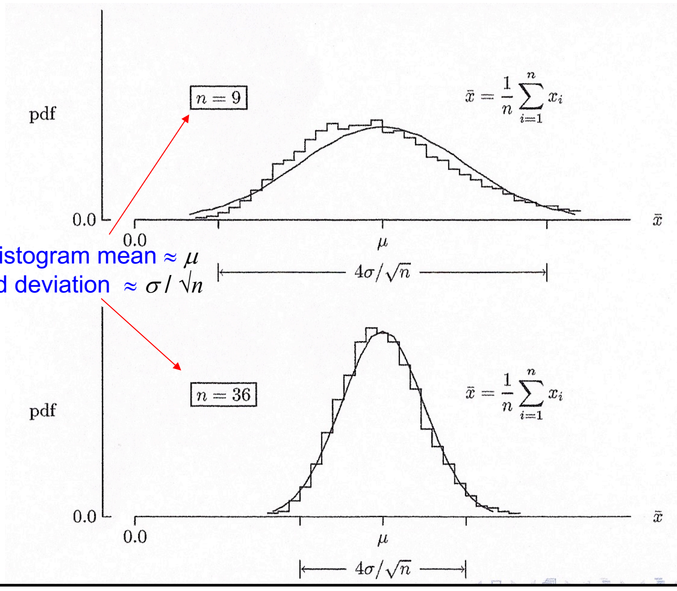

histogram mean $\approx \mu$
hist stand deviation $\approx \sigma / \sqrt{n}$

Example

- The histogram density corresponding to the 36-point sample means is closely matched by the pdf of a *Normal*(μ, $\sigma / \sqrt{n}$) RV

    for *Exponential*(μ) samples, n=36 is large enough for the sample mean to be approximately *Normal*(μ, $\sigma / \sqrt{n}$)

- The histogram density corresponding to the 9-point sample means matches relatively well, but with a skew to the left

    - n=9 is not large enough

Example (cont.)

- Essentially all of the sample means are within an interval of width of $4\sigma / \sqrt{n}$ centered about μ
- because $n \to \infty$ as $\sigma / \sqrt{n} \to 0$, if n is large, all the sample means will be close to μ
- In general:
  - the accuracy of the *Normal*(μ, $\sigma / \sqrt{n}$) pdf approximation is dependent on the shape of a fixed population pdf
  - If the samples are drawn from a population with
    - a highly asymmetric pdf (like the *Exponential*(μ) pdf): n may need to be as large as 30 or more for good fit
    - a pdf symmetric about the mean (like the *Uniform*(a,b) pdf): n as small as 10 or less may produce a good fit

Examples of Linear Data Transformations

- suppose $x_1, x_2, \ldots, x_n$ measured in seconds
  - to convert to minutes, let $x'_i = x_i/60$

  (a=1/60, b=0)

$$\bar{x}' = \frac{45}{60} = 0.75 \qquad s' = \frac{15}{60} = 0.25 \quad \text{(minutes)}$$

- *standardize* data (a=1/s, b= $-\bar{x}/s$)

$$x'_i = \frac{1}{s}x_i - \frac{\bar{x}}{s} \qquad x'_i = \frac{x_i - \bar{x}}{s}$$

Then

$$\bar{x}' = 0 \qquad s' = 1$$

Used to avoid problems with very large (or small) valued samples

Standardized Sample Mean Distribution

We can standardize the sample means $\bar{x}_1, \bar{x}_2, \bar{x}_3, \ldots$ by subtracting μ and dividing the result by $\sigma / \sqrt{n}$ to form the standardized sample means $z_1, z_2, z_3, \ldots$ defined by

$$z_j = \frac{\bar{x}_j - \mu}{\sigma / \sqrt{n}} \qquad j = 1, 2, 3, \ldots$$

- Generate a continuous-data histogram for the standardized sample means by program `cdh`

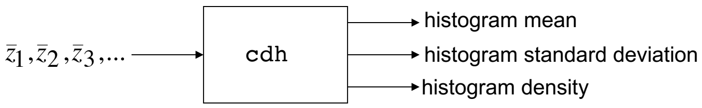

Properties of Standardized Sample Mean Histogram

- independent of n
  - the histogram mean is approximately 0
  - the histogram standard deviation is approximately 1
- if n is sufficiently large,
  - the histogram density approximates the *Normal*$(0,1)$ pdf

t-Statistic Distribution

Definition

- each sample mean $\bar{x}_j$ is a point estimate of μ
- each sample variance s_j^2 is a point estimate of σ^2
- each sample standard deviation s_j is a point estimate of σ

Want to replace *population* standard deviation σ with *sample* standard deviation s_j in standardization equation

$$z_j = \frac{\bar{x}_j - \mu}{\sigma / \sqrt{n}} \qquad j = 1,2,3,\ldots$$

$$\sigma / \sqrt{n} \rightarrow \frac{s_j}{\sqrt{n-1}}$$

- Calculate the *t*-statistic

$$t_j = \frac{\bar{x}_j - \mu}{s_j / \sqrt{n-1}} \qquad j = 1,2,3,\ldots$$

- Generate a continuous-data histogram using `cdh`

$$t_1, t_2, t_3, \ldots \longrightarrow \boxed{\texttt{cdh}} \longrightarrow \begin{cases} \text{histogram mean} \\ \text{histogram standard deviation} \\ \text{histogram density} \end{cases}$$

Properties of *t*-statistic histogram

- if $n > 2$, the histogram mean is approximately 0
- if $n > 3$, the histogram standard deviation is approximately $\sqrt{(n-1)/(n-3)}$
- if n is sufficiently large, the histogram density approximates the pdf of a *Student*$(n-1)$ random variable

Example (cont.)

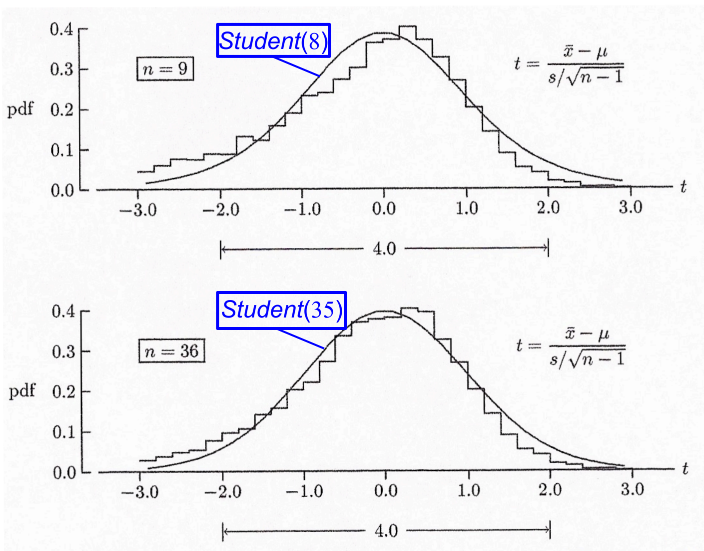

Example (cont.)

- The histogram mean and standard deviation are approximately 0.0 and $\sqrt{(n-1)/(n-3)} \cong 1.0$ respectively
- The histogram density corresponding to the 36-point sample means matches the pdf of a *Student*(35) RV relatively well
- The histogram density corresponding to the 9-point sample means matches the pdf of a *Student*(8) RV, but not as well

19

RIPARTI DA QUI.

Interval Estimation

Theorem 2

If $x_1, x_2, \ldots, x_n$ is an independent random sample from a "source" of data with unknown mean μ, if $\bar{x}$ and s are the mean and standard deviation of this sample, and if n is large, it is approximately true that

$$t = \frac{\bar{x} - \mu}{s / \sqrt{n-1}}$$

is a *Student*$(n-1)$ random variate

- provides the justification for estimating an interval that is likely to contain the mean μ
- as $n \to \infty$, the *Student*$(n-1)$ distribution becomes indistinguishable from *Normal*$(0,1)$

tolgo dipendenza da 'n'.

20

Se prendo campione random, dove i suoi elementi sono indipendenti, di dimensione 'n', media è mu IGNOTA, se calcolo media e dev.std. campionaria. Se 'n' è grande, possiamo dire che questa variabile che fuoriesce è una Student(n-1).
L'idea è stimare intervallo che contenga la media mu.

skip

Suppose

- T is a *Student*(n-1) random variable
- α is a "confidence parameter" with $0.0 < \alpha < 1.0$

Then there exists a corresponding positive real number t^*

$$\Pr(-t^* \le T \le t^*) = 1 - \alpha$$

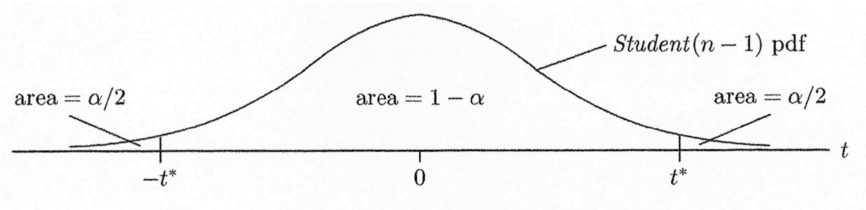

21

Interval Estimation

- suppose μ is unknown. Since $t \approx$ *Student*(n-1)

$$-t^* \le \frac{\bar{x}-\mu}{s/\sqrt{n-1}} \le t^*$$

will be approximately true with probability 1- α

- right inequality: $\frac{\bar{x}-\mu}{s/\sqrt{n-1}} \le t^* \Leftrightarrow \bar{x}-\mu \le \frac{t^* s}{\sqrt{n-1}} \Leftrightarrow \bar{x} - \frac{t^* s}{\sqrt{n-1}} \le \mu$
- left inequality: $-t^* \le \frac{\bar{x}-\mu}{s/\sqrt{n-1}} \Leftrightarrow -\frac{t^* s}{\sqrt{n-1}} \le \bar{x}-\mu \Leftrightarrow \mu \le \bar{x} + \frac{t^* s}{\sqrt{n-1}}$

So, with probability 1- α (approximately),

$$\bar{x} - \frac{t^* s}{\sqrt{n-1}} \le \mu \le \bar{x} + \frac{t^* s}{\sqrt{n-1}}$$

22

Se questo campione estratto è indipendente, calcolo media e std dev campionaria, n grande, allora posso fissare livello di confidenza/affidabilità con cui voglio fare questa stima 'alfa', allora posso associare t* tale che la probabilità di cadere in un intorno della media campionaria sia 1-alfa.
Tutto ciò a livello approssimativo.

Theorem 3

If

$x_1, x_2, \ldots, x_n$ is an independent random sample from a "source" of data with unknown mean μ

- if $\bar{x}$ and s are the sample mean and sample standard deviation
- n is large

Then, given a confidence parameter α with $0.0 < \alpha < 1.0$, there exists an associated positive real number t^* such that

$$Pr\left(\bar{x} - \frac{t^* s}{\sqrt{n-1}} \leq \mu \leq \bar{x} + \frac{t^* s}{\sqrt{n-1}}\right) \cong 1 - \alpha$$

è un numero!

Example margine errore: 5%

- If $\alpha = 0.05$, we are 95% confident that μ lies somewhere between

$$\bar{x} - \frac{t^* s}{\sqrt{n-1}} \quad \text{and} \quad \bar{x} + \frac{t^* s}{\sqrt{n-1}}$$

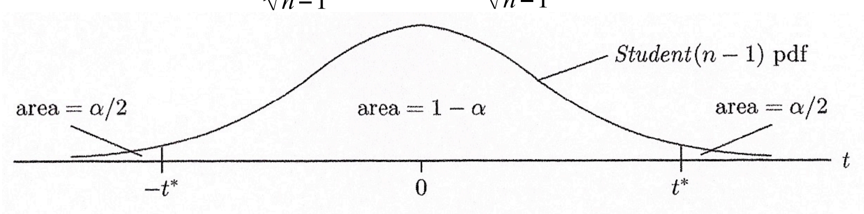

- for a fixed sample size n and level of confidence $1-\alpha$, use `rvms` to determine t^* = *idfStudent*($n-1, 1-\alpha/2$) inversa student
- ex. $n = 30$, $\alpha = 0.05$ $\rightarrow$ t^* = *idfStudent*(29, 0.975) $\cong$ 2.045

Definition

- The interval defined by the two endpoints $\bar{x} \pm \frac{t^* s}{\sqrt{n-1}}$

  is a $(1-\alpha)$x100% confidence *interval estimate* for μ

- $(1-\alpha)$ is the *level of confidence* associated with this interval estimate and t^* is the *critical value* of t

Algorithm

x1 viene fuori da un run
x2 viene fuori da un altro run,
ogni x_i nasce da un run diverso indipendente.
Ogni elemento del campione viene fuori da un run.

To calculate an interval estimate for the unknown mean μ of the population from which a random sample $x_1, x_2, \ldots, x_n$ was drawn:

campione 'ben costruito'

- pick a level of confidence $1-\alpha$ (tipically α =0.05)
- calculate the sample mean $\bar{x}$ and standard deviation s (use Welford's algorithm)
- calculate the critical value t^* = *idfStudent*$(n-1, 1-\alpha/2)$
- calculate the interval endpoints $\bar{x} \pm \frac{t^* s}{\sqrt{n-1}}$

If n is sufficiently large, then you are $(1-\alpha)$x100% confident that the mean μ lies within the interval. The midpoint of the interval is $\bar{x}$

Il campione x1,x2,...,xn deve ben rappresentare l'intera popolazione, in modo da poter dire che ciò che osservo nel campione vale per tutti! Per far questo, le componenti x1,x2,..,xn devono essere indipendenti.

Tradeoff - Confidence Versus Sample Size

- For a fixed sample size
  - More confidence can be achieved only at the expense of a larger interval
  - A smaller interval can be achieved only at the expense of less confidence

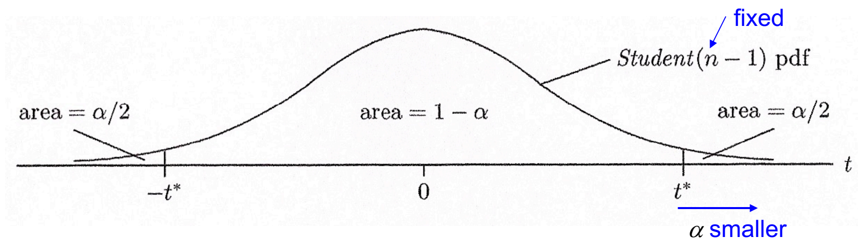

Muovendo alfa, ed essendo 'n' fisso, la
curva si alza o si abbassa a seconda dell'operazione.

α larger

27

Per essere più affidabile, allargo alfa, l'indicazione che ho è poco significativa.
(è come dire che al 100% la media cade in (-infinito, + infinito), poco utile.

Example

- The random sample of size $n = 10$:

1.051	6.438	2.646	0.805	1.505
0.546	2.281	2.822	0.414	1.307

is drawn from a population with unknown mean μ

$$\bar{x} = \frac{1}{n}\sum_{i=1}^{n} x_i \qquad s^2 = \frac{1}{n}\sum_{i=1}^{n}(x_i - \bar{x})^2 \qquad s = \sqrt{s^2}$$

$$\bar{x} = 1.982$$

28

Example

- The random sample of size $n = 10$:

1.051	6.438	2.646	0.805	1.505
0.546	2.281	2.822	0.414	1.307

  is drawn from a population with unknown mean μ

- $\bar{x} = 1.982$ and $s = 1.690$
- to calculate a 90% confidence interval estimate:
  - determine $t^* =$ *idfStudent*(9, 0.95) $\simeq 1.833$
  - interval: $1.982 \pm (1.833)(1.690/\sqrt{9}) = 1.982 \pm 1.032$

$$\bar{x} \pm \frac{t^* s}{\sqrt{n-1}}$$

- we are approximately 90% confident that μ is between 0.950 and 3.014

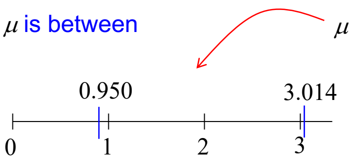

Example (cont.)

- To calculate a 95% confidence interval estimate:
  - determine: $t^* =$ *idfStudent*(9, 0.975) $\simeq 2.262$
  - interval: $1.982 \pm (2.262)(1.690/\sqrt{9}) = 1.982 \pm 1.274$

più affidabilità?
campione più largo,
ma meno indicativo

- We are approximately 95% confident that μ is between 0.708 and 3.256

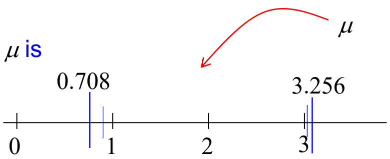

- To calculate a 99% confidence interval estimate:
  - determine: $t^* =$ *idfStudent*(9, 0.995) $\simeq 3.250$
  - interval: $1.982 \pm (3.250)(1.690/\sqrt{9}) = 1.982 \pm 1.832$
- We are approximately 99% confident that μ is between 0.150 and 3.814
- Note: n=10 is <u>not</u> large

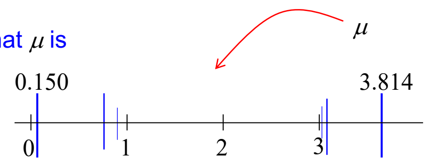

1. starting from a sample $x_1, x_2, \ldots, x_n$

- Program `estimate` automates the interval estimation process
- A typical application: estimate the value of an unknown population mean μ by using n replications to generate an independent random variate sample $x_1, x_2, \ldots, x_n$
- Function Generate() represents a discrete-event or Monte Carlo simulation program that returns a random variate output x

Using the Generate Method

```
for (i = 1; i <= n; i++)
        xi = Generate();
return x1, x2, . . . , xn;
```

ci genera un certo campione come risultato di una simulazione Montecarlo.

- Given a level of confidence $1 - \alpha$, program `estimate` can be used with $x_1, x_2, \ldots, x_n$ to compute an interval estimate for μ

estimate.c

formule di Welford

$$\bar{x}_i = \bar{x}_{i-1} + \frac{1}{i}(x_i - \bar{x}_{i-1}) \qquad v_i = v_{i-1} + \left(\frac{i-1}{i}\right)(x_i - \bar{x}_{i-1})^2$$

```
#include <math.h>
#include <stdio.h>
#include "rvms.h"
#define LOC 0.95                /* level of confidence, */
                                           /* use 0.95 for
95% confidence */
  int main(void)
{ long   n     = 0;            /* counts data points */
  double sum  = 0.0;
  double mean = 0.0;
  double data;
  double stdev;
  double u, t, w;
  double diff;
  while (!feof(stdin)) {   /* use Welford's one-pass method */
    scanf("%lf\n", &data); /* to calculate the sample mean  */
    n++;                   /* and standard deviation        */
    diff  = data - mean;
    sum  += diff * diff * (n - 1.0) / n;
    mean += diff / n;
  }
  stdev  = sqrt(sum / n)
```

$t^* = idfStudent(n - 1, 1 - \alpha/2)$

```
if (n > 1) {
  u = 1.0 - 0.5 * (1.0 - LOC);      /* interval parameter  */
  t = idfStudent(n - 1, u);         /* critical value of t */
  w = t * stdev / sqrt(n - 1);      /* interval half width */
  printf("\nbased upon %ld data points", n);
  printf(" and with %d%% confidence\n", (int) (100.0 * LOC + 0.5));
  printf("the expected value is in the interval");
  printf("%10.2f +/- %6.2f\n", mean, w);  }
else
  printf("ERROR - insufficient data\n");
return (0);}
```


33

Vorrei buon livello di confidenza, e dimensione del campione idonea per avere intervallo abbastanza stretto.

Tradeoff - Confidence Versus Sample Size

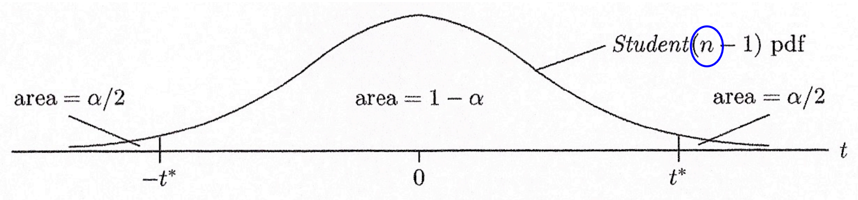

The only way to make the interval smaller without lessening the level of confidence is to increase the sample size

$$\bar{x} \pm \frac{t^* s}{\sqrt{n-1}}$$

basta fare simulazioni più lunghe

- Good news: with simulation, we can collect more data
- Bad news: interval size decreases with $\sqrt{n}$, not n

decrescita più lenta dell'intervallo, rispetto a quello del campione.

34

Discrete Simulation
Interval Estimation

How Much More Data Is Enough?

- How large should n be to achieve an interval estimate $\bar{x} \pm w$ where w is user-specified? fisso ampiezza, cerco 'n'. Questo posso farlo.

- Answer: Use Welford's Algorithm with the algorithm p. 28 to iteratively collect data until a specified interval width is achieved

Note: if n is large then t^* is essentially independent of n

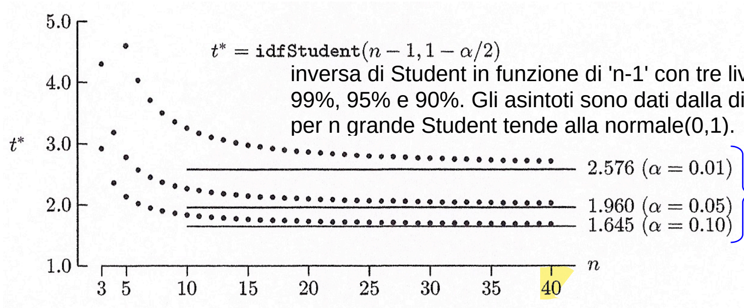

inversa di Student in funzione di 'n-1' con tre livelli di confidenza diversi, 99%, 95% e 90%. Gli asintoti sono dati dalla distribuzione normale, perchè per n grande Student tende alla normale(0,1).

asymptotic value t^*

Prof. Vittoria de Nitto Personè 35

35

Come vediamo, già con dimensione n=40 posso usare una Normale rispetto ad una Student. Nei programmi abbiamo anche la Student, potremmo usare direttamente quella, ma Student dipende da 'n', mentre la Normale no!
Con Student >40, possiamo liberarci dal vincolo di sapere 'n' e passare direttamente alla Normale.

11/05/2023

Discrete Simulation
Interval Estimation

Asymptotic Value of t^*

- The asymptotic (large n) value of t^* is
$$t^*_\infty = \lim_{n\to\infty} idfStudent(n-1, 1-\alpha/2) = idfNormal(0.0, 1.0, 1-\alpha/2)$$
- Unless α is very close to 0.0, if $n > 40$, the asymptotic value t^*_∞ can be used

- If $n > 40$ and wish to construct a 95% confidence interval estimate, $t^*_\infty = 1.960$ can be used in the algorithm on p.28

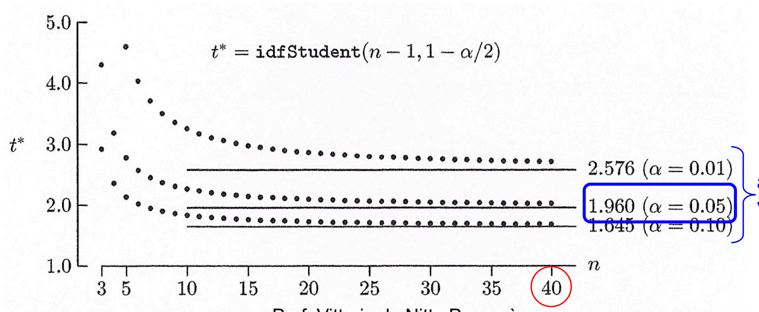

asymptotic value t*

Prof. Vittoria de Nitto Personè 36

36

Example

- Given a reasonable guess for s and a user-specified half-width parameter w, if t^*_∞, is used in place of t^*

  n can be determined by solving $w = \frac{t^* s}{\sqrt{n-1}}$ for n:

$$n = \left\lfloor \left(\frac{t^*_\infty s}{w} \right)^2 \right\rfloor + 1$$

provided $n > 40$

- For example, if s=3.0 and want to estimate μ with 95% confidence to within ± 0.5, a value of $n = 139$ should be used

Qui ciò che cerco è 'n' in funzione di condizioni che impongo.

Example

$$n = \left\lfloor \left(\frac{t^*_\infty s}{w} \right)^2 \right\rfloor + 1$$

- If a reasonable guess for s is not available, w can be specified as a proportion of s thereby eliminating s from the previous equation

- For example, if w is 10% of s and 95% confidence is desired, $n = 385$ should be used to estimate μ to within $\pm w$

($w/s = 0.1$)

Se non sono in grado di calcolare bene 's', allora impongo una proporzione fissata di w/s, cioè scrivo il rapporto, non mi serve sapere s.

See in the book algorithm 8.1.2 to obtain confidence interval starting from the sample $x_1, x_2, \ldots, x_n$ or from the half-width parameter w respectively

The meaning of confidence

Incorrect:

"For this 95% confidence interval, the probability that μ is within this interval is 0.95"

- Why incorrect? ciò che varia è media campionaria x barra, non mu.
  - μ is not a random variable; it is constant (but unknown)
  - the interval endpoints are random

Correct:

"If I create many 95% confidence intervals, approximately 95% of them should contain μ"

E' come lo definivamo in CPS. Ovvero, se creo 100 intervalli, 95 conterranno la media.

Example

- 100 samples of size n=9 drawn from *Normal*(6,3) population
- For each sample, construct a 95% confidence interval
- 95 intervals contain μ=6
- Three intervals "too low", two intervals "too high"

La media è nota perchè conosco la distribuzione, normalmente noi non la sappiamo quando facciamo i nostri studi.

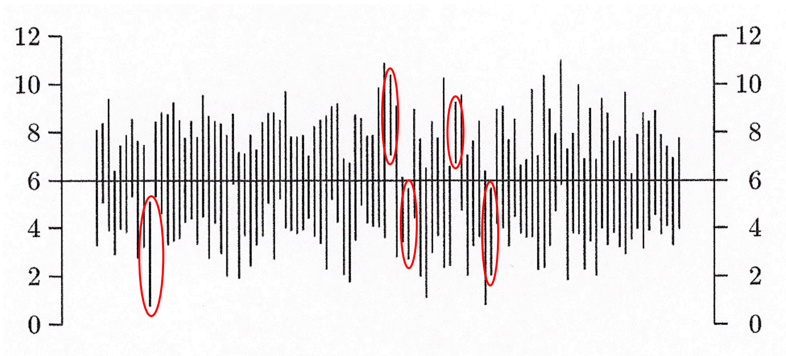

Come vediamo, per molte "linee" vediamo che non sono centrate nella media 6 (ovvero, l'asse y=6 non le divide equamente, anche se 6 è la media). Non è rilevante il punto centrale, bensì la confidenza associata al punto centrale.

Se faccio run, non devo fare intervallo di confidenza su quel run (vedremo i batch means a tal proposito, che in pratica da un run genera più campioni). Non significa nulla, perchè quel run ci fornisce 1 media campionaria significata, 1 elemento del campione. Nell'esempio di prima ne ho tirati 100, da questi 100 calcolo poi l'intervallo.

Exercise

• Exercises 8.1.1, 8.1.5

• Consider case study 1 or case study 2, at your choice. Derive the sample mean histogram from one run (as in the picture in slide 8) and for two different sizes for the samples. Compare the obtained results with reference to the Exponential sample mean histograms seen in this lecture (slide p.10).